

Global Illumination



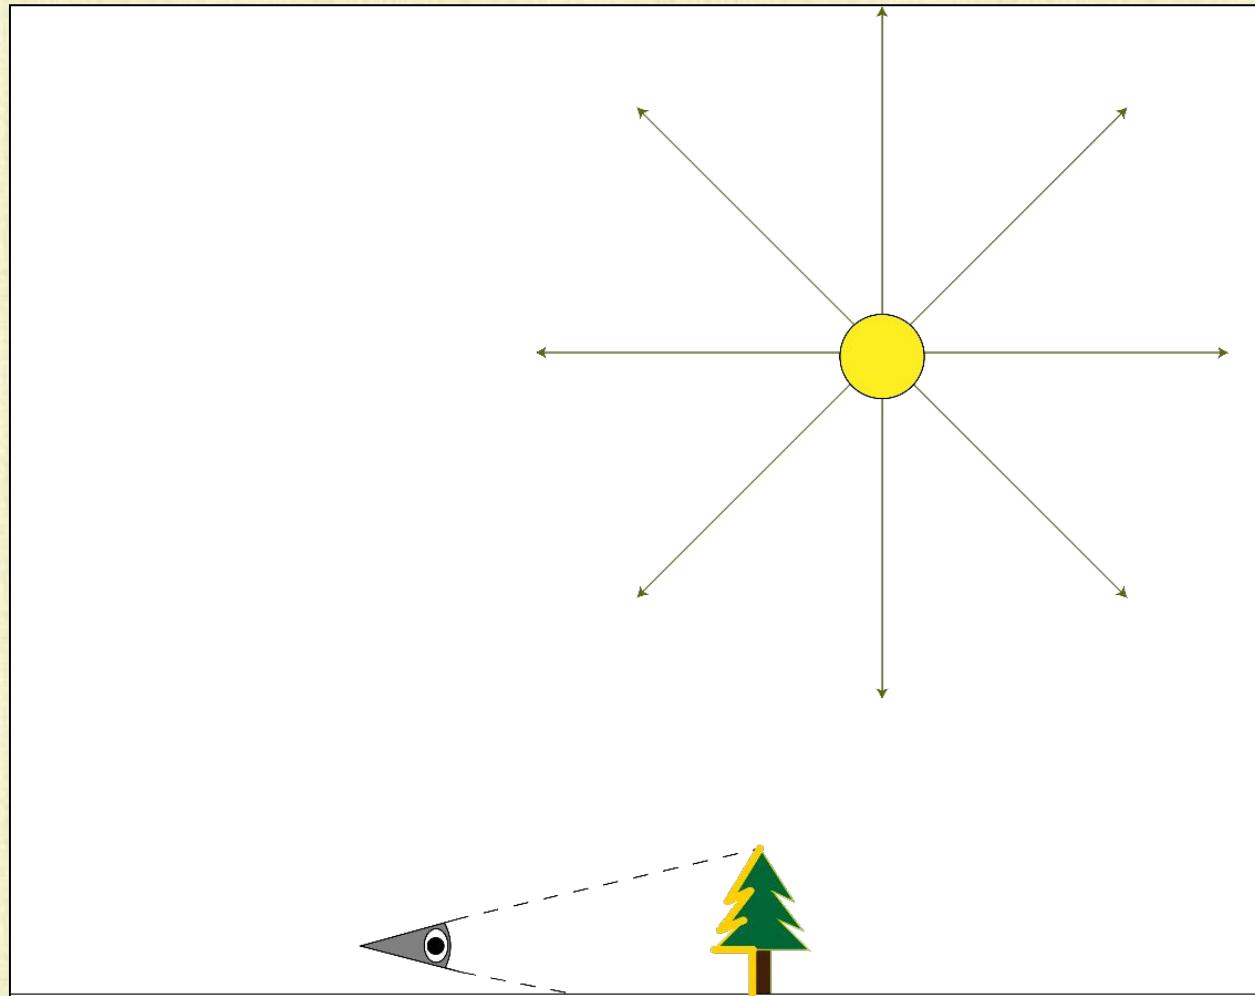


Photon Tracing

- For each light, choose a number of directions on the outgoing hemisphere (or sphere); emit a photon in each direction
- Each photon travels in a straight line, until it intersects an object
- If Absorbed: terminate photon (doesn't get to the film)
- If Reflected/Transmitted/Scattered: photon goes off in a new direction (until it again intersects an object)
- If a photon goes through the camera aperture and hits the film, it contributes to the final image

Photon Tracing

- Most of the light never hits the film (far too inefficient)

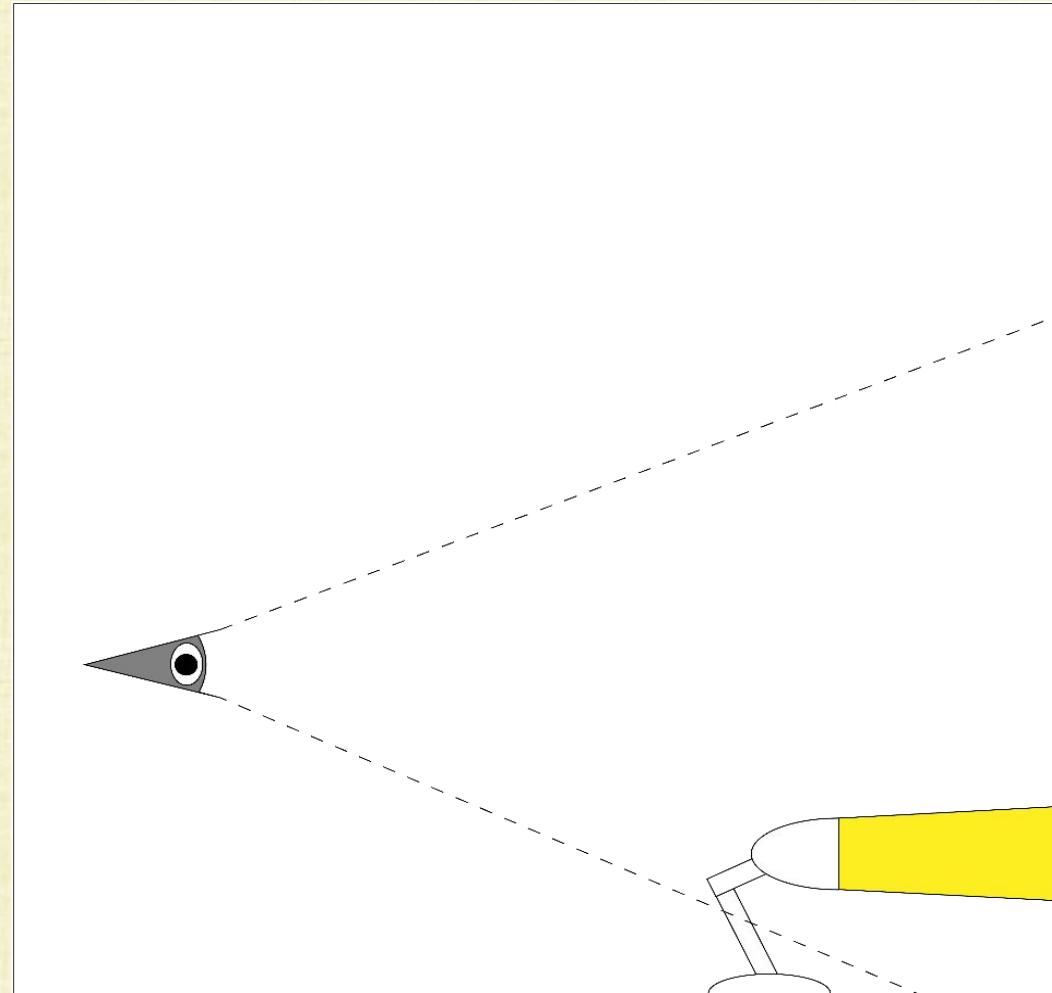


Path Tracing

- For each pixel, send a ray through the aperture to backward trace a photon that would hit the pixel (same as ray tracing)
 - If the ray hits an object, send out reflected/transmitted rays (same as ray tracing)
 - Follow all rays until they hit a light source
-
- When a ray hits a light source, use the outgoing light in the ray direction to determine how many photons hit the pixel
 - A terminated ray **only** gives a path from the pixel to a light source
 - Emit photons along this path, bounce them off all the objects along the path, check to see if absorbed, otherwise continue on to the pixel
 - Some percentage of the photons are absorbed resulting in a specific color/brightness of light hitting the pixel (along that path)

Path Tracing

- Most paths never find their way to the light source (inefficient)



BRDF Path Tracing

- Reflected/Transmitted ray directions **do not** adequately represent the backwards path of incoming photons
- Photons bombard a point on a surface **from every direction** of the hemisphere
- Cast rays in **all directions of the hemisphere**, in order to backwards trace incoming photons incoming from every direction
- Spawning so many rays at every intersection point is impractical
- Every spawned ray that hits a surface spawns an entire hemisphere of rays of its own (exponential growth)

Ray Tracing (an optimization of BRDF Path Tracing)

- Ignore most incoming directions on the hemisphere, only keeping **the most important** ones
- Rays incoming directly from the light source have a lot of photons
 - Shadow rays are used to track this incoming light
 - Using only shadow rays is called direct illumination (only light directly from light sources is considered)
- Reflective objects bounce a lot of photons in the mirror reflection direction
 - Thus, incoming light from the mirror reflection direction is accounted for
- Transparent objects transmit a lot of photons along the transmitted ray direction
 - Thus, incoming light from the transmitted ray direction is accounted for

Bidirectional Ray Tracing

- Combine Photon Tracing and Path Tracing
- Step 1: Emit photons from the light, bathe objects in those photons, record the result in a light map
 - Photons bounce around illuminating shadows, bleeding color, etc.
- Step 2: Ray trace the scene, using the light maps to estimate indirect light (from directions of the hemisphere typically ignored by ray tracers)
- IMPORTANT: Still deal with the most important directions (on the hemisphere) explicitly, for increased accuracy
 - Shadow rays for direct illumination
 - Reflected rays
 - Transmitted rays

Light Maps

- Light maps work great for soft shadows, color bleeding, etc.
- But they can also generate many other interesting effects:



Light Maps

- Advantage:
 - Can track the full path of light, accounting for the multiple bounces in indirect lighting
 - This results in far more realistic images
 - Light maps aren't camera dependent, so they can be pre-computed
 - The light map remains fixed, as the camera moves around
- Disadvantage:
 - Requires consideration of all incoming/outgoing light at every point in the (relevant) world
 - It's a lot more work than just computing lighting at points visible to the camera
 - Computationally expensive

Lighting Equation

- Explicitly add the dependence of L_i on the surface location x and incoming angle ω_i
- Change $i \in in$ for “incoming directions” to $i \in hemi$ for “hemisphere”
- Add an emission term $L_e(x, \omega_o)$, so x can be on the surface of actual lights too

$$L_o(x, \omega_o) = L_e(x, \omega_o) + \int_{i \in hemi} BRDF(x, \omega_i, \omega_o) L_i(x, \omega_i) \cos \theta_i d\omega_i$$

- Incoming light from direction ω_i left some other surface point x' going in direction $-\omega_i$
- So, can replace $L_i(x, \omega_i)$ with $L_o(x', -\omega_i)$

$$L_o(x, \omega_o) = L_e(x, \omega_o) + \int_{i \in hemi} BRDF(x, \omega_i, \omega_o) L_o(x', -\omega_i) \cos \theta_i d\omega_i$$

Implicit Equation

- Computing the outgoing radiance $L_o(x, \omega_o)$ on a particular surface requires knowing the outgoing radiance $L_o(x', -\omega_i)$ from all the other (relevant) surfaces
 - But the outgoing radiance from those other surfaces (typically) depends on the outgoing radiance from the surface under consideration

$$L_o(x, \omega_o) = L_e(x, \omega_o) + \int_{i \in hemi} L_o(x', -\omega_i) BRDF(x, \omega_i, \omega_o) \cos \theta_i d\omega_i$$

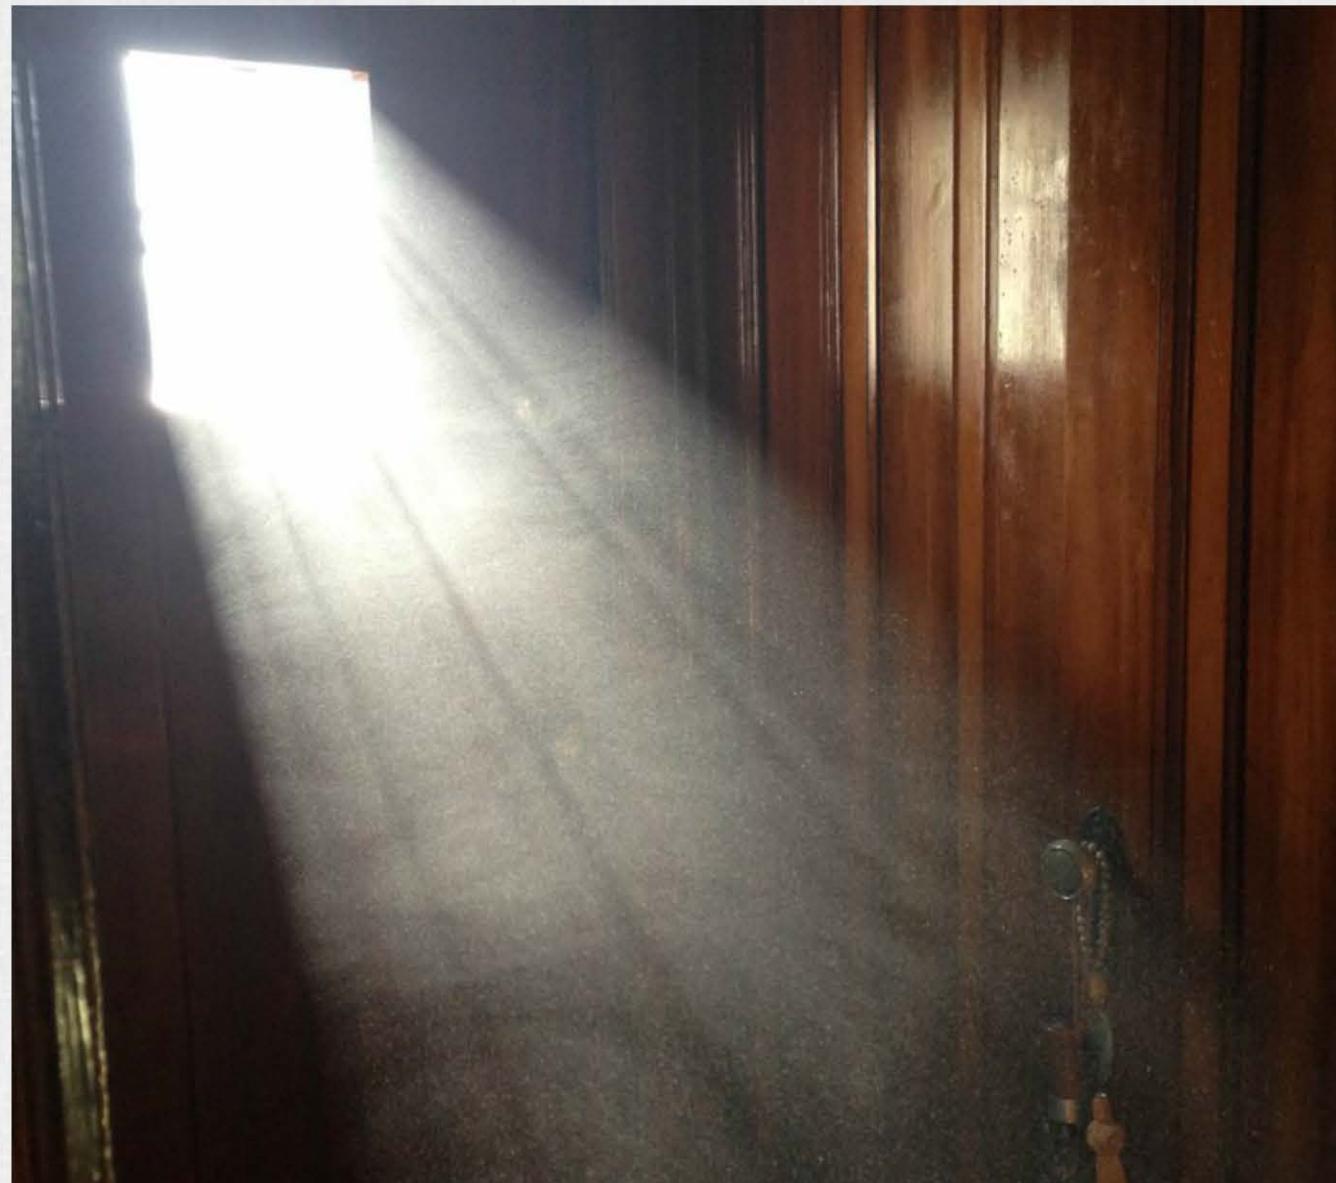
Reflected Light UNKNOWN	Emission KNOWN	Reflected Light UNKNOWN	BRDF KNOWN	incident angle KNOWN
----------------------------	-------------------	----------------------------	---------------	-------------------------

- Fredholm Integral Equation of the second kind (extensively studied) given in canonical form with kernel $k(u, v)$ by:

$$l(u) = e(u) + \int l(v) k(u, v) dv$$

Aside: Participating Media

- The “air” typically contains participating media (e.g. dust, droplets, smoke, etc.)
- So, L should be defined over all of 3D space
- Then, the incoming light would be considered in a sphere centered around each point in 3D
- Neglecting this makes the assumption that “air” is a vacuum
- This restricts L to surfaces

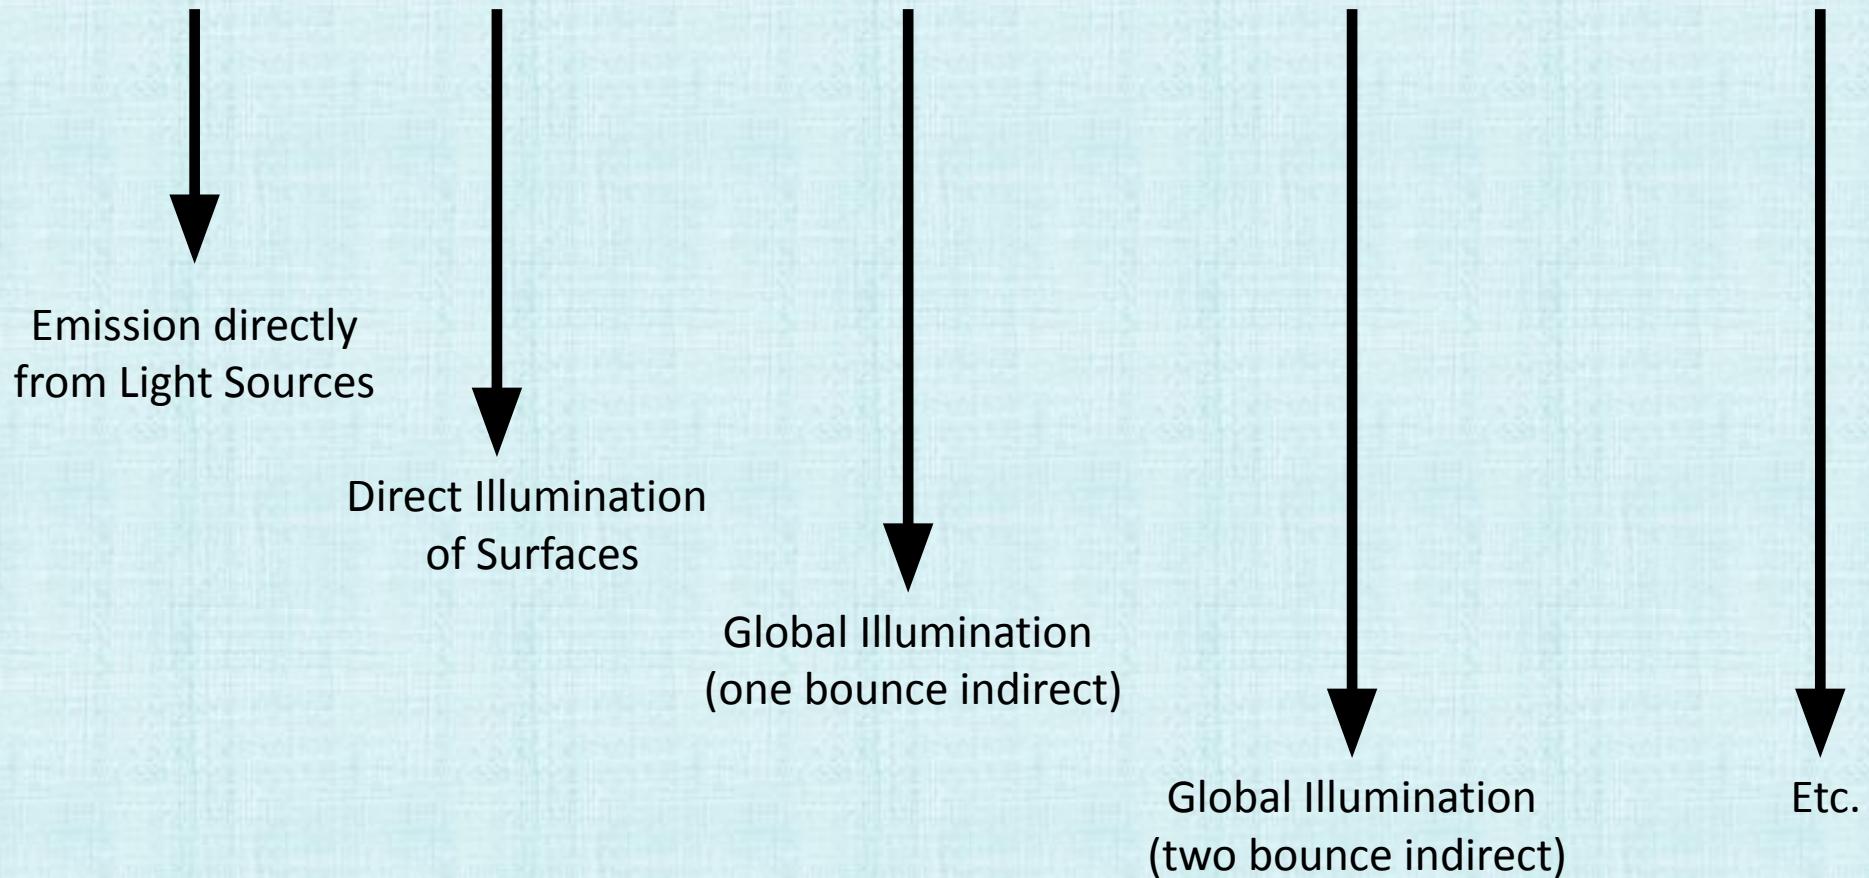


Discretization

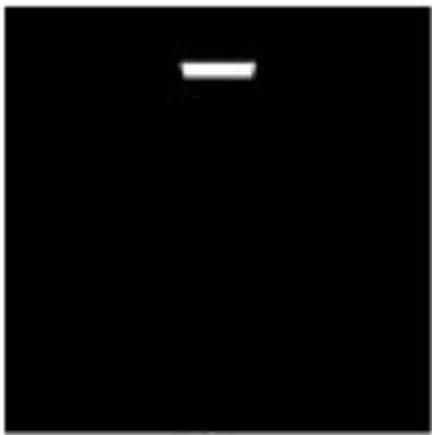
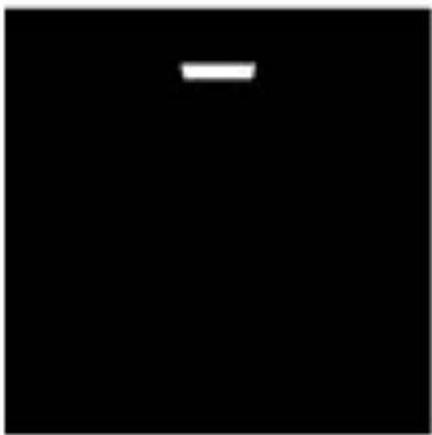
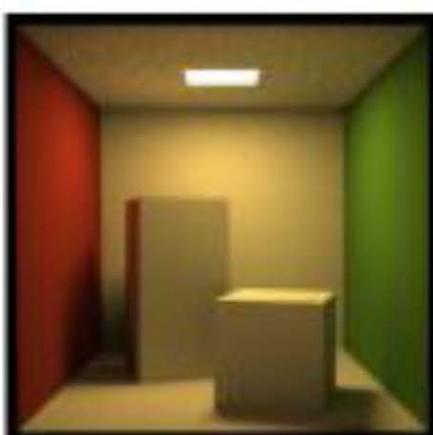
- Choose p points, each representing a chunk of surface area (or chunk of volume for participating media), which is a 2D (or 3D) discretization
- For each of the p points, choose q outgoing directions (each representing a chunk of solid angles of the hemisphere/sphere), which is a 2D discretization
 - q can vary from surface chunk to surface chunk (if desired)
- Then, L_o and L_e each have $p * q$ unknowns (a 4D or 5D discretization)
- So, they can be represented by vectors: L and E (each with length $p * q$)
- The light transport (kernel) matrix K has size $(p * q)$ by $(p * q)$
- The (linear) system of equations is: $L = E + KL$ or $(I - K)L = E$
- Solution: $L = (I - K)^{-1}E$
- Using the Binomial Theorem: $L = (I + K + K^2 + \dots)E$
- K only bounces a fraction of the light (the rest is absorbed), so higher powers are smaller (and the series can be truncated)

Power Series

$$L = E + KE + K^2E + K^3E + \dots$$



Power Series

 L_e  $K \circ L_e$  $K \circ K \circ L_e$  $K \circ K \circ K \circ L_e$  L_e  $L_e + K \circ L_e$  $L_e + \dots K^2 \circ L_e$  $L_e + \dots K^3 \circ L_e$

Tractability

- One might (typically) consider thousands or tens of thousands of area chunks
 - So, p could be $1e3, 1e4, 1e5, 1e6$, etc.
 - Incoming light could vary significantly across the hemisphere
 - So, q might need to be $1e2, 1e3, 1e4$, etc.
 - Then, L and E would range in length from $1e5$ to $1e10$
 - The matrix K would range in size from $1e5$ by $1e5$ up to $1e10$ by $1e10$
- K would have between $1e10$ and $1e20$ entries!
- This tractability analysis is for the 4D problem (5D is even worse)
- The curse of dimensionality makes problems in 4 and 5 dimensions (and higher) hard to discretize (with straightforward quadrature)

Addressing Tractability

- Idea: separate the diffuse and specular contributions (to be treated separately)

Diffuse:

- First, assuming all materials are purely diffuse (i.e. no specular contributions)
- Compute the view-independent global illumination for the entire scene
- This can be done in a pre-processing step

Specular:

- Compute (view-dependent) specular illumination on-the-fly as the camera moves
- Using Phong (or any other model)

Radiosity and Albedo

- **Radiosity**: power per unit surface area leaving a surface (similar to irradiance, but outgoing instead of incoming):

$$B(x) = \frac{d\Phi}{dA} = \int_{hemi} L_o(x, \omega_o) \cos \theta_o d\omega_o$$

- When L_o is independent of ω_o (i.e. diffuse):

$$B(x) = \frac{d\Phi}{dA} = L(x) \int_{hemi} \cos \theta_o d\omega_o = \pi L(x)$$

- **Albedo**: a “reflection coefficient” relating incoming light hitting a surface patch (irradiance E_i) to outgoing light emitted in all possible directions

$$\rho(x) = \int_{hemi} BRDF(x, \omega_o, \omega_i) \cos \theta_o d\omega_o$$

- When the BRDF is independent of ω_o and ω_i (i.e. diffuse):

$$\rho(x) = BRDF(x) \int_{hemi} \cos \theta_o d\omega_o = \pi BRDF(x)$$

(Diffuse) Lighting Equation

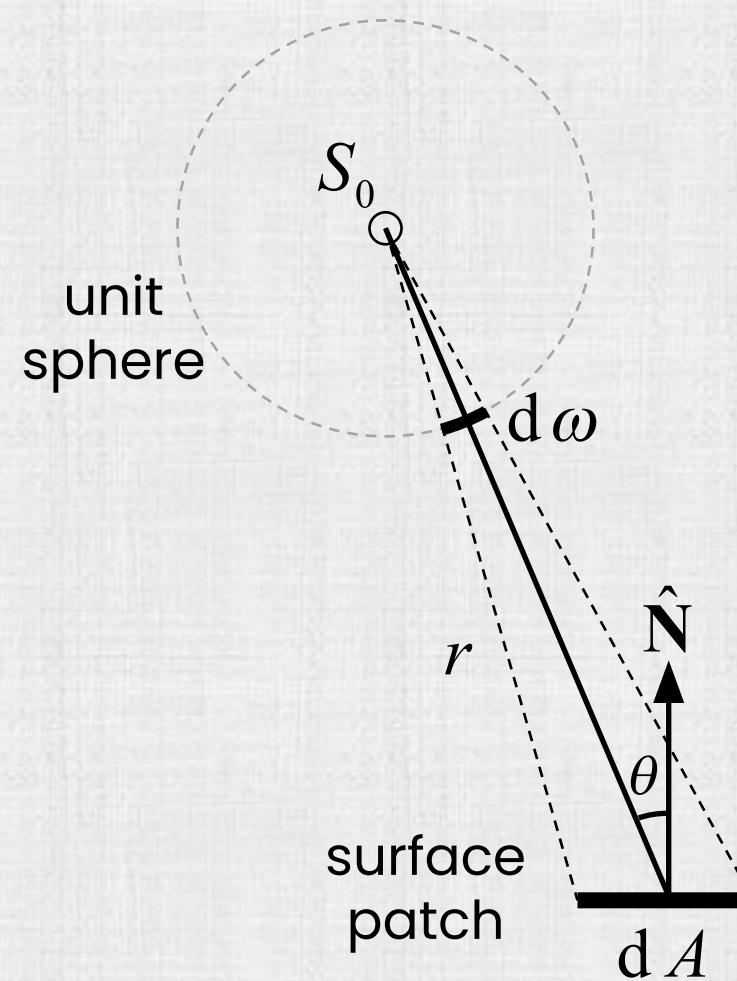
- Given $L_o(x, \omega_o) = L_e(x, \omega_o) + \int_{i \in hemi} L_o(x', -\omega_i) BRDF(x, \omega_i, \omega_o) \cos \theta_i d\omega_i$, multiply through by $\cos \theta_o d\omega_o$ and integrate over the hemisphere (i.e. $d\omega_o$) to obtain:

$$B(x) = E(x) + \int_{i \in hemi} B(x') BRDF(x, \omega_i, \omega_o) \cos \theta_i d\omega_i$$

- B is a 2D function (of x), whereas L was a 4D function (of x and ω_o)
- Then, assume that all surfaces have a diffuse BRDF independent of angle:

$$B(x) = E(x) + \frac{\rho(x)}{\pi} \int_{i \in hemi} B(x') \cos \theta_i d\omega_i$$

Recall: Solid Angle vs. Cross-Sectional Area



Interchange Solid Angle and Surface Area

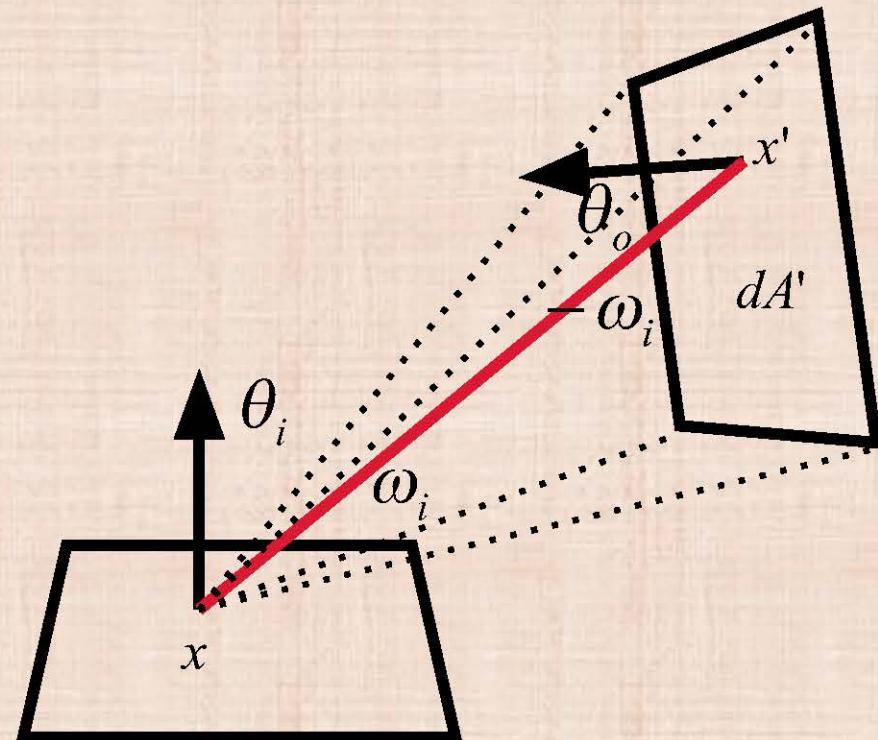
- Note: $d\omega = \frac{dA \cos\theta}{r^2}$ gives $d\omega_i = \frac{dA' \cos\theta_o}{\|x-x'\|_2^2}$

- Then, $B(x) = E(x) + \frac{\rho(x)}{\pi} \int_{i \in hemi} B(x') \cos\theta_i d\omega_i$ is:

$$B(x) = E(x) + \rho(x) \int_{i \in hemi} B(x') \frac{\cos\theta_i \cos\theta_o}{\pi \|x - x'\|_2^2} dA'$$

- Let $V(x, x') = 1$ when x and x' are mutually visible (and $V(x, x') = 0$ otherwise), so:

$$B(x) = E(x) + \rho(x) \int_{all \ x'} B(x') V(x, x') \frac{\cos\theta_i \cos\theta_o}{\pi \|x - x'\|_2^2} dA'$$



A Tractable Discretization

- Choose p points, each representing a chunk of surface area (a 2D discretization)
- Then $B_i = E_i + \rho_i \sum_{j \neq i} B_j F_{i,j}$ with a purely geometric $F_{i,j} = V(x_i, x_j) \frac{\cos \theta_i \cos \theta_j}{\pi \|x_i - x_j\|_2^2} A_j$
- Rearrange to $B_i - \rho_i \sum_{j \neq i} B_j F_{i,j} = E_i$ and put into matrix form:

$$\begin{pmatrix} 1 & -\rho_1 F_{12} & \cdots & -\rho_1 F_{1p} \\ -\rho_2 F_{21} & 1 & \cdots & -\rho_2 F_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ -\rho_p F_{p1} & -\rho_p F_{p2} & \cdots & 1 \end{pmatrix} \begin{pmatrix} B_1 \\ B_2 \\ \vdots \\ B_p \end{pmatrix} = \begin{pmatrix} E_1 \\ E_2 \\ \vdots \\ E_p \end{pmatrix}$$

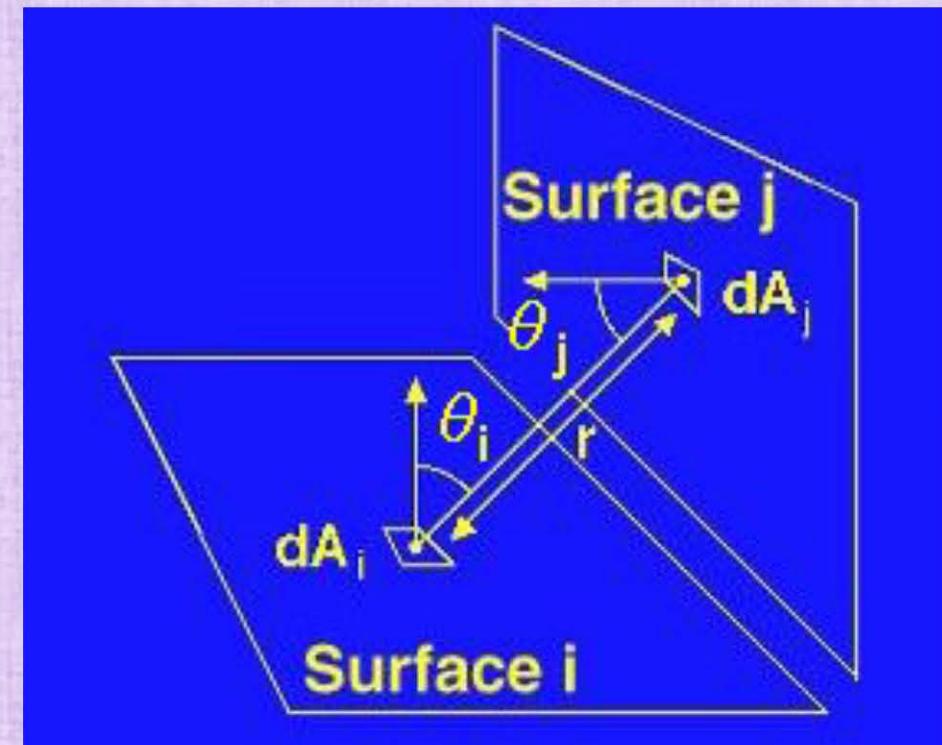
- For p ranging from 1e3 to 1e6, B and E have the same range (while the matrix has 1e6 to 1e12 entries, still big but 1e4 to 1e8 times smaller than before)

Form Factor

- Computing $F_{i,j} = V(x_i, x_j) \frac{\cos \theta_i \cos \theta_j}{\pi \|x_i - x_j\|_2^2} A_j$ requires checking visibility (between x_i and x_j) and evaluating the (symmetric) form factor:

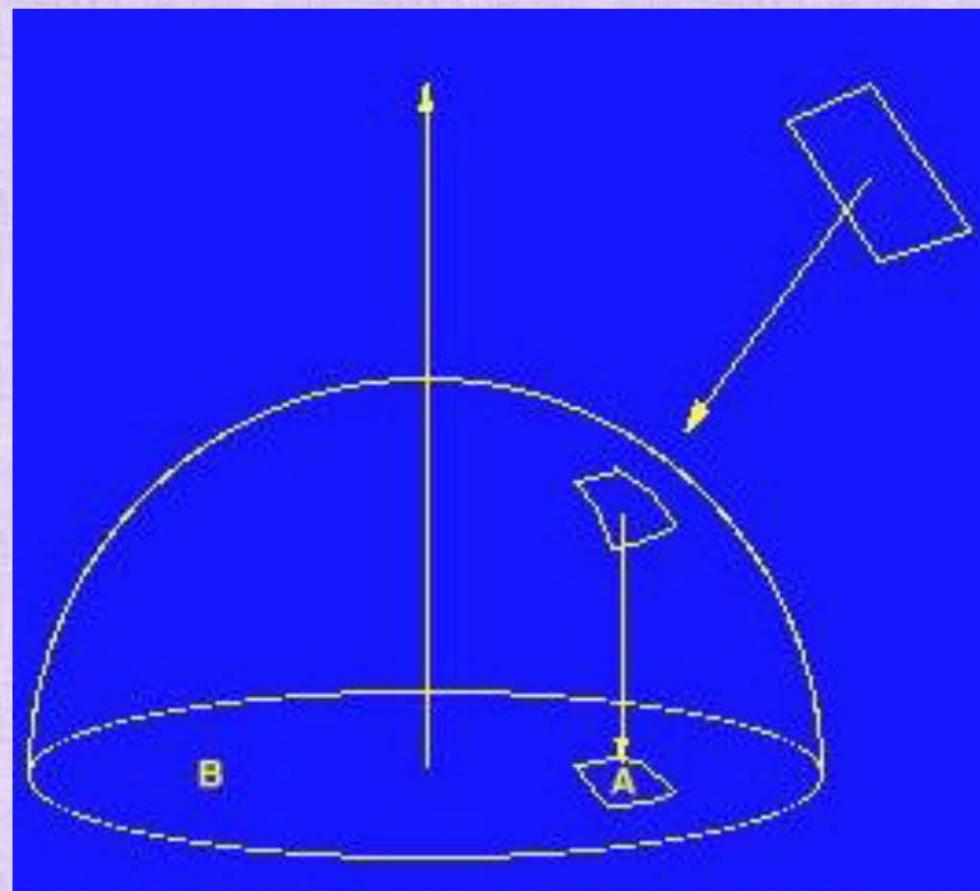
$$\hat{F}_{i,j} = \frac{\cos \theta_i \cos \theta_j}{\pi \|x_i - x_j\|_2^2} A_i A_j$$

- Then $F_{i,j} = V(x_i, x_j) \frac{\hat{F}_{i,j}}{A_i}$ and $F_{j,i} = V(x_i, x_j) \frac{\hat{F}_{i,j}}{A_j}$
- $\hat{F}_{i,j}$ is the fraction of energy leaving one surface that reaches another (and only depends on the geometry, not the light)
- Note: $V(x_i, x_j)$ can be included in $\hat{F}_{i,j}$ (if desired)



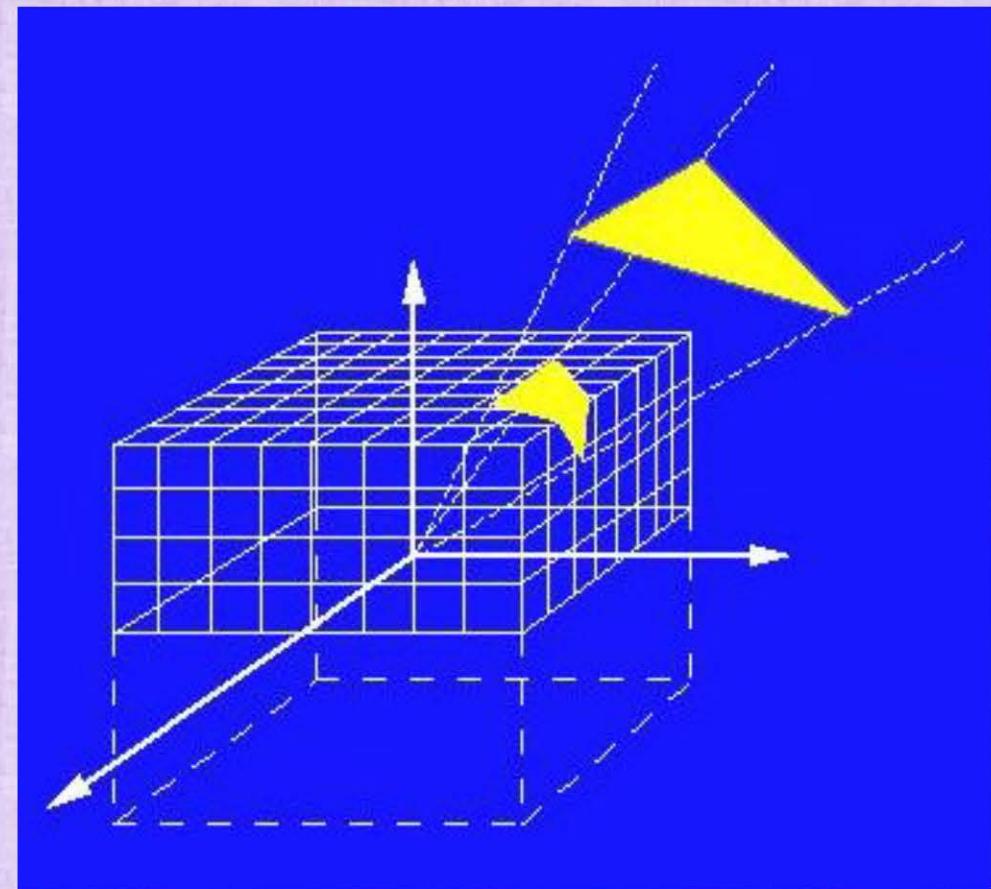
Understanding the Form Factor

- Place a unit hemisphere at a surface point x_i
- Project the other surface onto the hemisphere, noting that $d\omega = \frac{dA \cos\theta}{r^2}$ gives $\frac{A_j \cos\theta_j}{\|x_i - x_j\|_2^2}$ as the result
 - Project the result downwards onto the circular base of the hemisphere, which multiples by $\cos\theta_i$ (recall $\int_{i \in hemi} \cos\theta_i d\omega_i = \pi$, the area of the unit circle)
 - Divide the result by the total area π to get the fraction of the circle occupied
- Overall, this gives: $F_{ij} = \frac{\cos\theta_i \cos\theta_j}{\pi \|x_i - x_j\|_2^2} A_j$

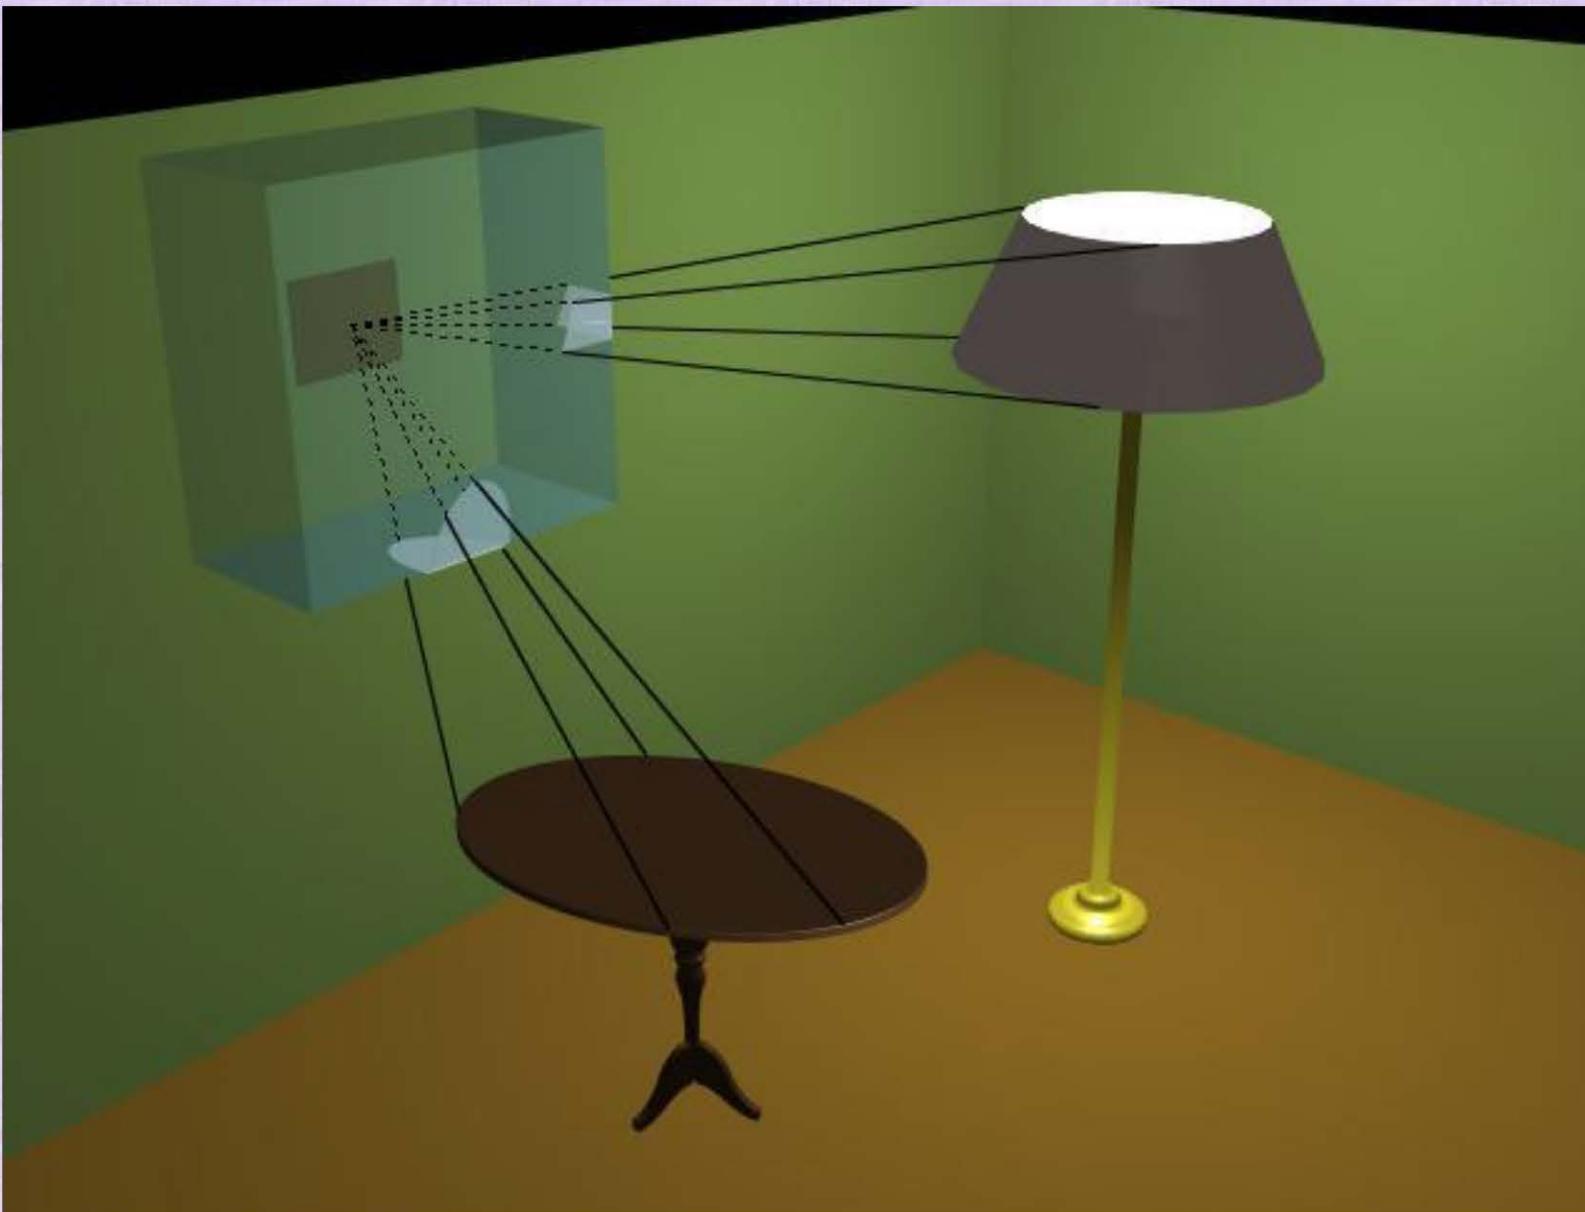


Implementation

- Create a hemicube, and divide each face into sub-squares (as small as desired)
 - For each sub-square, use hemisphere projection (from the last slide) to pre-compute its contribution to F_{ij}
- Place the hemicube at a surface point x_i
- A surface patch (from another object) is projected onto the hemicube, and F_{ij} is approximated, using the pre-computed values for the sub-squares
- The hemicube faces can be treated as image planes and the sub-squares as pixels, making hemicube projection equivalent to scanline rasterization
- The depth buffer can be used to detect occlusions making the visibility term readily computable



Hemicube Scanline Rasterization



Iterative Solvers

- For large matrices, iterative solvers are typically far more accurate than direct methods (that compute an inverse)
- Iterative methods start with an initial guess, and subsequently iteratively improve it

- Consider $\begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} 8 \\ 10 \end{pmatrix}$ with exact solution $\begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} 2 \\ 4 \end{pmatrix}$

- Start with an initial guess of $\begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$

- Jacobi iteration (solve both equations using the current guess):

- $x^{new} = \frac{8-y^{old}}{2}$ and $y^{new} = \frac{10-x^{old}}{2}$

- Gauss Seidal iteration (always use the most up to date values):

- $x^{current} = \frac{8-y^{current}}{2}$ and $y^{current} = \frac{10-x^{current}}{2}$

Jacobi vs. Gauss-Seidal

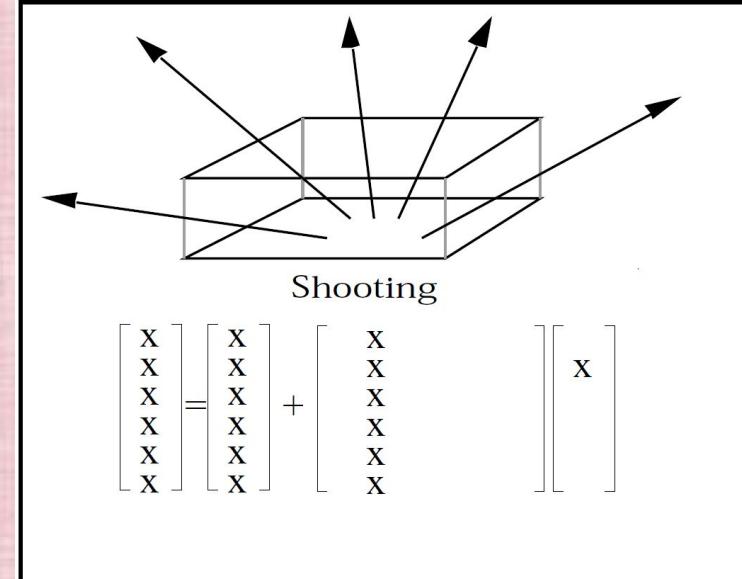
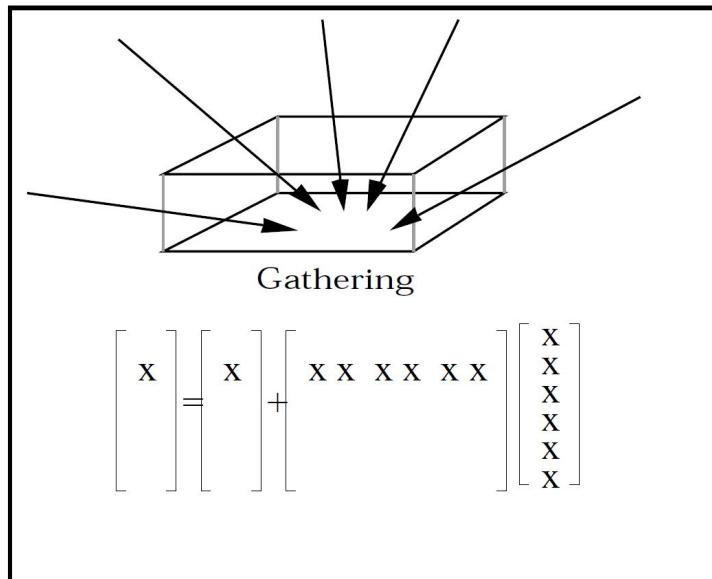
Iteration	Jacobi		Gauss Seidel	
	x	y	x	y
1	0	0	0	0
2	4	5	4	3
3	1.5	3	2.5	3.75
4	2.5	4.25	2.125	3.9375
5	1.875	3.75	2.03125	3.984375
6	2.125	4.0625	2.007813	3.996094
7	1.96875	3.9375	2.001953	3.999023
8	2.03125	4.015625	2.000488	3.999756
9	1.9921875	3.984375	2.000122	3.999939
10	2.0078125	4.00390625	2.000031	3.999985
11	1.998046875	3.99609375	2.000008	3.999996
12	2.001953125	4.000976563	2.000002	3.999999
13	1.999511719	3.999023438	2	4
14	2.000488281	4.000244141	2	4
15	1.99987793	3.999755859	2	4
16	2.00012207	4.000061035	2	4
17	1.999969482	3.999938965	2	4
18	2.000030518	4.000015259	2	4
19	1.999992371	3.999984741	2	4
20	2.000007629	4.000003815	2	4

Better Initial Guess

Iteration	Jacobi		Gauss Seidal	
	x	y	x	y
1	2	3	2	3
2	2.5	4	2.5	3.75
3	2	3.75	2.125	3.9375
4	2.125	4	2.03125	3.984375
5	2	3.9375	2.007813	3.996094
6	2.03125	4	2.001953	3.999023
7	2	3.984375	2.000488	3.999756
8	2.0078125	4	2.000122	3.999939
9	2	3.99609375	2.000031	3.999985
10	2.001953125	4	2.000008	3.999996
11	2	3.999023438	2.000002	3.999999
12	2.000488281	4	2	4
13	2	3.999755859	2	4
14	2.00012207	4	2	4
15	2	3.999938965	2	4
16	2.000030518	4	2	4
17	2	3.999984741	2	4
18	2.000007629	4	2	4
19	2	3.999996185	2	4
20	2.000001907	4	2	4

Iterative Radiosity

- Gathering - update one surface by collecting light energy from all surfaces
- Shooting - update all surfaces by distributing light energy from one surface
- Sorting and Shooting - choose the surface with the greatest un-shot light energy and use shooting to distribute it to other surfaces
 - start by shooting light energy out of the lights onto objects (the brightest light goes first)
 - then the object that would reflect the most light goes next, etc.
- Sorting and Shooting with Ambient - start with an initial guess for ambient lighting and do sorting and shooting afterwards



Iterative Radiosity

