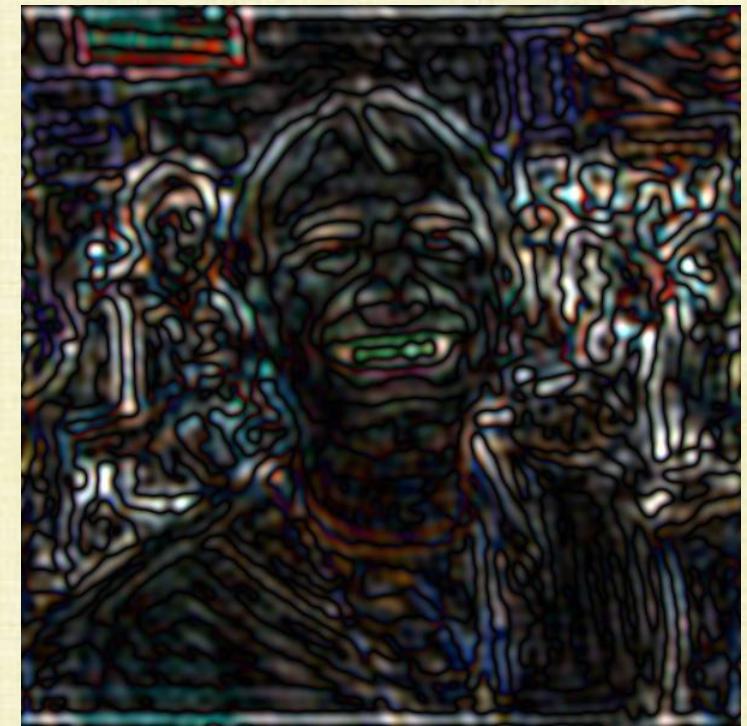


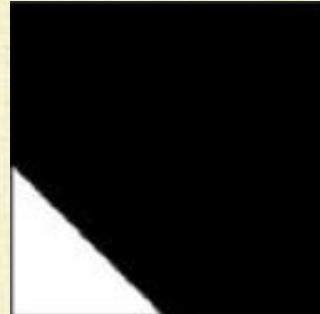
# Sampling



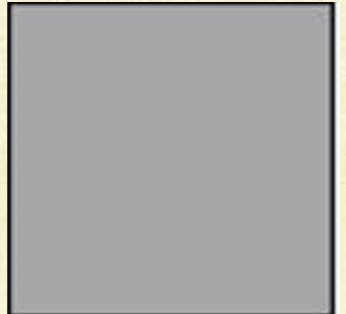
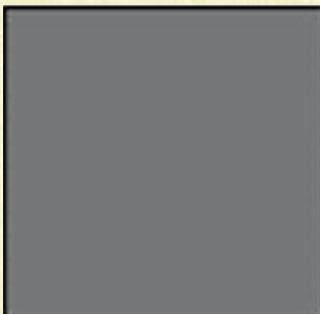
# Area-Coverage

- Real-world sensors get a signal based on the area fractions of the sensor “covered” by objects

Coverage:



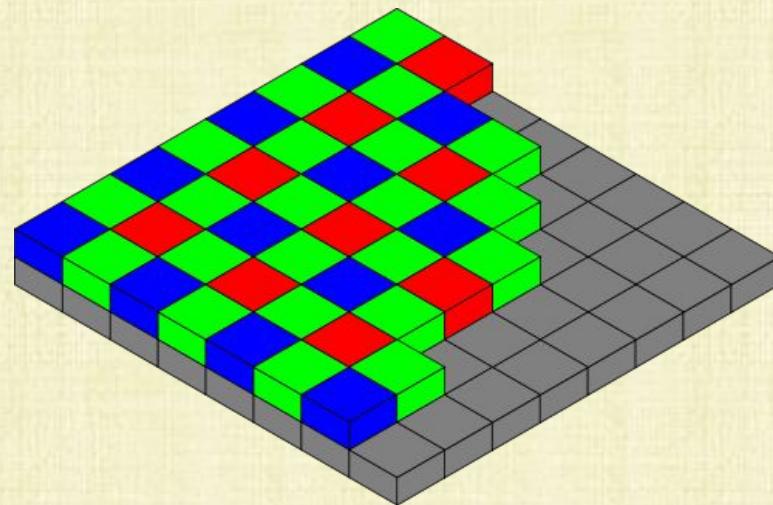
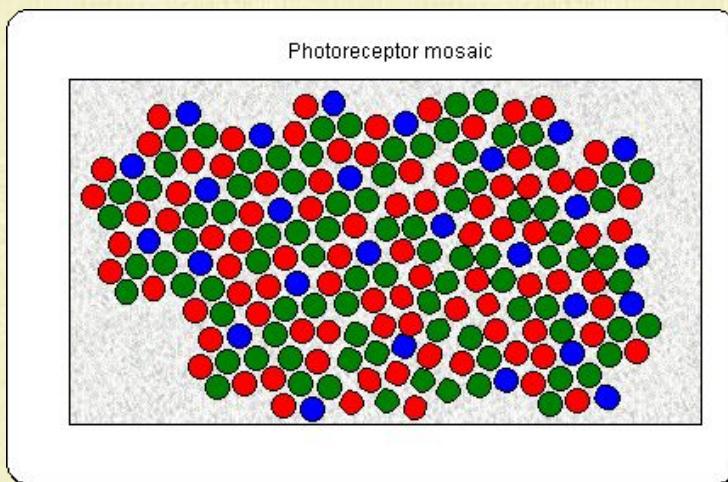
Signal:



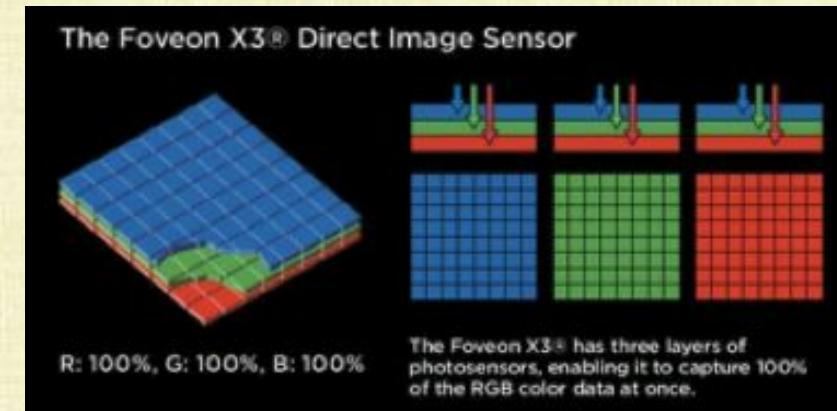
- A ray tracer **only** gets samples of the geometry (using ray-geometry intersection points)
- A scanline renderer projects the entire triangle onto the image plane
  - Testing pixel centers against triangles **only** collects sample information on geometry
  - Computing area overlap between triangles and (square) pixels would better mimic real-world sensors

# Missing Information

- Eyes/cameras don't collect all of the information either
- The staggered spatial layout of real-world sensors means that large regions lack information for certain wavelengths (layered approaches can help to circumvent this)

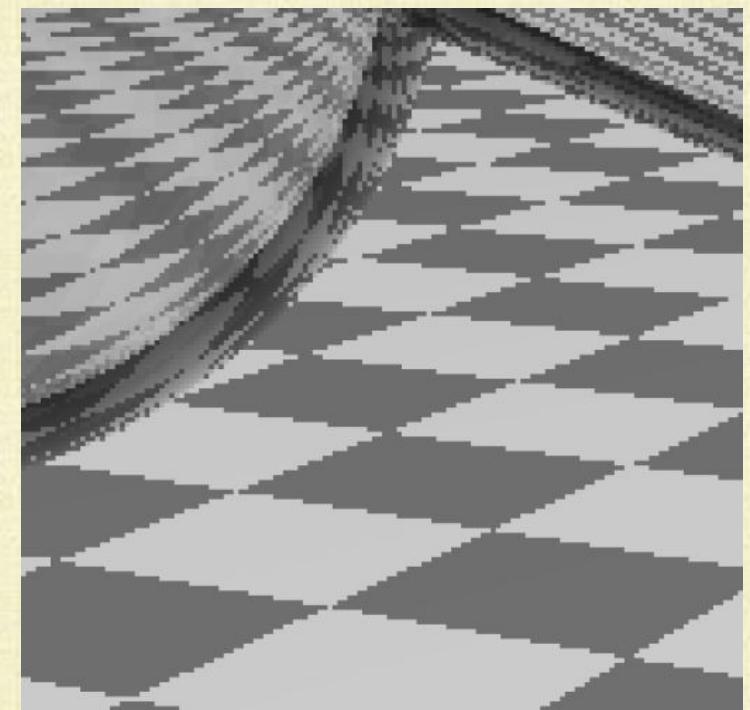
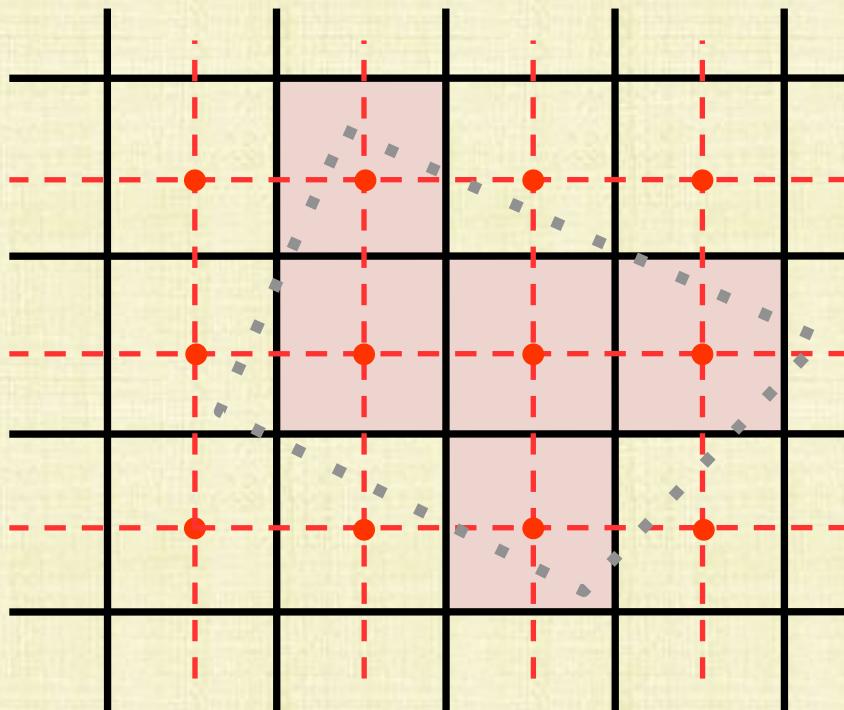


layered approaches can help to circumvent this:



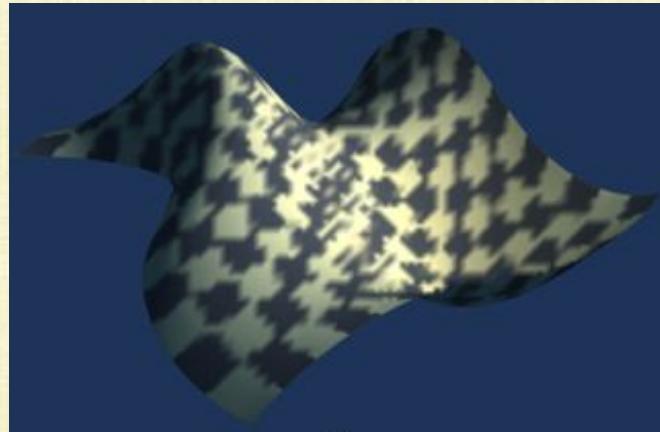
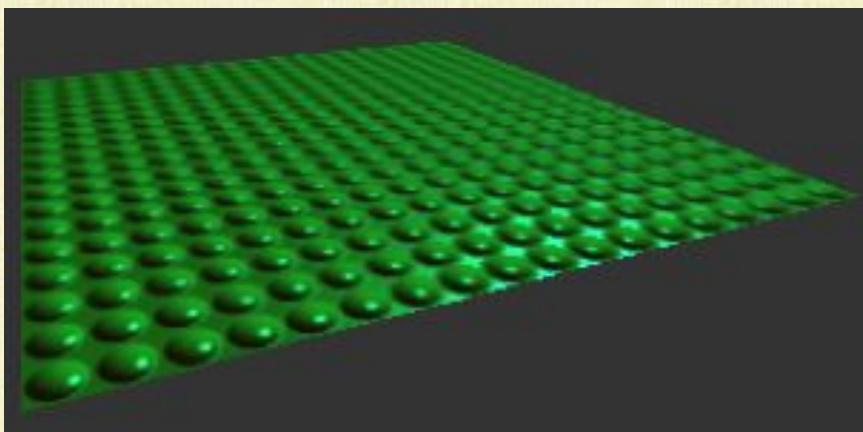
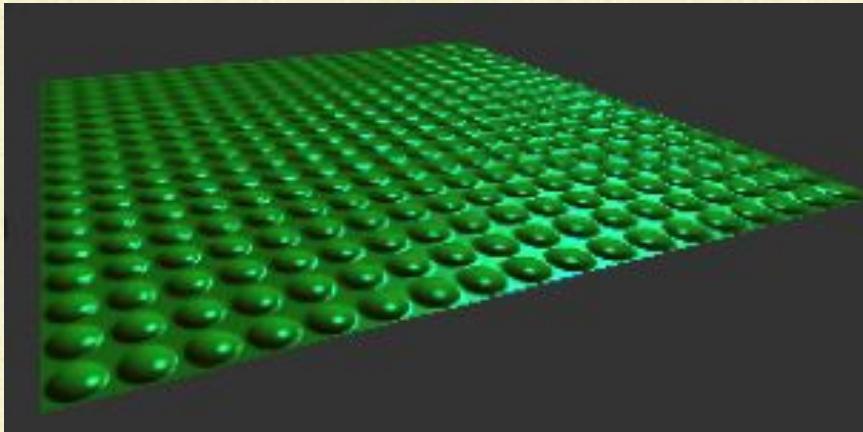
# Aliasing

- Testing **only** the pixel center (with ray-tracing or rasterization) leads to jagged edges
- This causes aliasing artifacts (an alias/impostor replaces the correct feature)
- A jagged line appears instead of the correct straight line
- Anti-aliasing strategies reduce aliasing artifacts caused by sampling information



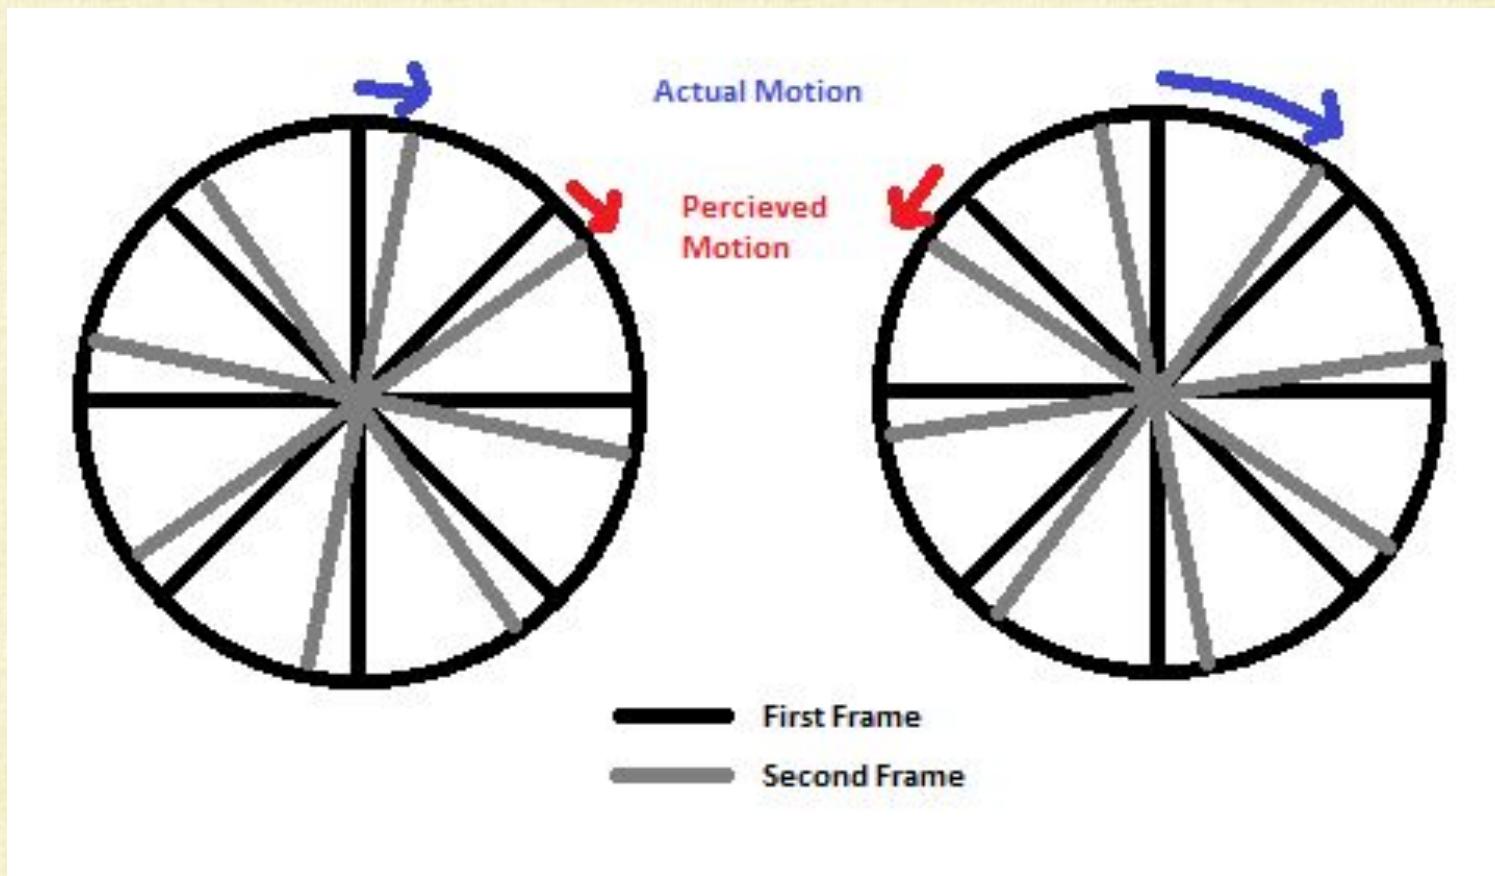
# Aliasing: Shaders & Textures

- Aliased normal vectors can cause erroneous sparkling highlights (top left)
- Aliasing can occur when texture mapping objects too (top right)



# Temporal Aliasing

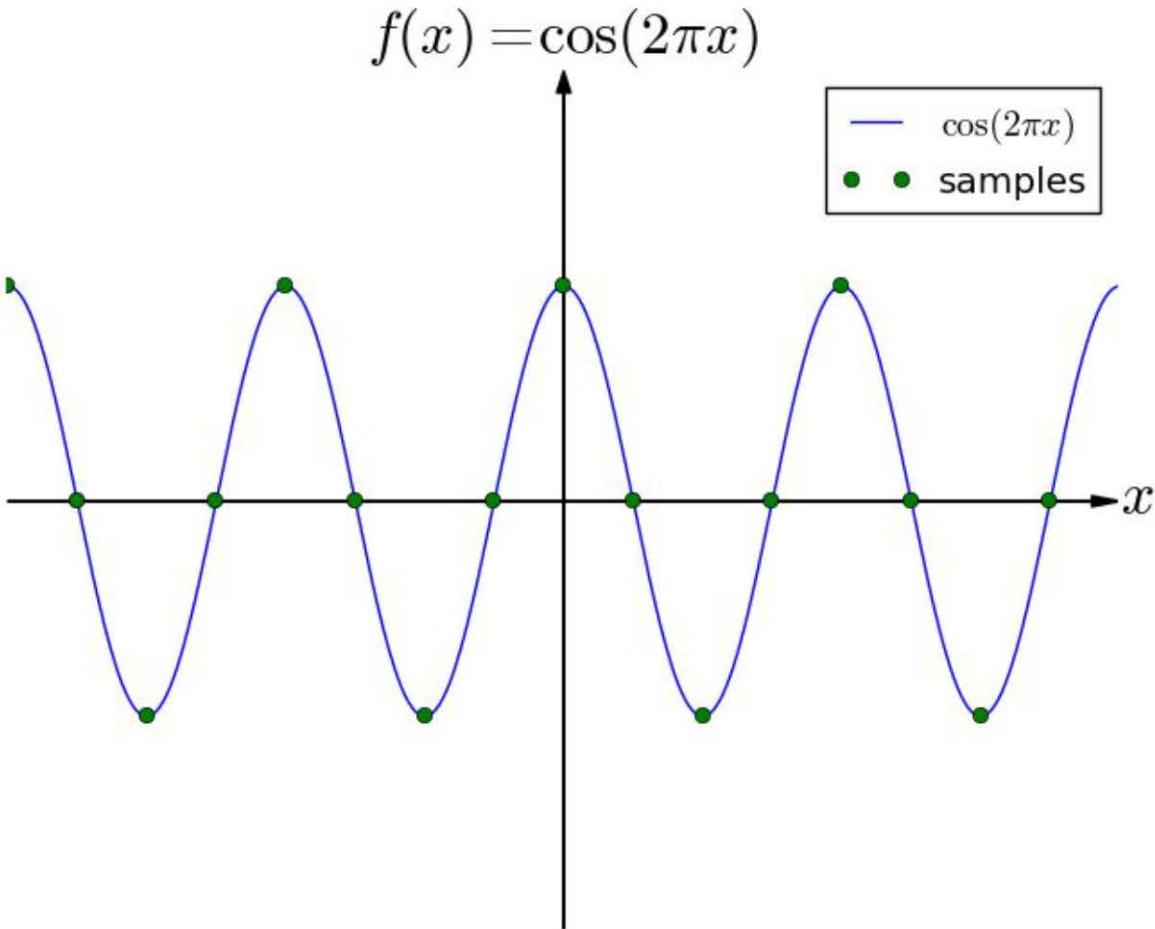
- A spinning wheel can appear to spin backwards, when the motion is insufficiently sampled in time (“wagon wheel” effect)



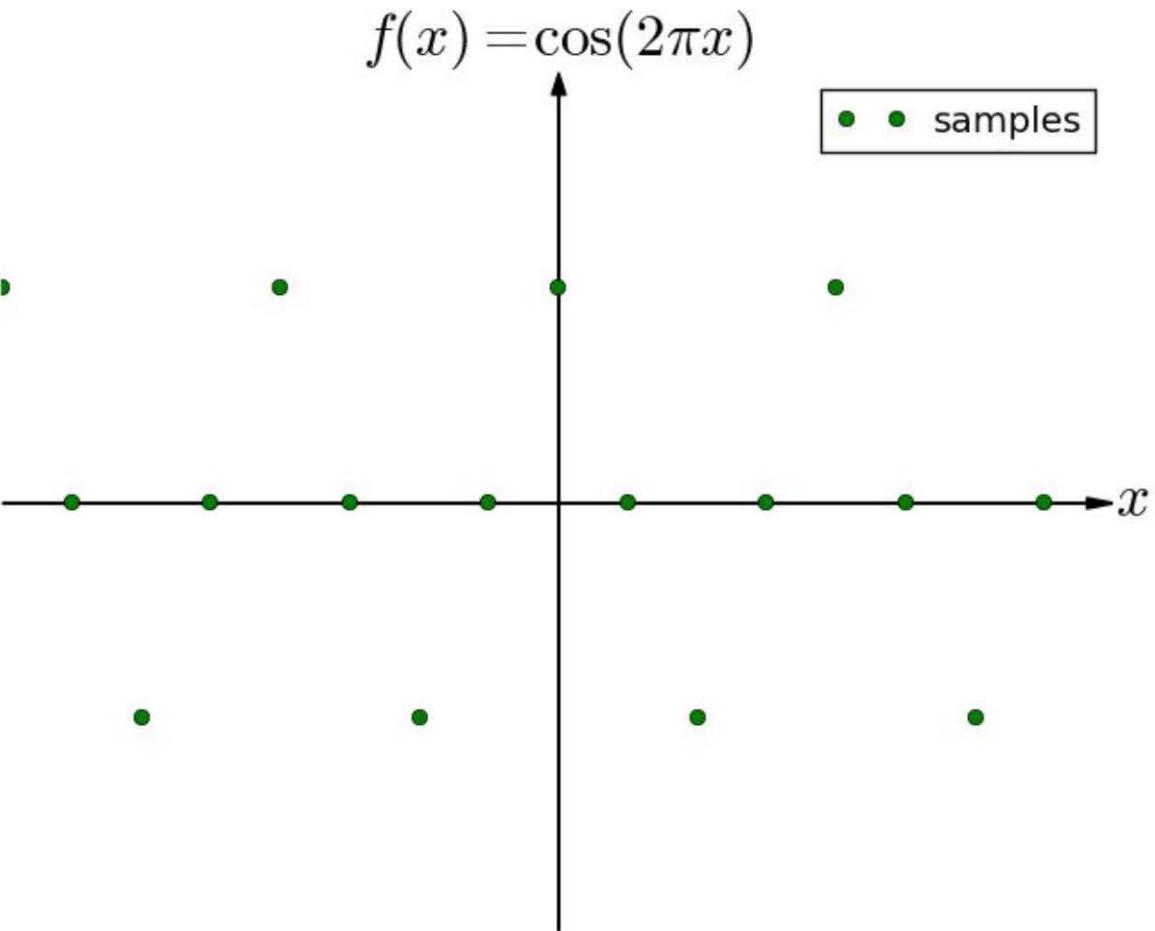
# Sampling Rate

- Artifacts can be reduced by increasing the number of samples (per unit area)
- Thus, one might increase the number of pixels in the image; but:
  - It takes longer to render the scene (because the number of rays increases)
  - Displaying higher-resolution images requires additional storage/computation
- **Optimize the Sample Rate!**
- Use the lowest possible sampling rate that does not result in noticeable artifacts
- What is the optimal sampling rate?

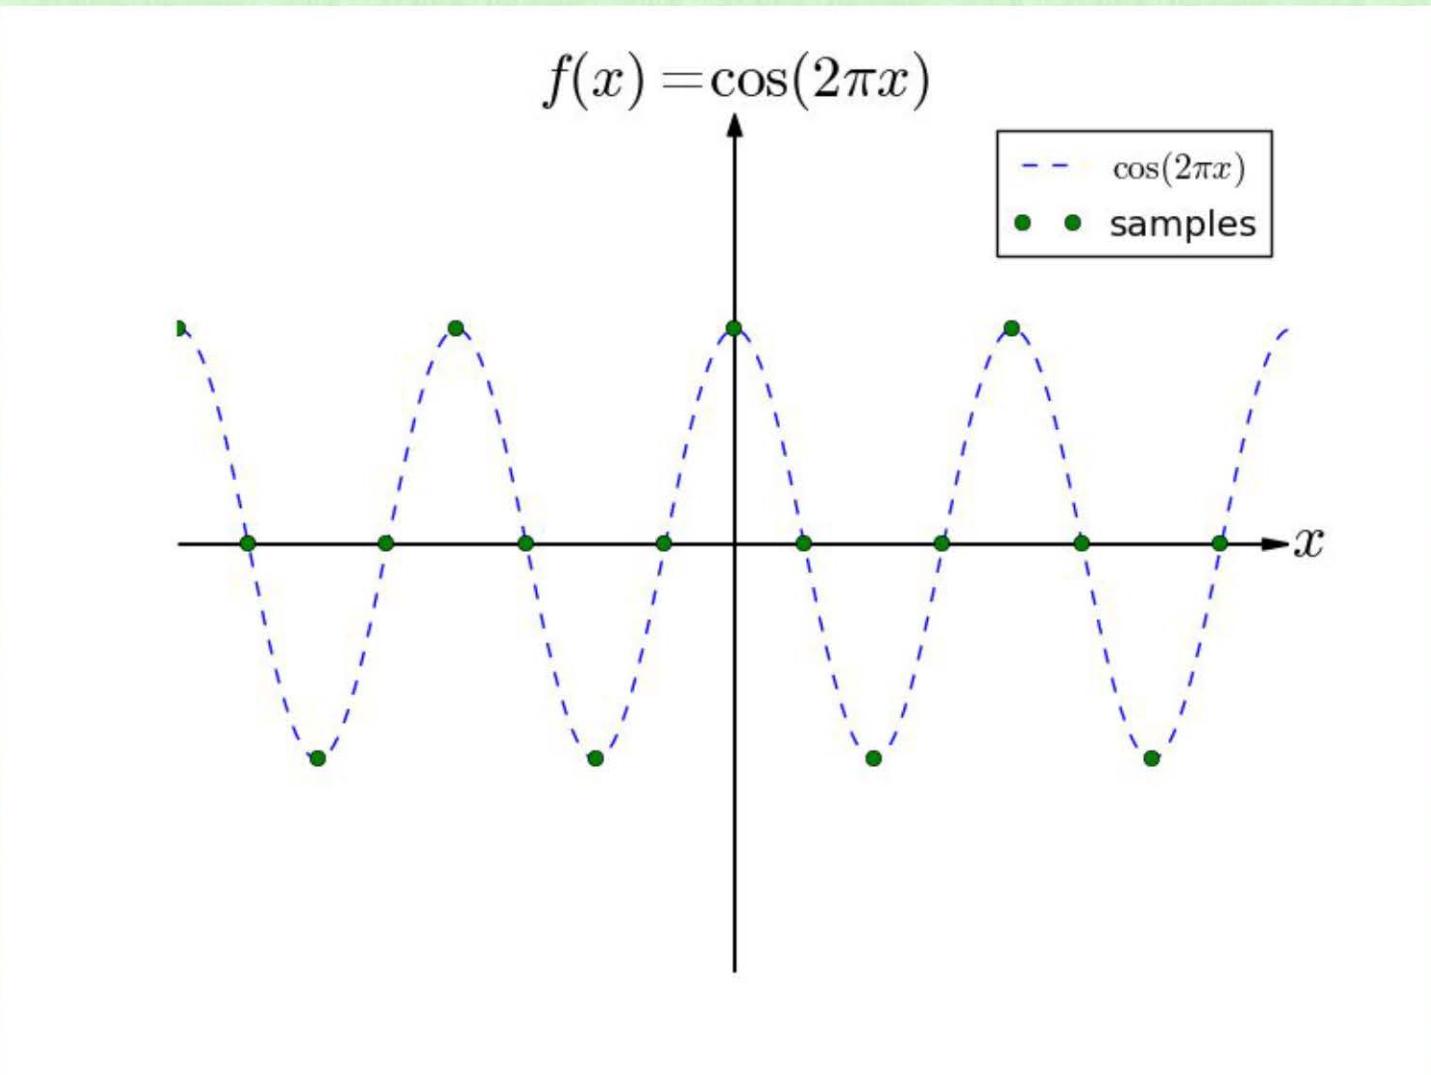
# 4 samples per period



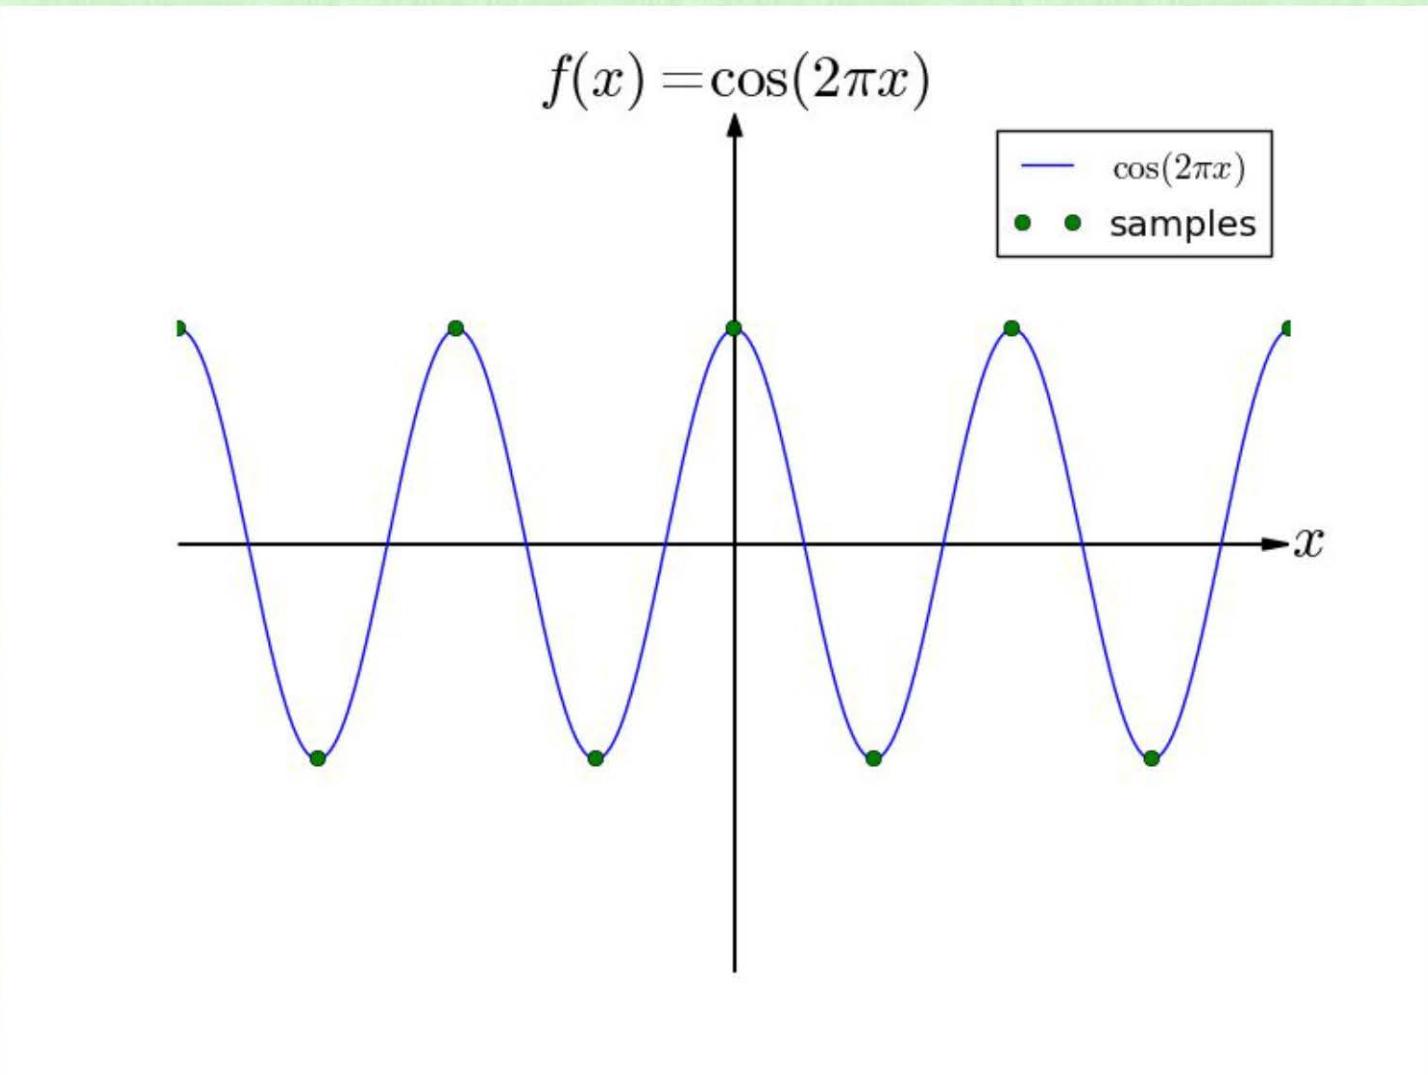
# samples



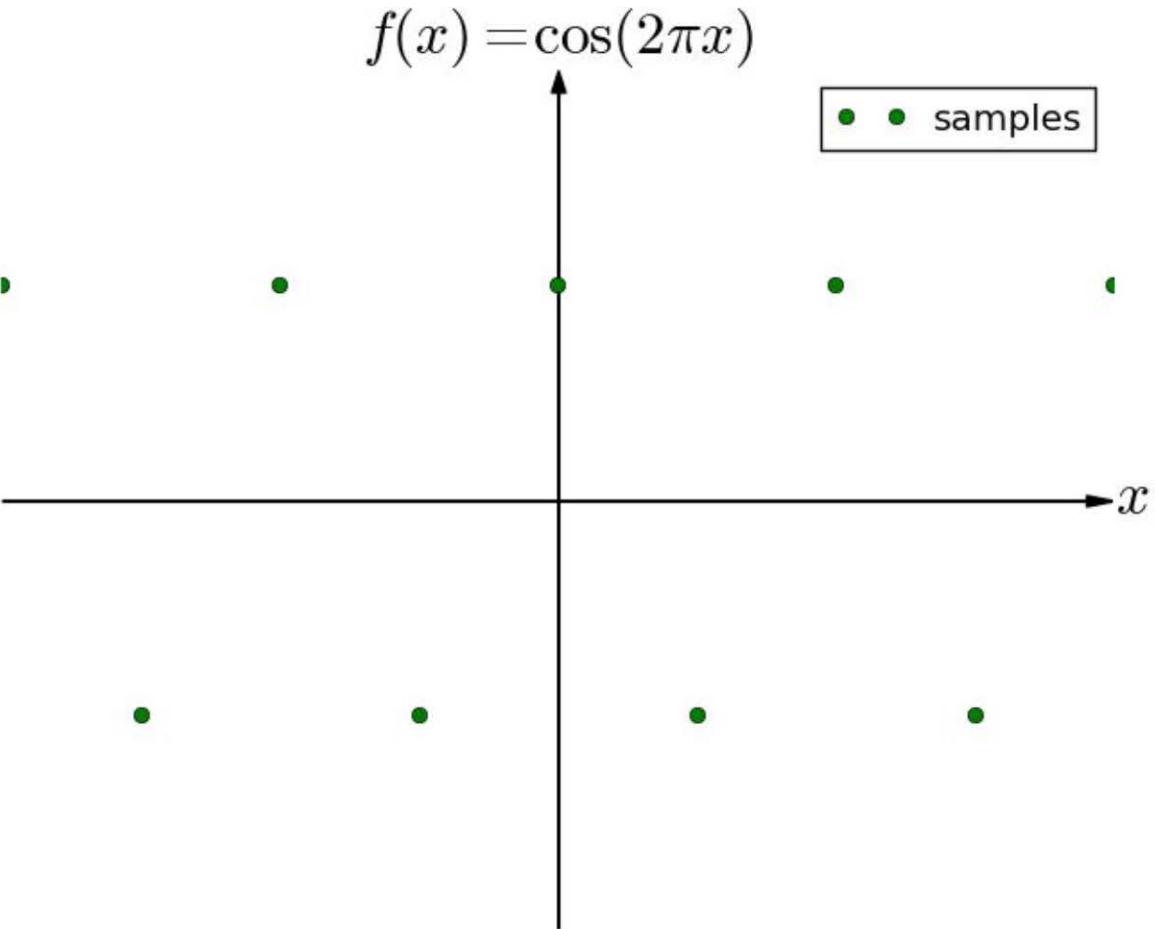
# reconstruction



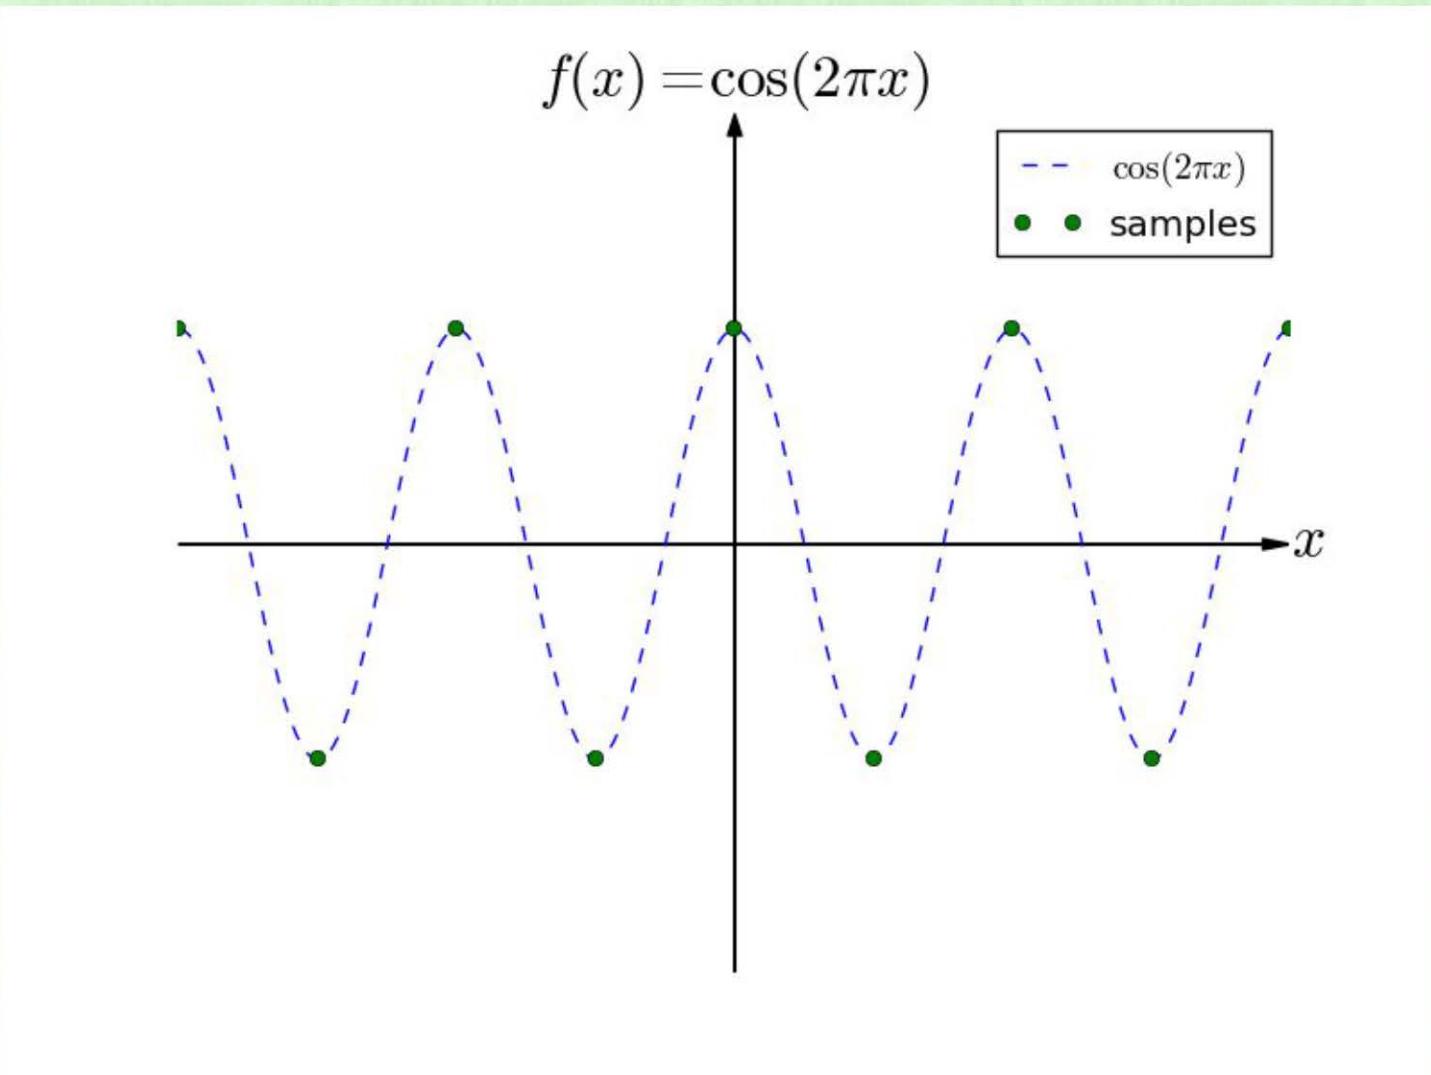
# 2 samples per period



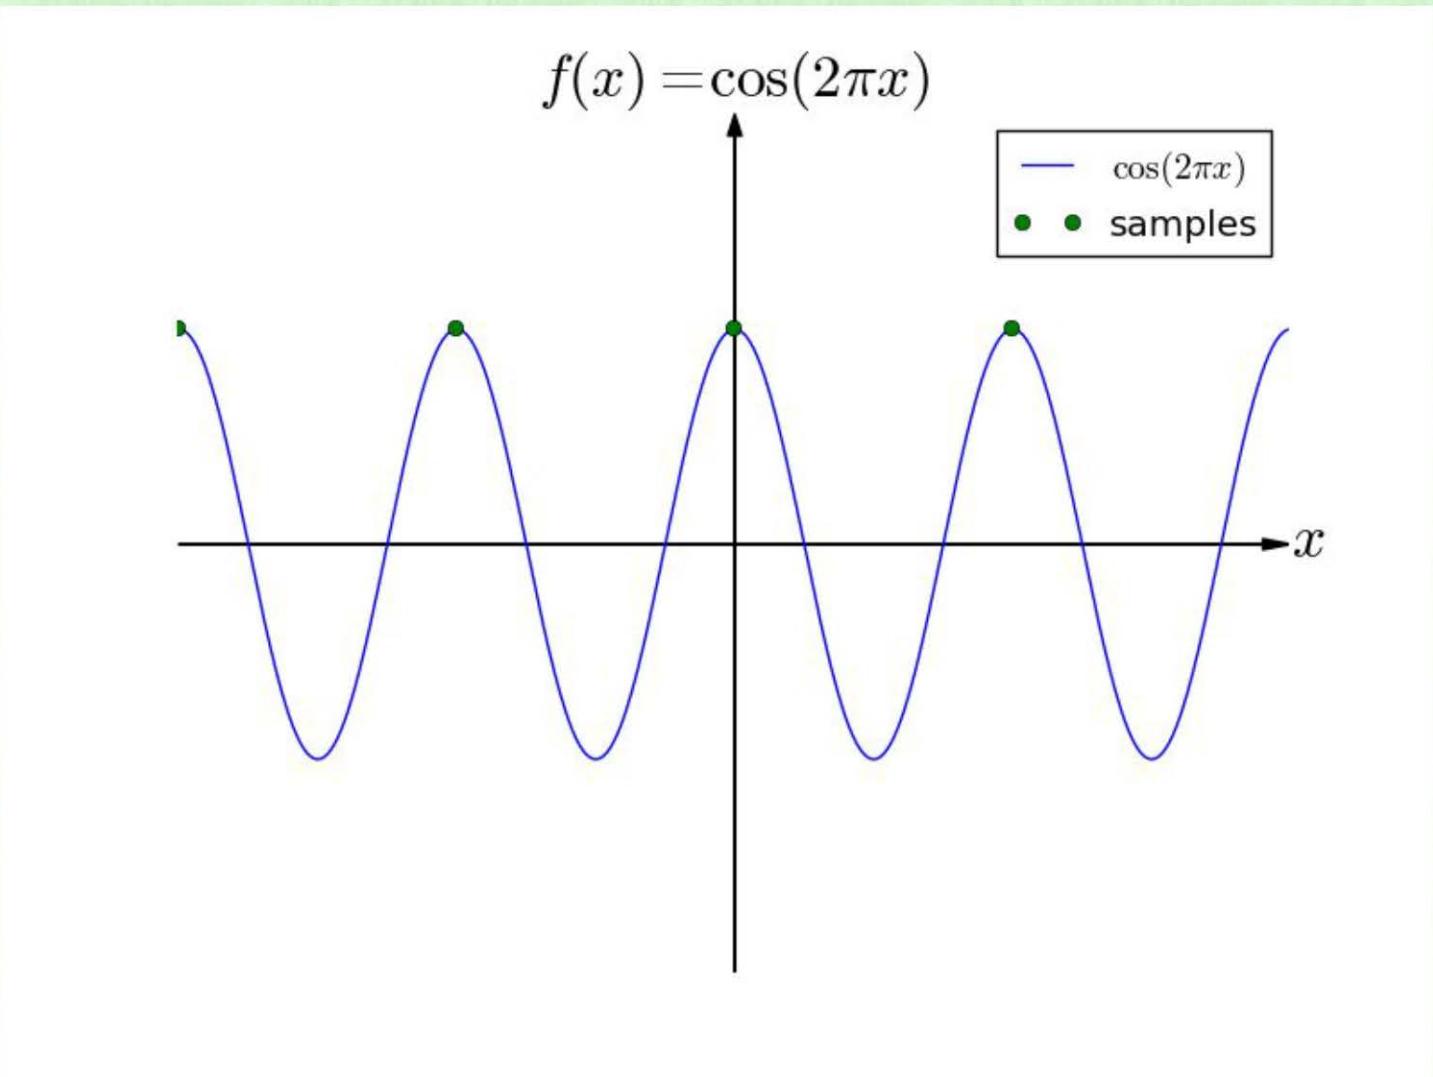
# samples



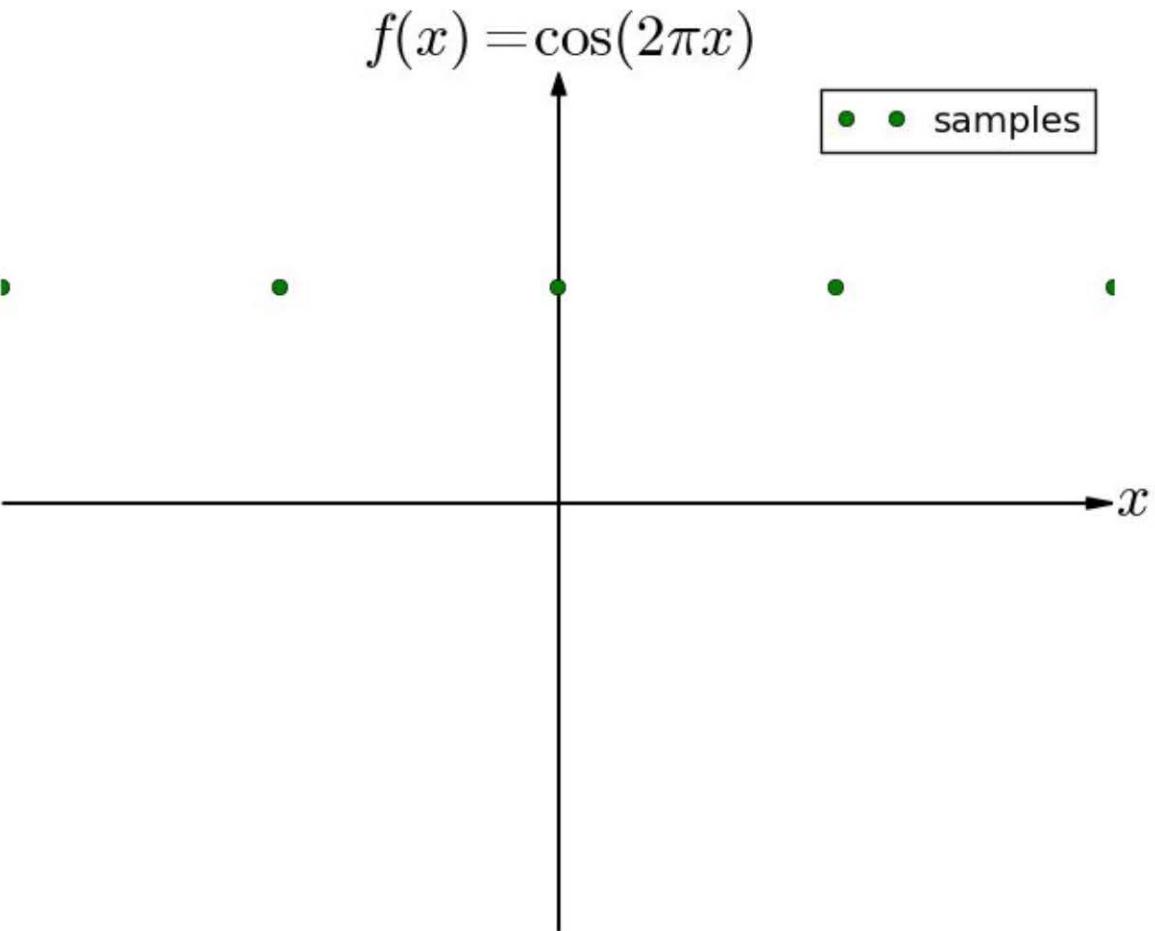
# reconstruction



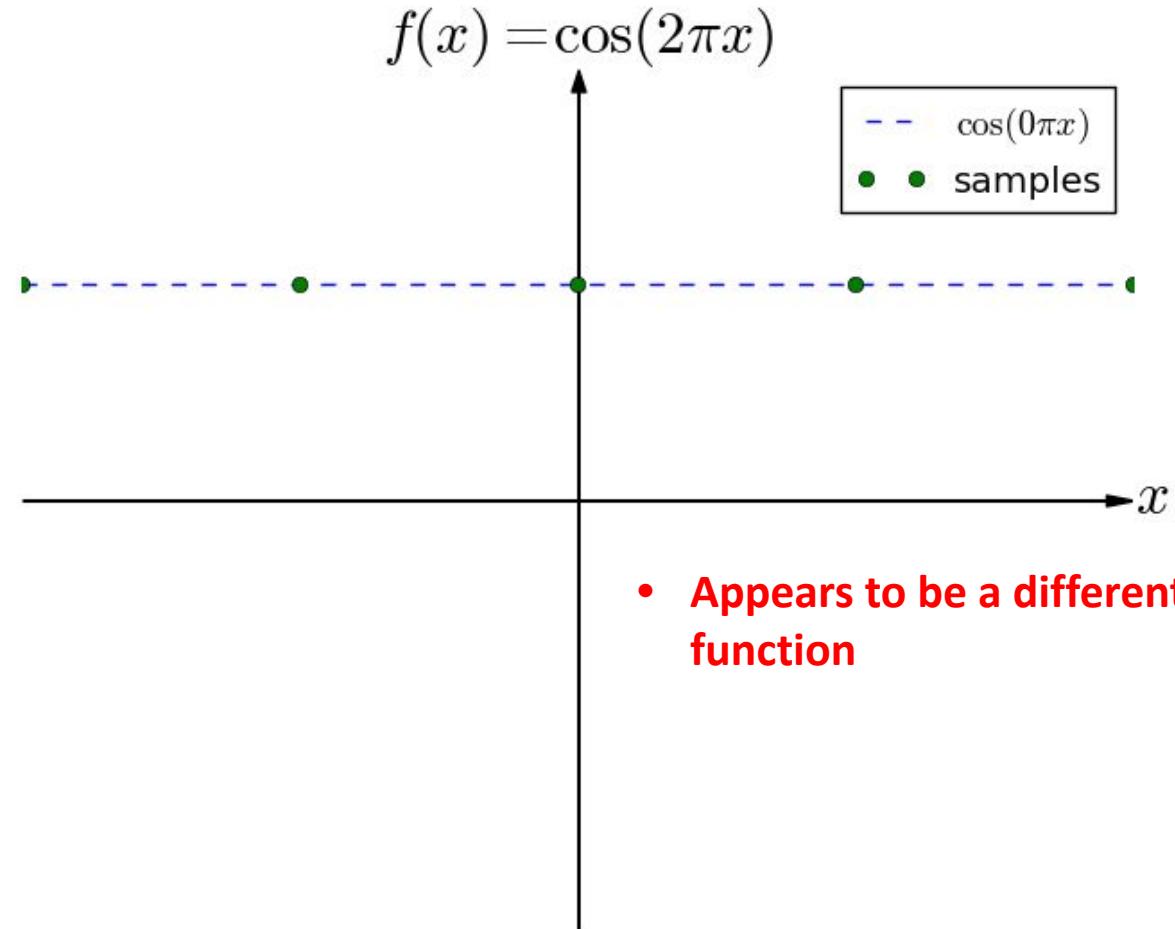
# 1 sample per period



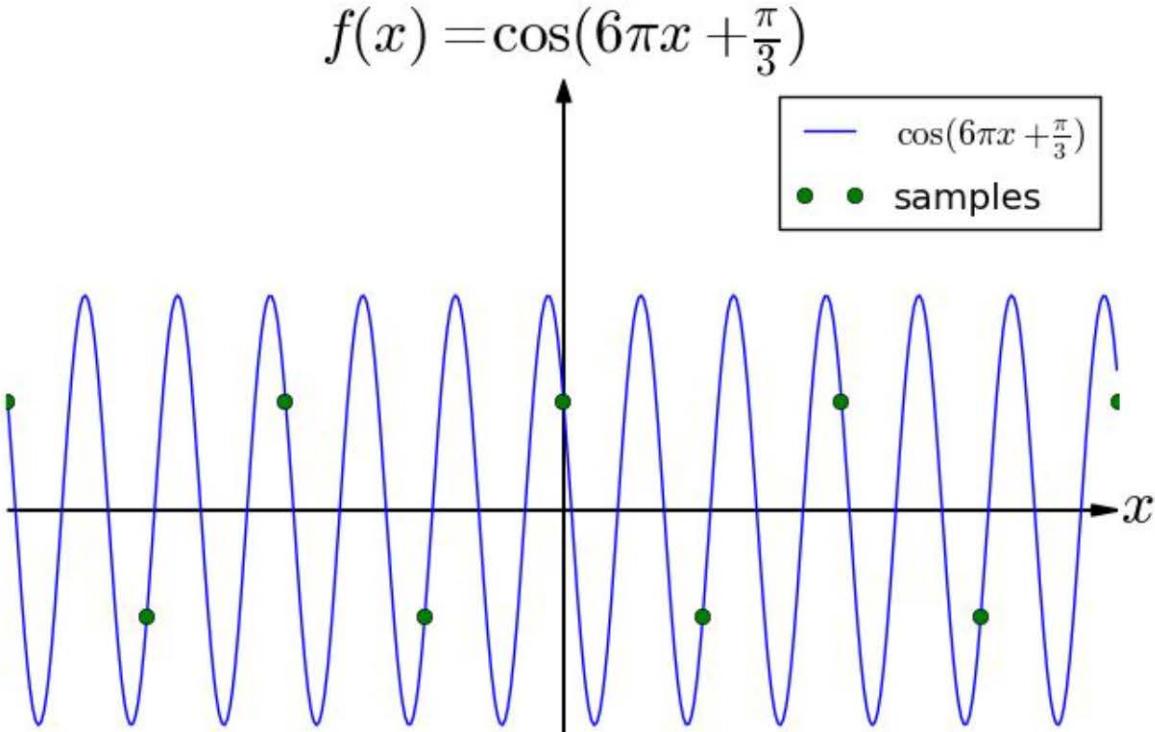
# samples



# reconstruction

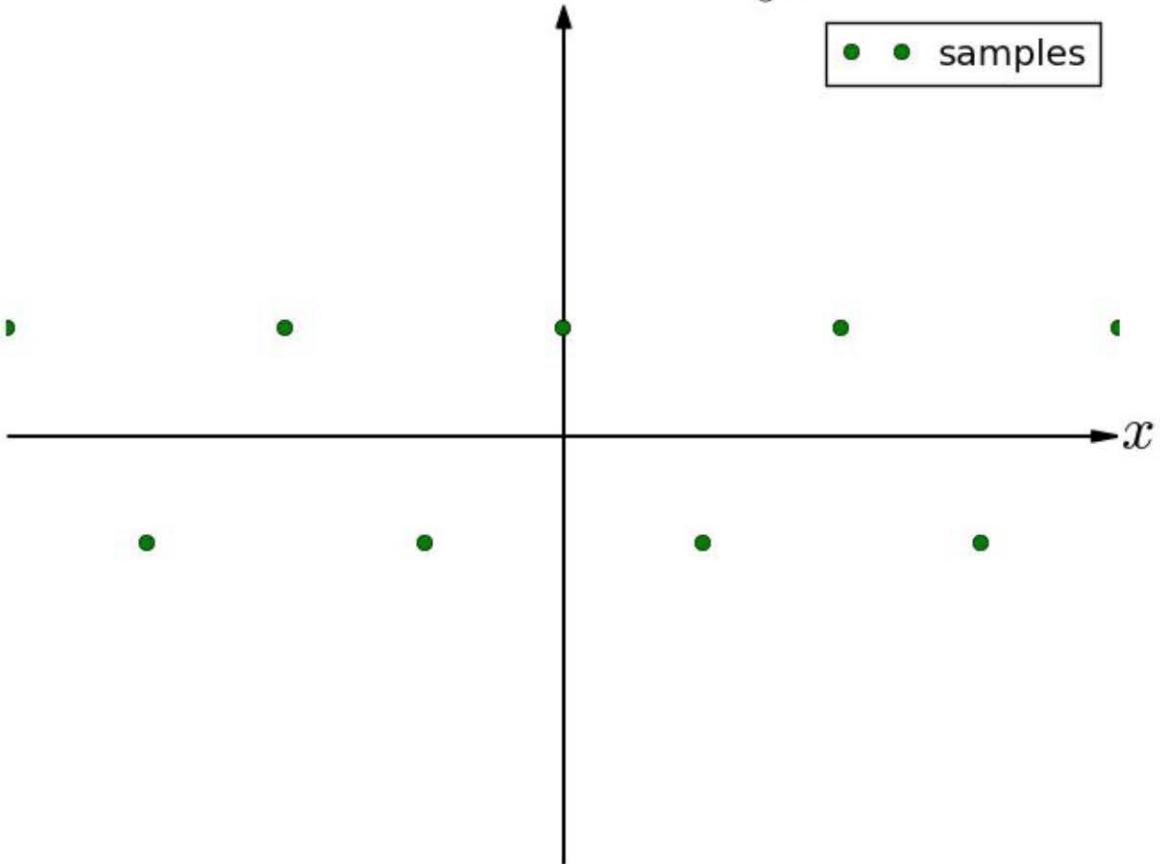


# 2/3 sample per period

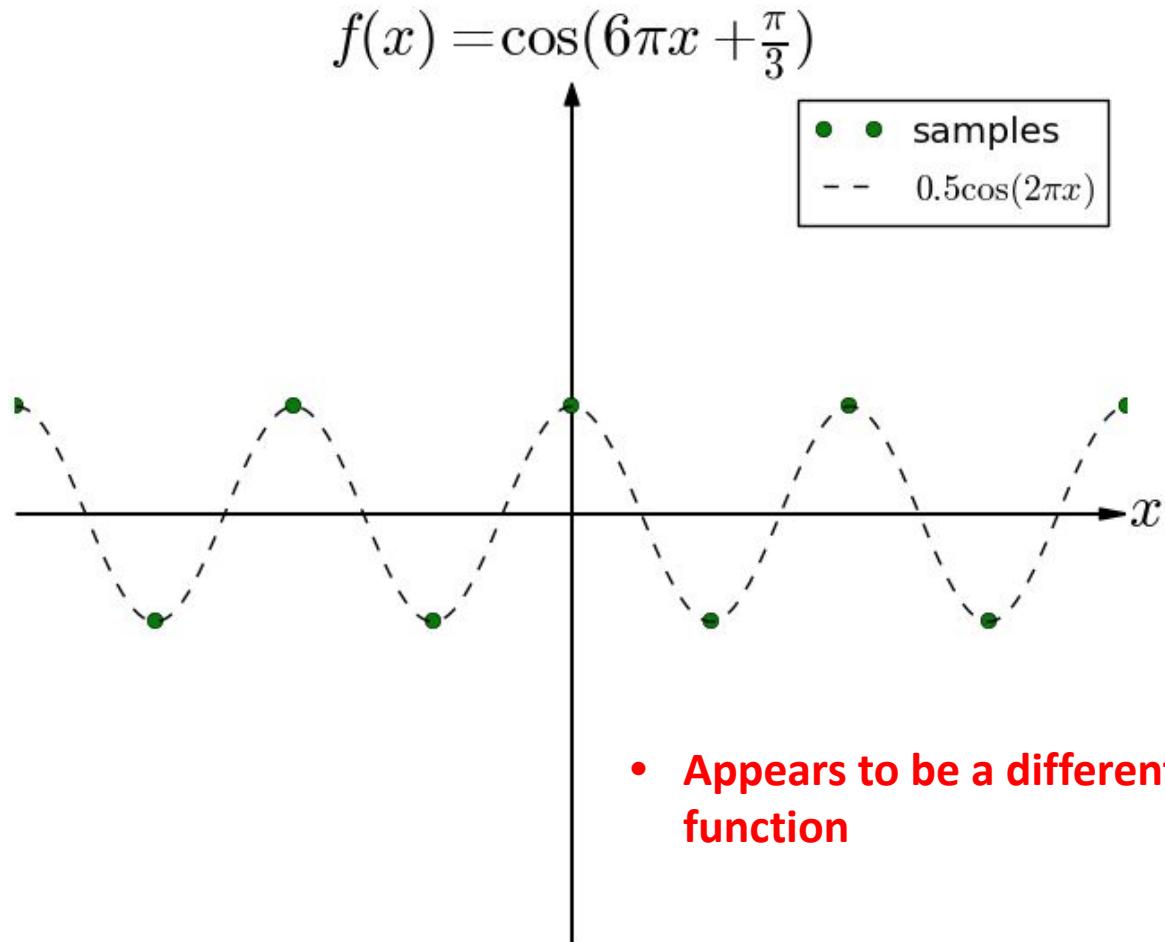


# samples

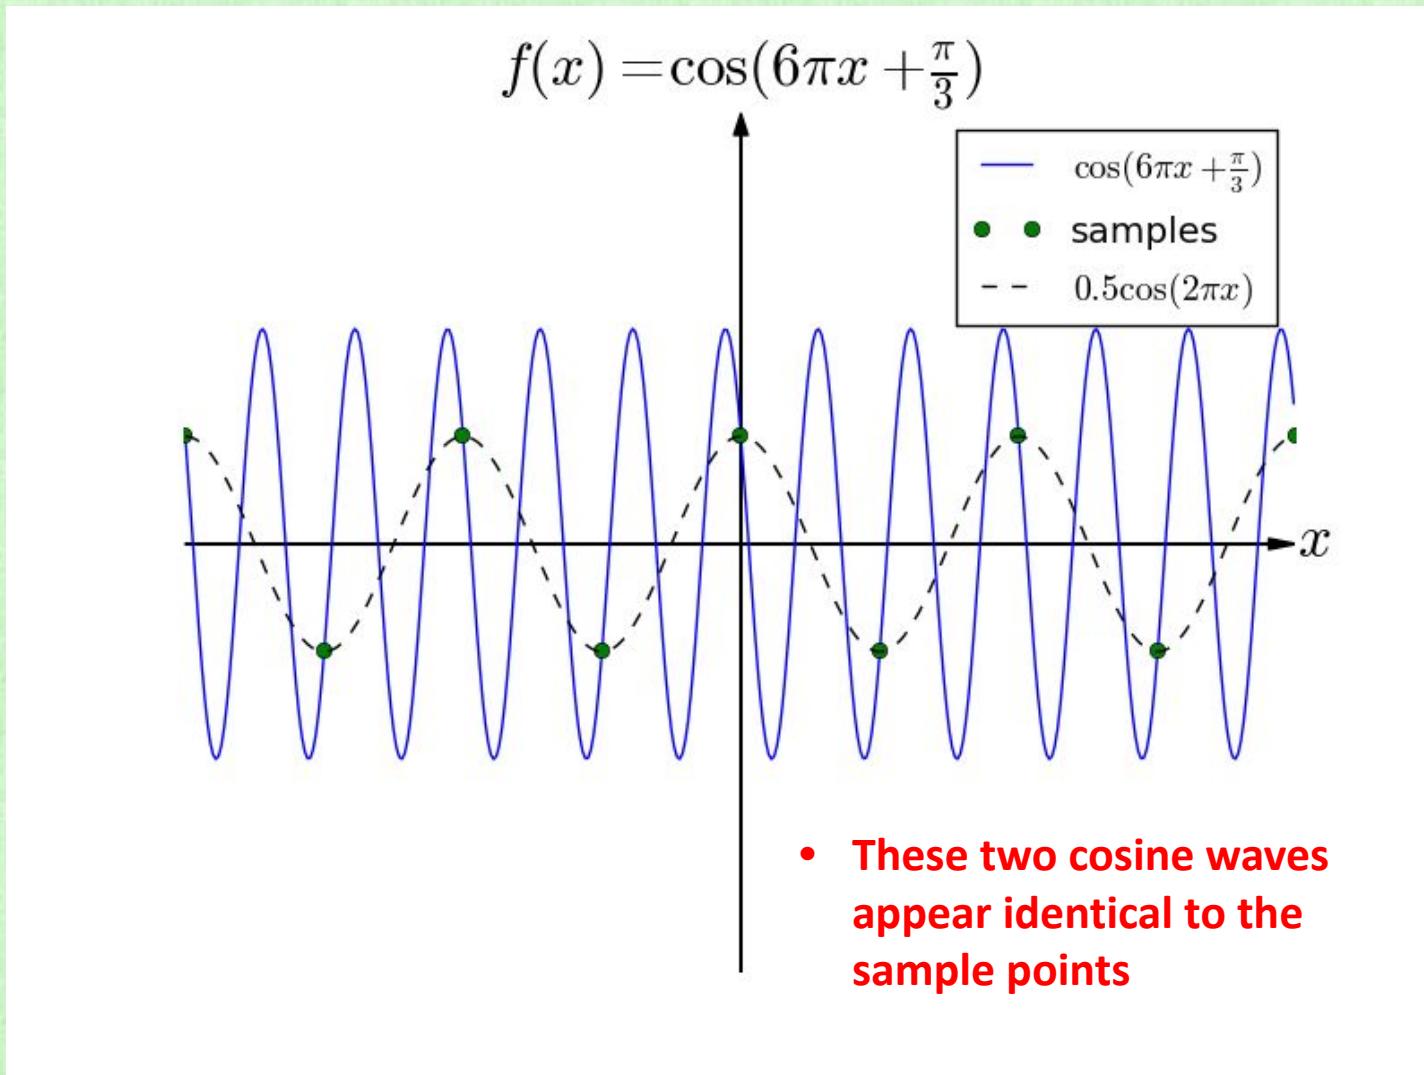
$$f(x) = \cos(6\pi x + \frac{\pi}{3})$$



# reconstruction



# Aliasing



# Sampling Rate

- Sampling at too low a rate results in aliasing
- Two different signals become indistinguishable (or aliased)
- Nyquist-Shannon Sampling Theorem
  - If  $f(t)$  contains no frequencies higher than  $W$  hertz, it can be completely determined by samples spaced  $1/(2W)$  seconds apart
  - That is, a minimum of 2 samples per period are required to prevent aliasing

# Anti-Aliasing

- The Nyquist frequency is defined as half the sampling frequency
- If the function being sampled has no frequencies above the Nyquist frequency, then no aliasing occurs
- *Real world frequencies above the Nyquist frequency appear as aliases to the sampler*
- **Before sampling, remove frequencies higher than the Nyquist frequency**

# Fourier Transform

- Transform between the spatial domain  $f(x)$  and the frequency domain  $F(k)$

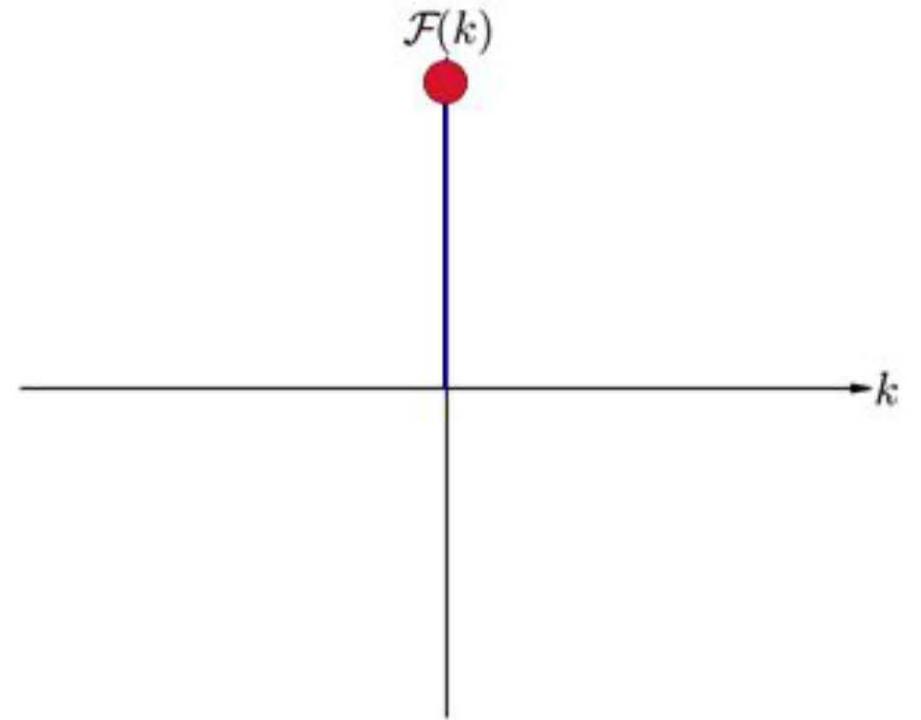
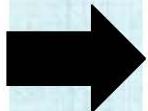
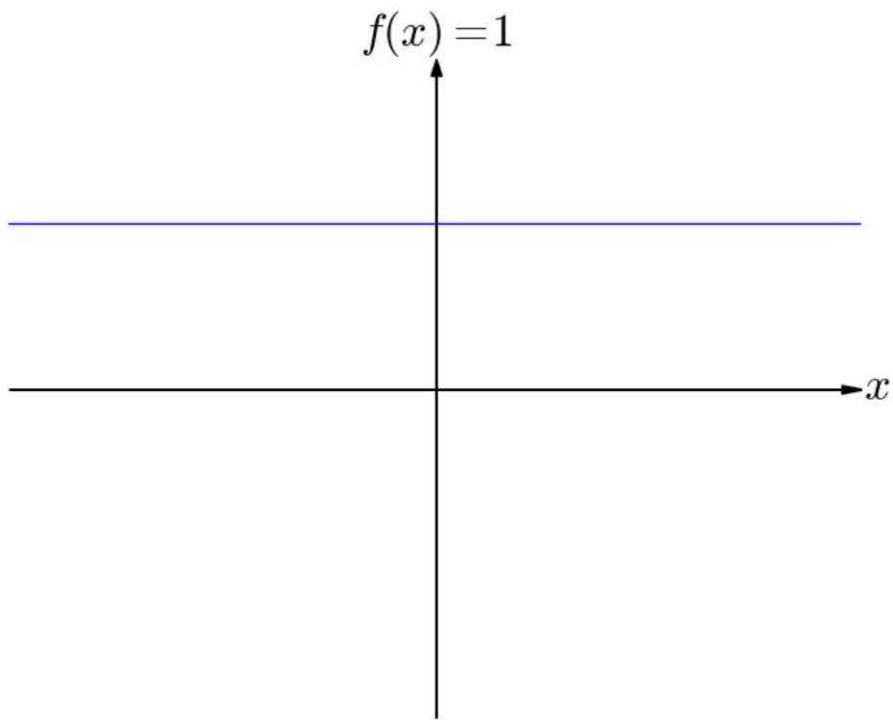
Frequency Domain: 
$$F(k) = \int_{-\infty}^{\infty} f(x)e^{-2\pi i k x} dx$$

Spatial Domain: 
$$f(x) = \int_{-\infty}^{\infty} F(k)e^{2\pi i k x} dk$$

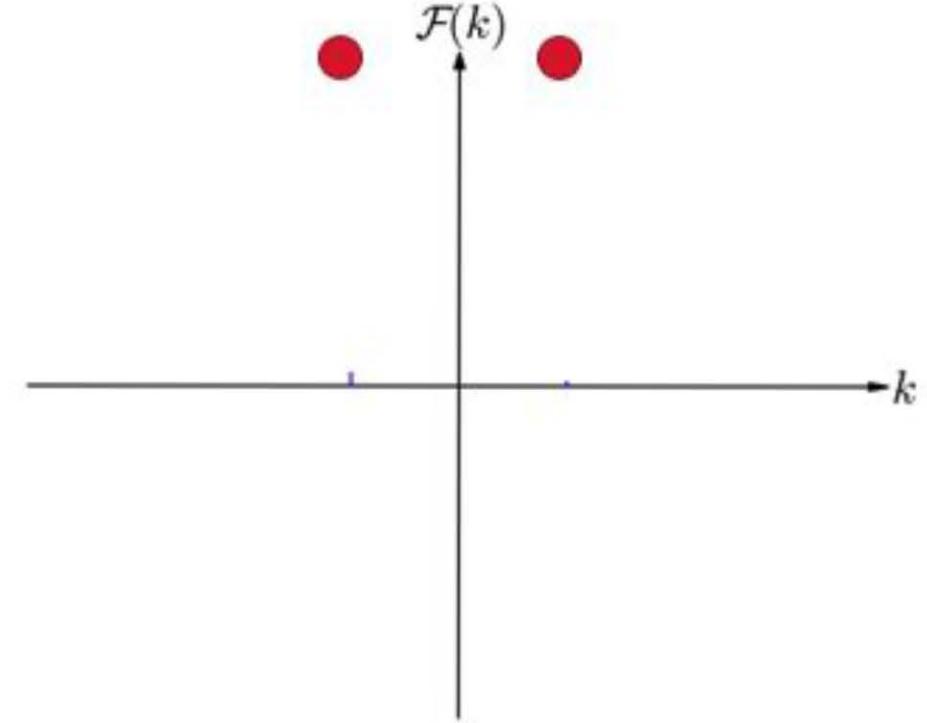
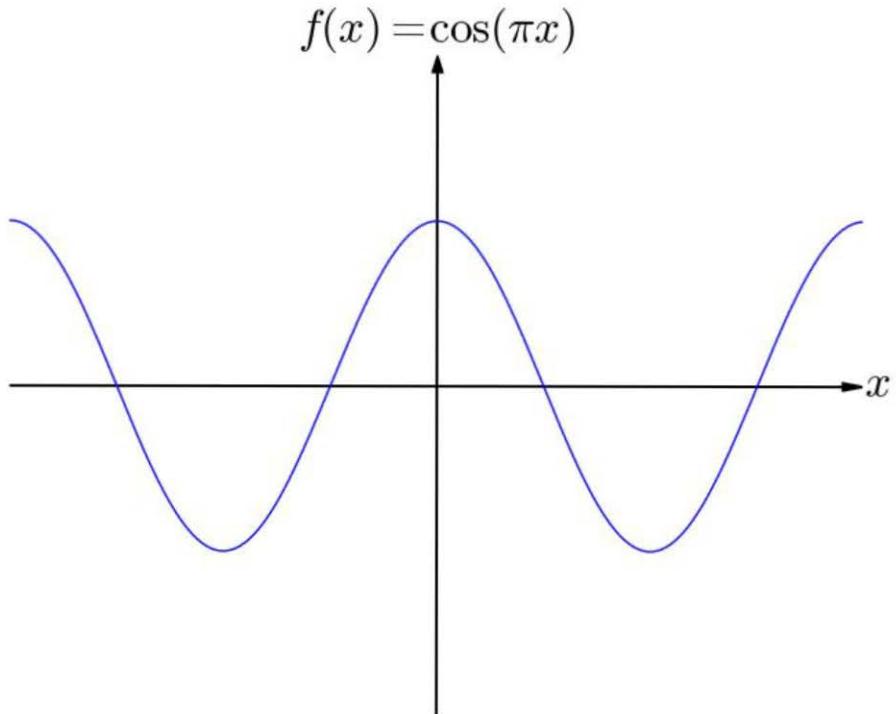
$$e^{i\theta} = \cos \theta + i \sin \theta$$

$$\cos \theta = \frac{e^{i\theta} + e^{-i\theta}}{2} \qquad \sin \theta = \frac{e^{i\theta} - e^{-i\theta}}{2i}$$

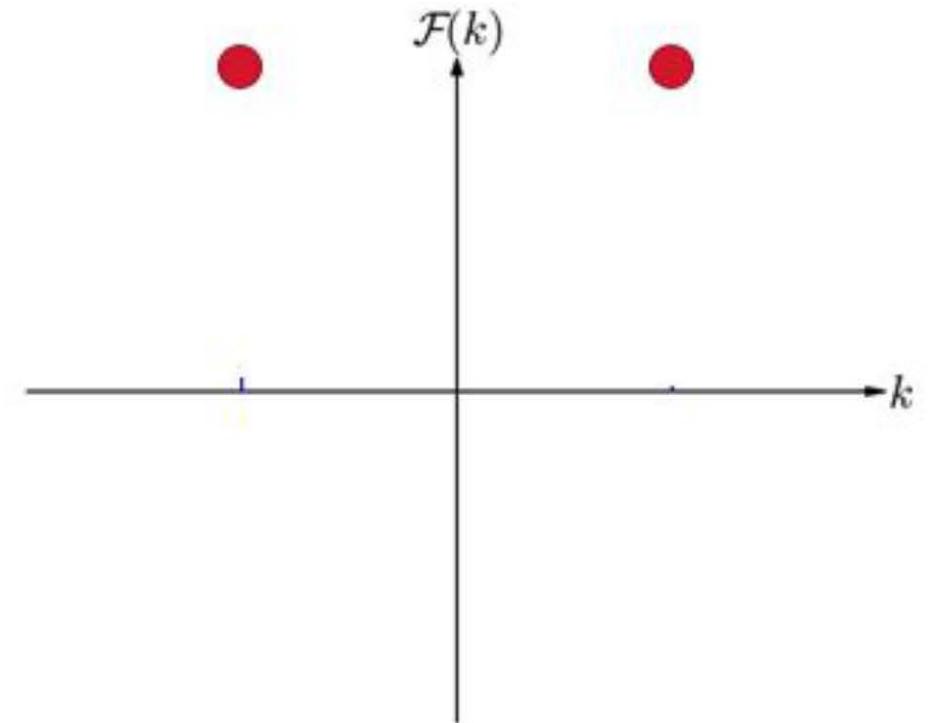
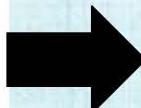
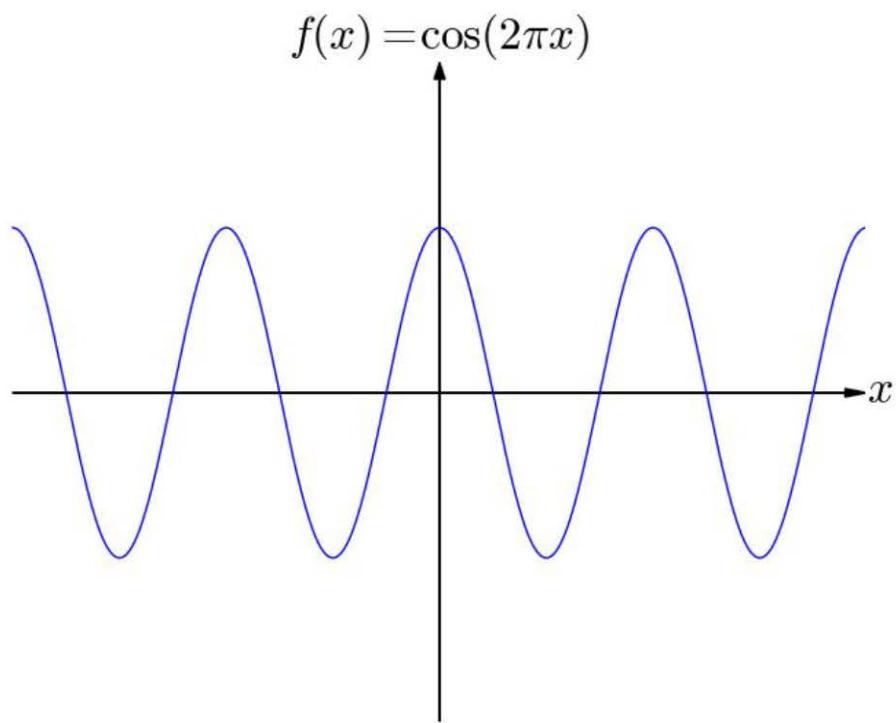
# Constant Function



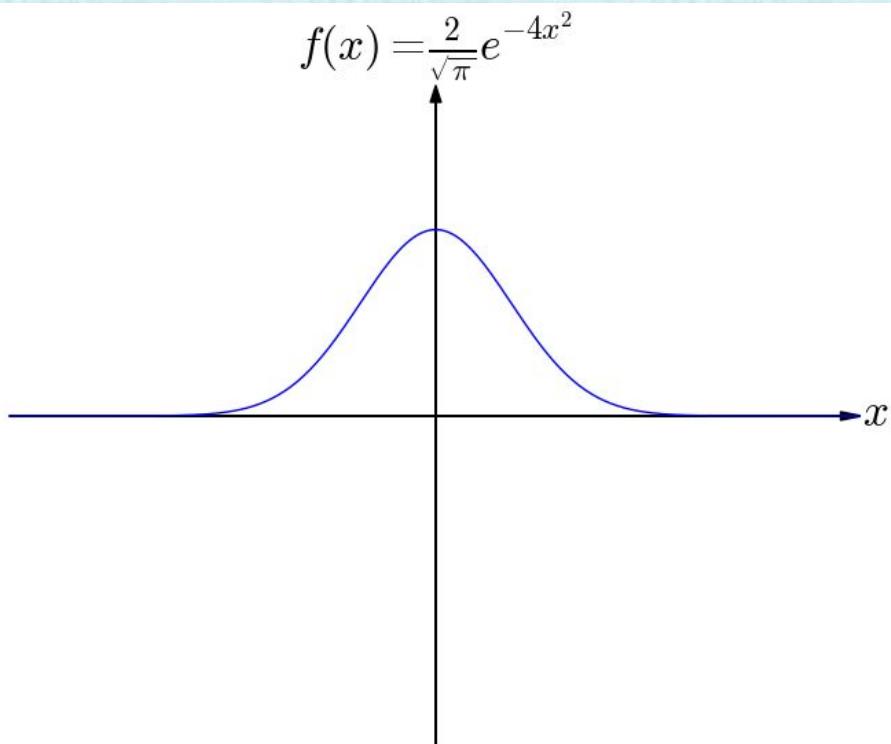
# Low Frequency Cosine



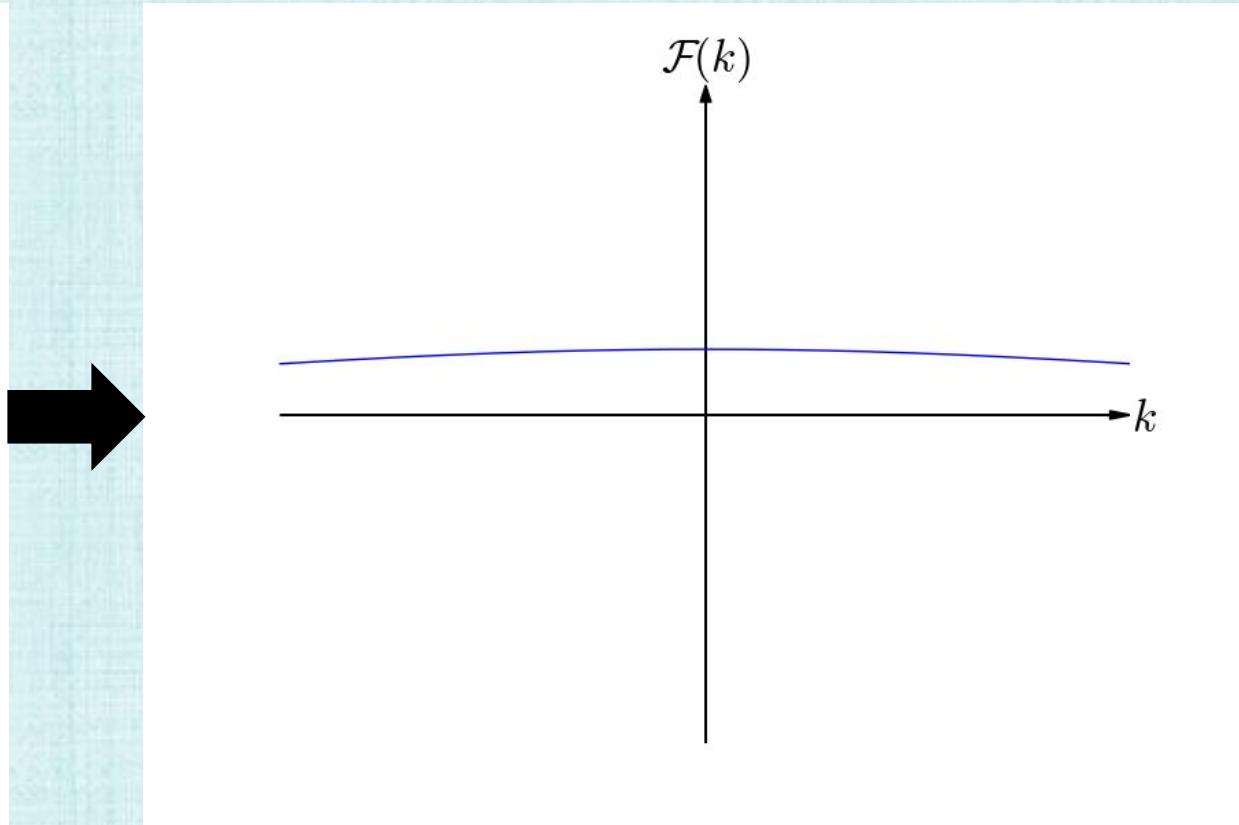
# High Frequency Cosine



# Narrow Gaussian

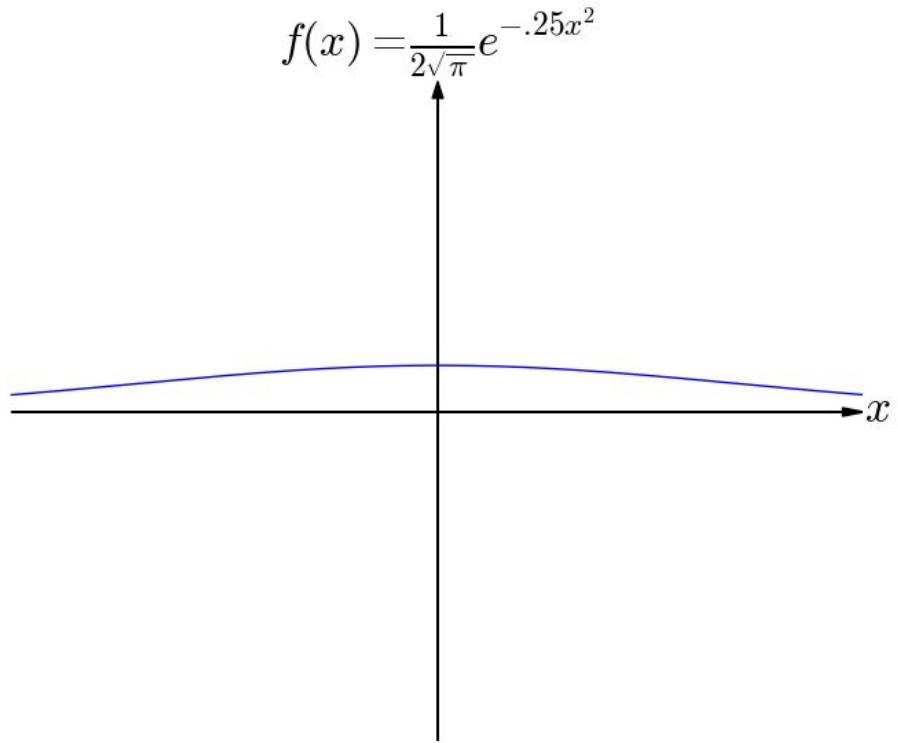


Narrow

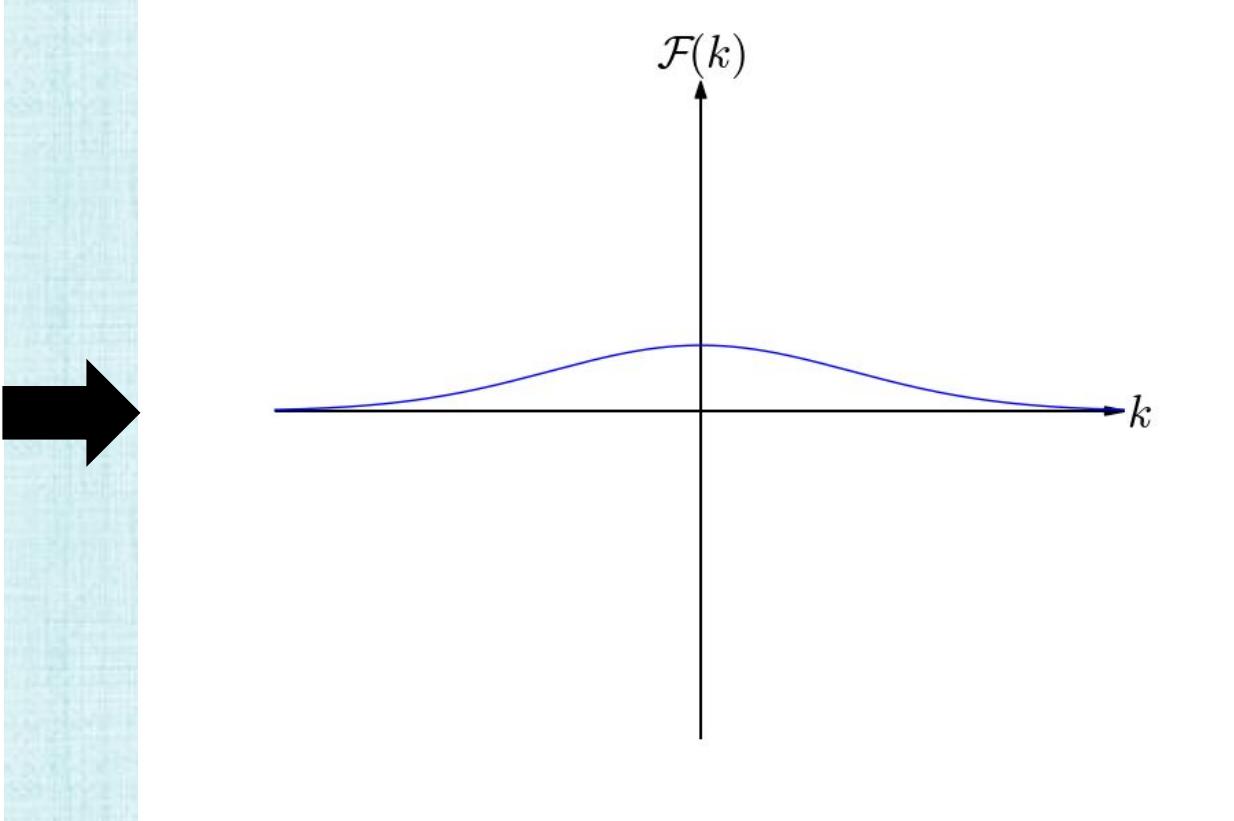


Wide

# Wider Gaussian

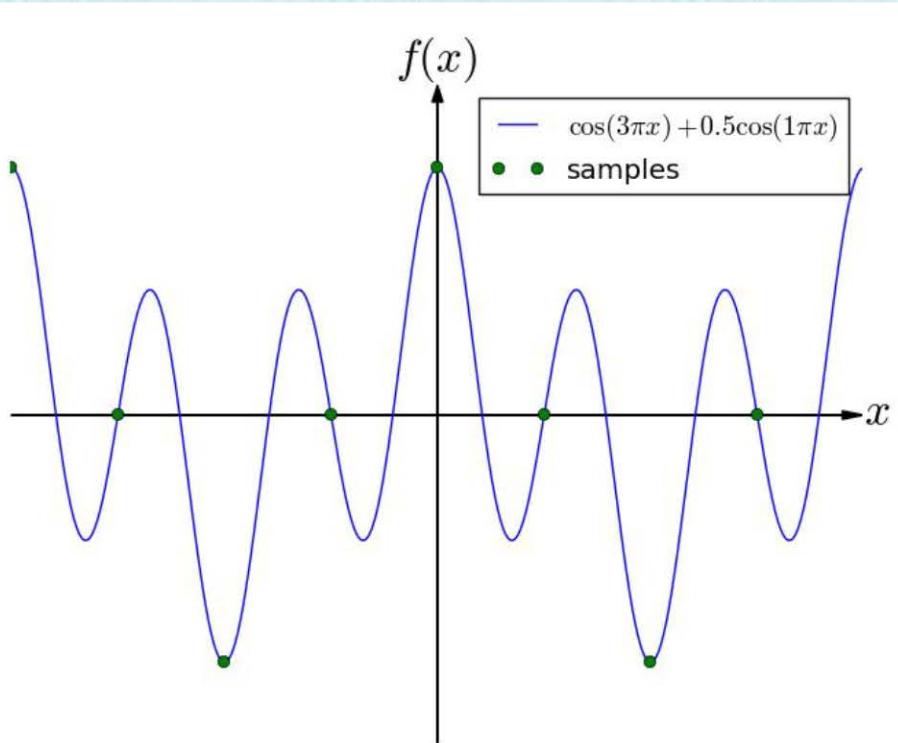


Wider

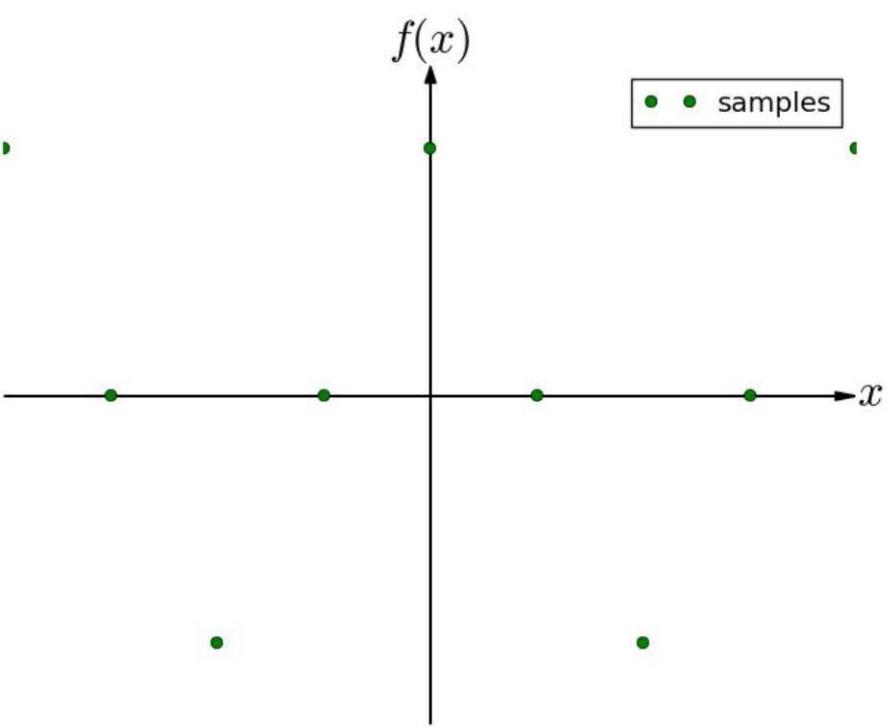


Narrower

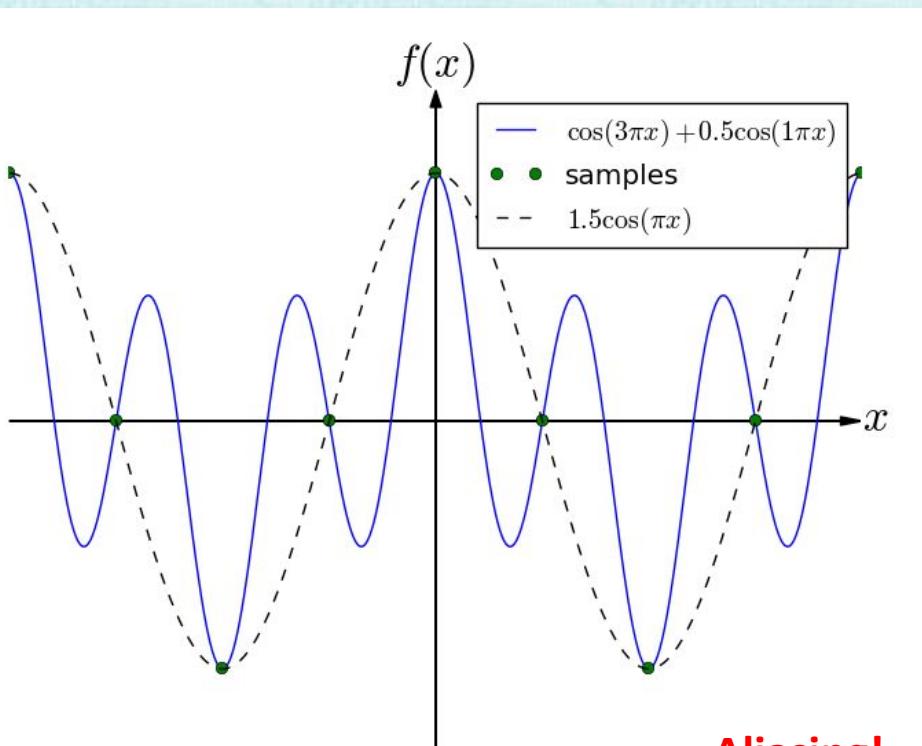
# sum of two different cosine functions



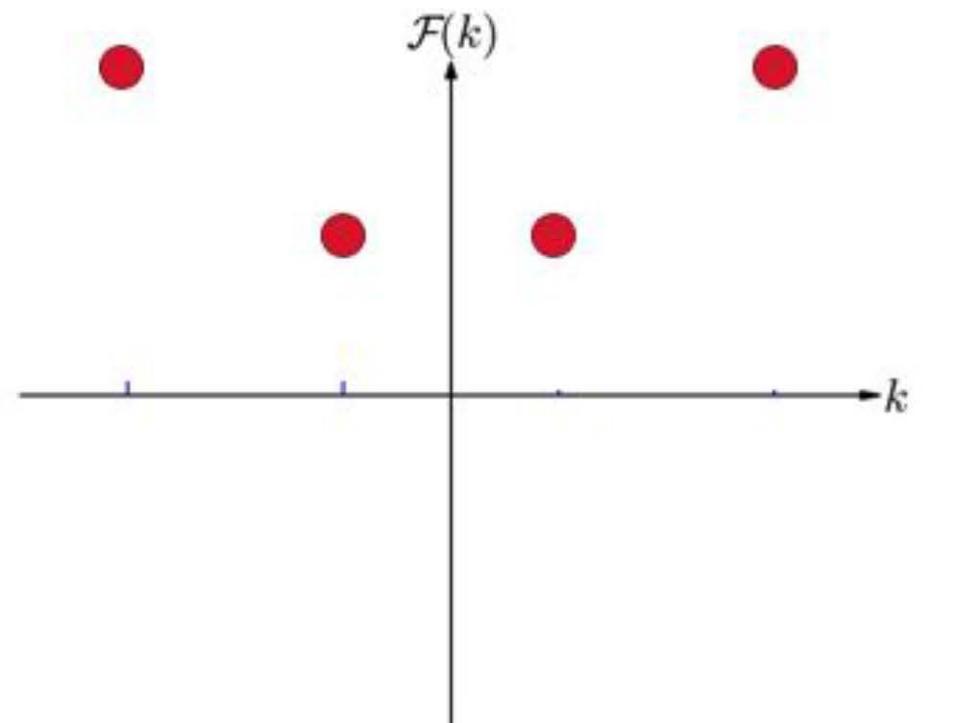
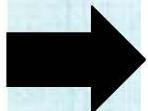
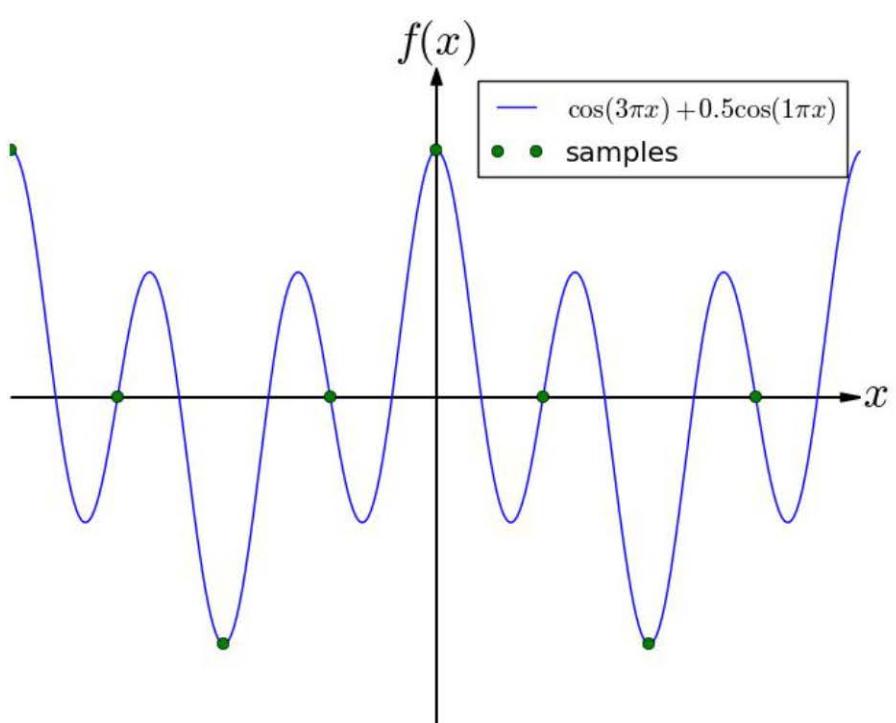
# samples



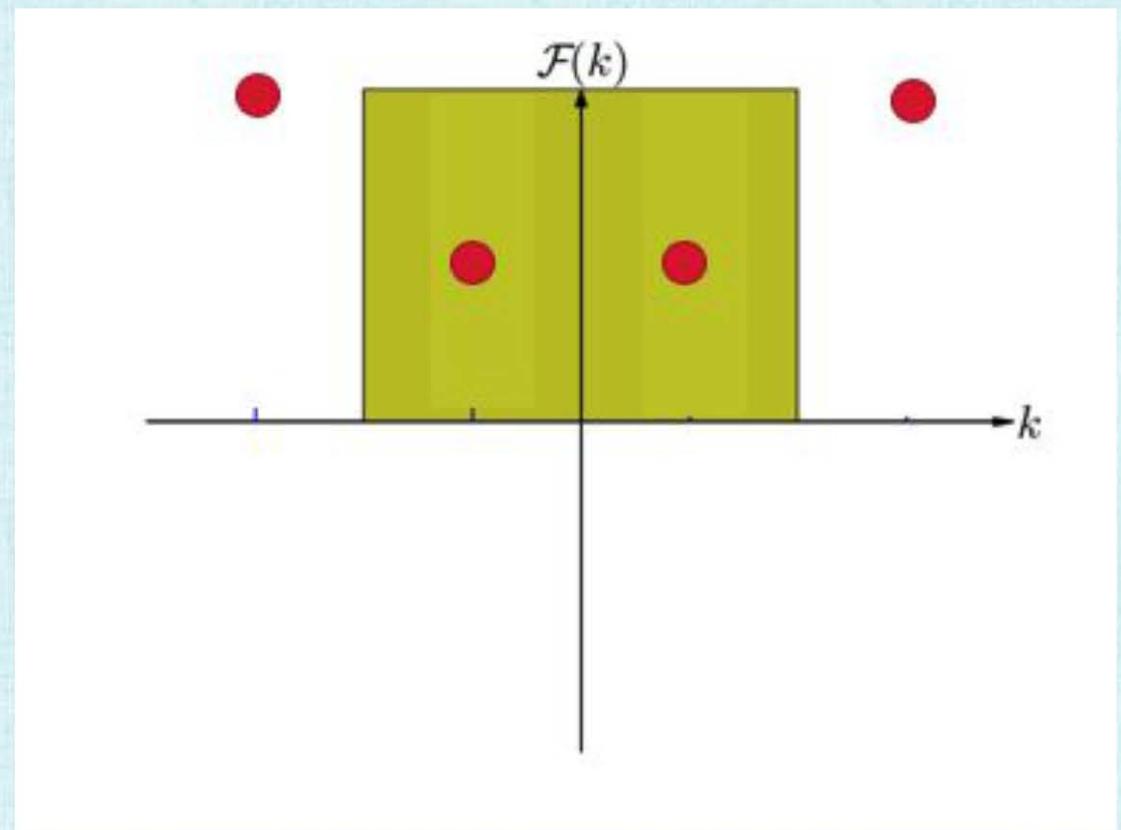
# reconstruction



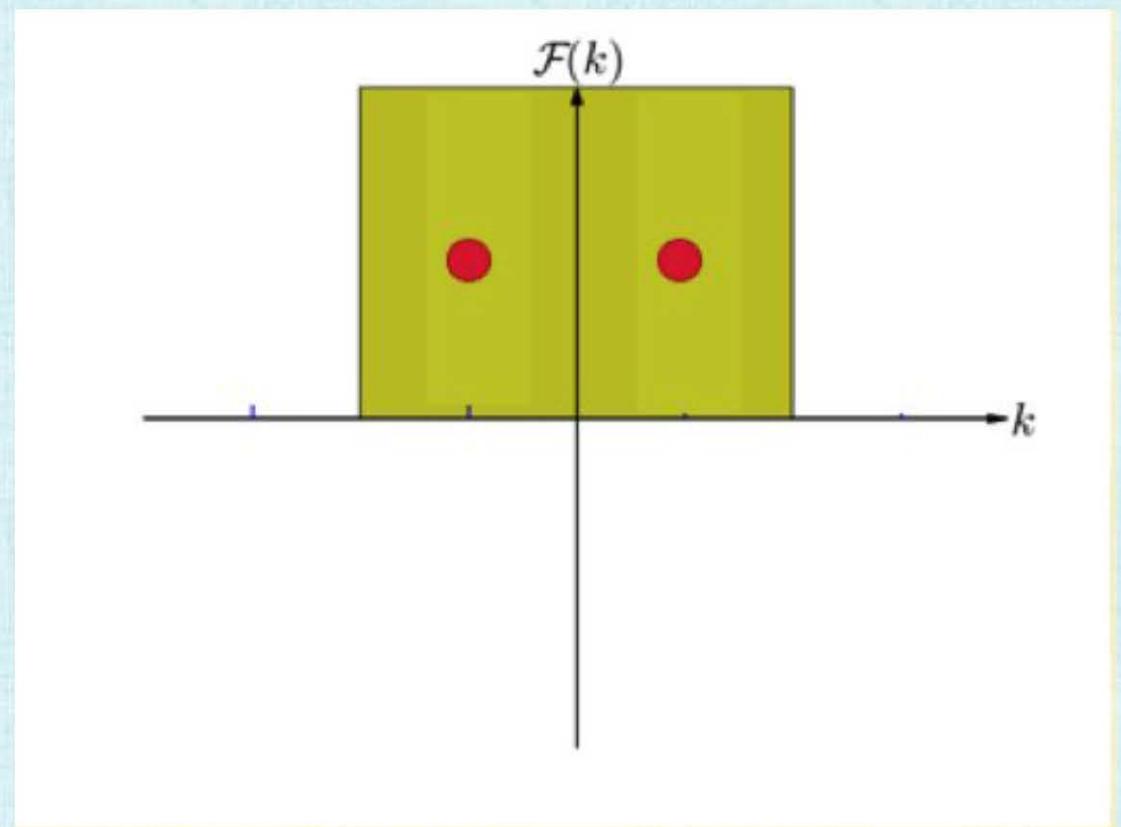
# Fourier transform



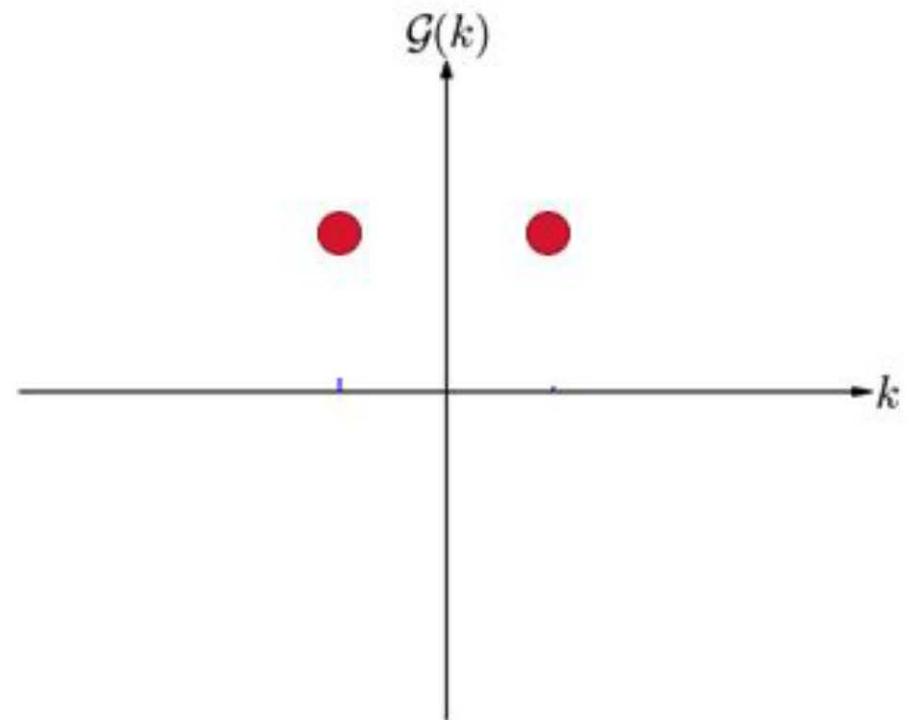
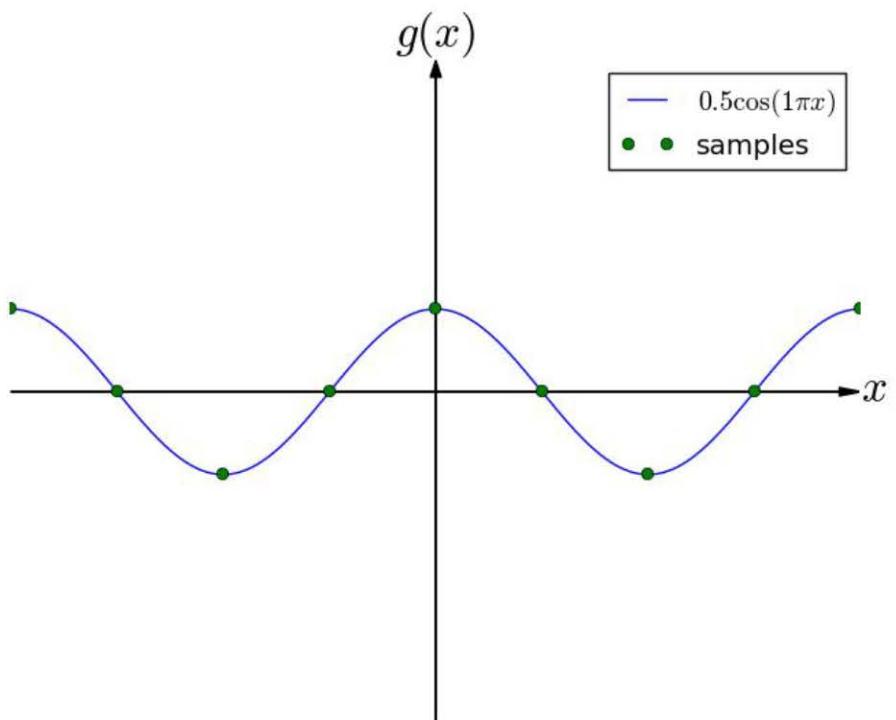
# identify Nyquist frequency bounds



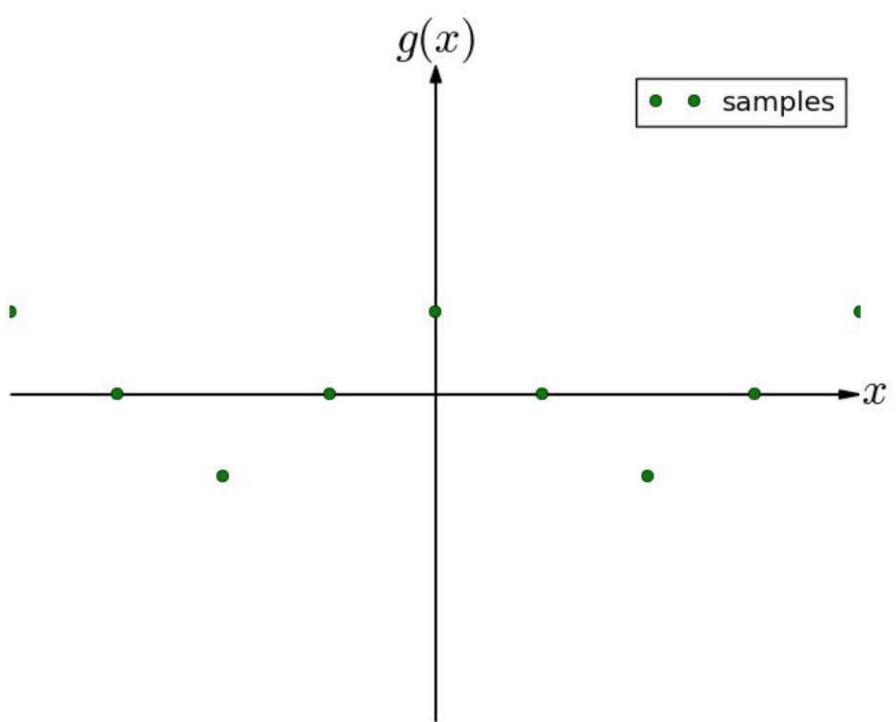
# remove high frequencies



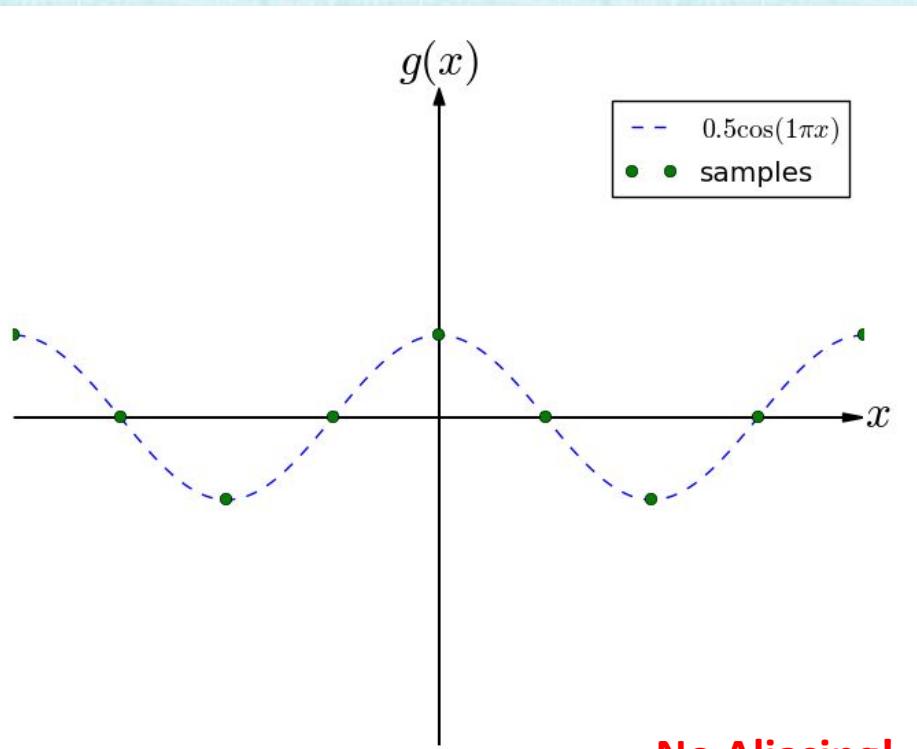
# inverse Fourier transform



# samples



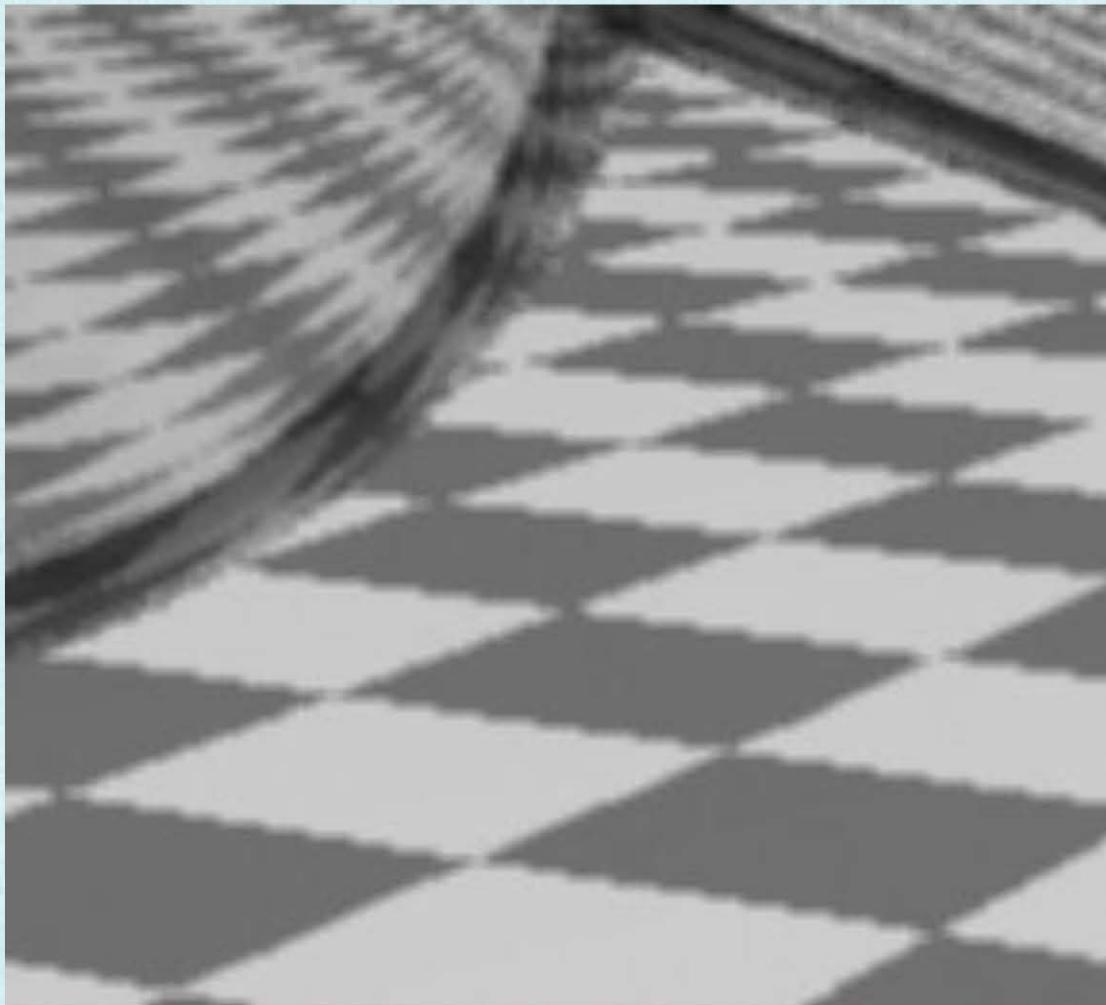
# reconstruction



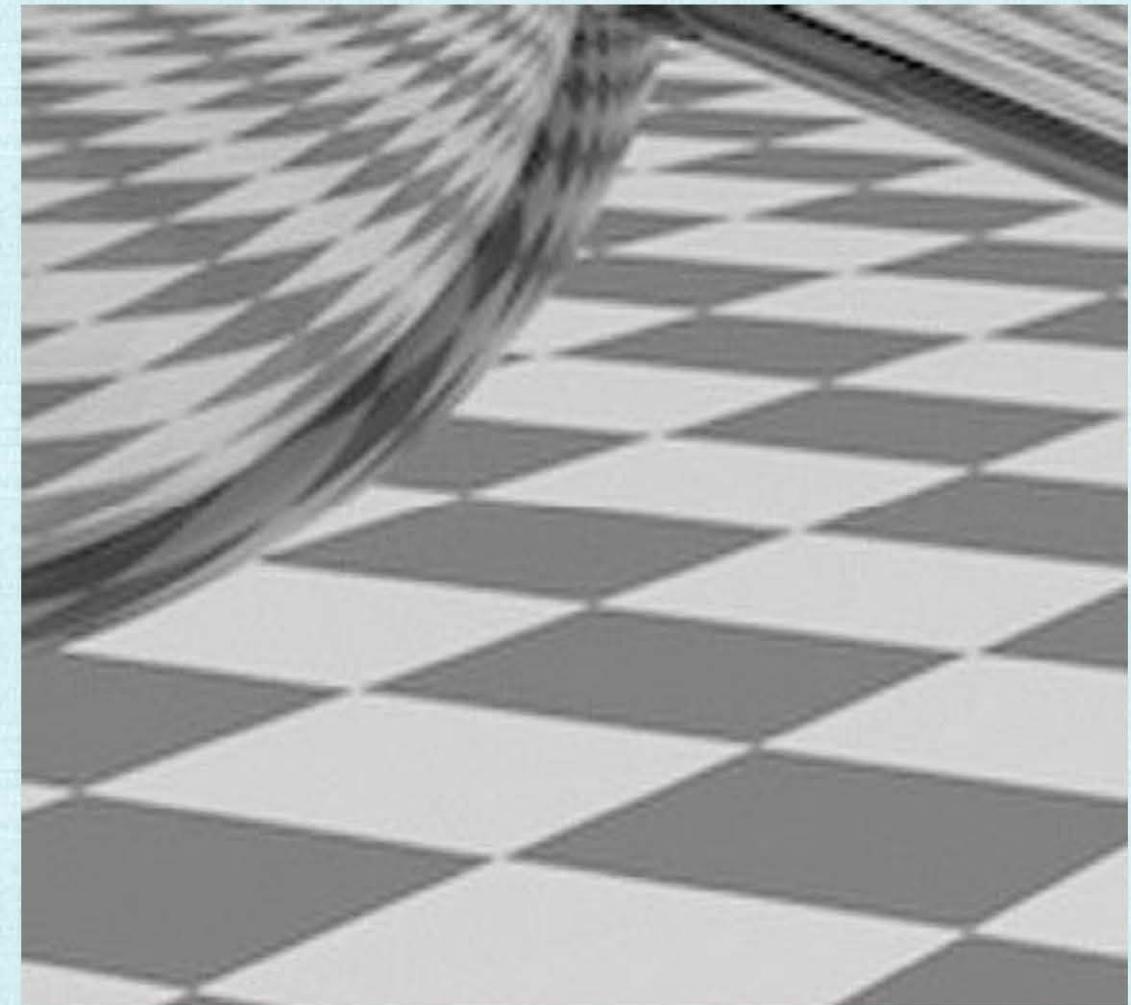
# Anti-Aliasing

- Sampling causes higher frequencies to masquerade as lower frequencies
- After sampling, cannot untangle the mixed high/low frequencies
- Remove the high frequencies **before** sampling (in order to avoid aliasing)
- Part of the signal is lost
- But, that part of the signal was not representable by the sampling rate anyways

# Blurring vs. Anti-Aliasing



blurring jaggies after sampling

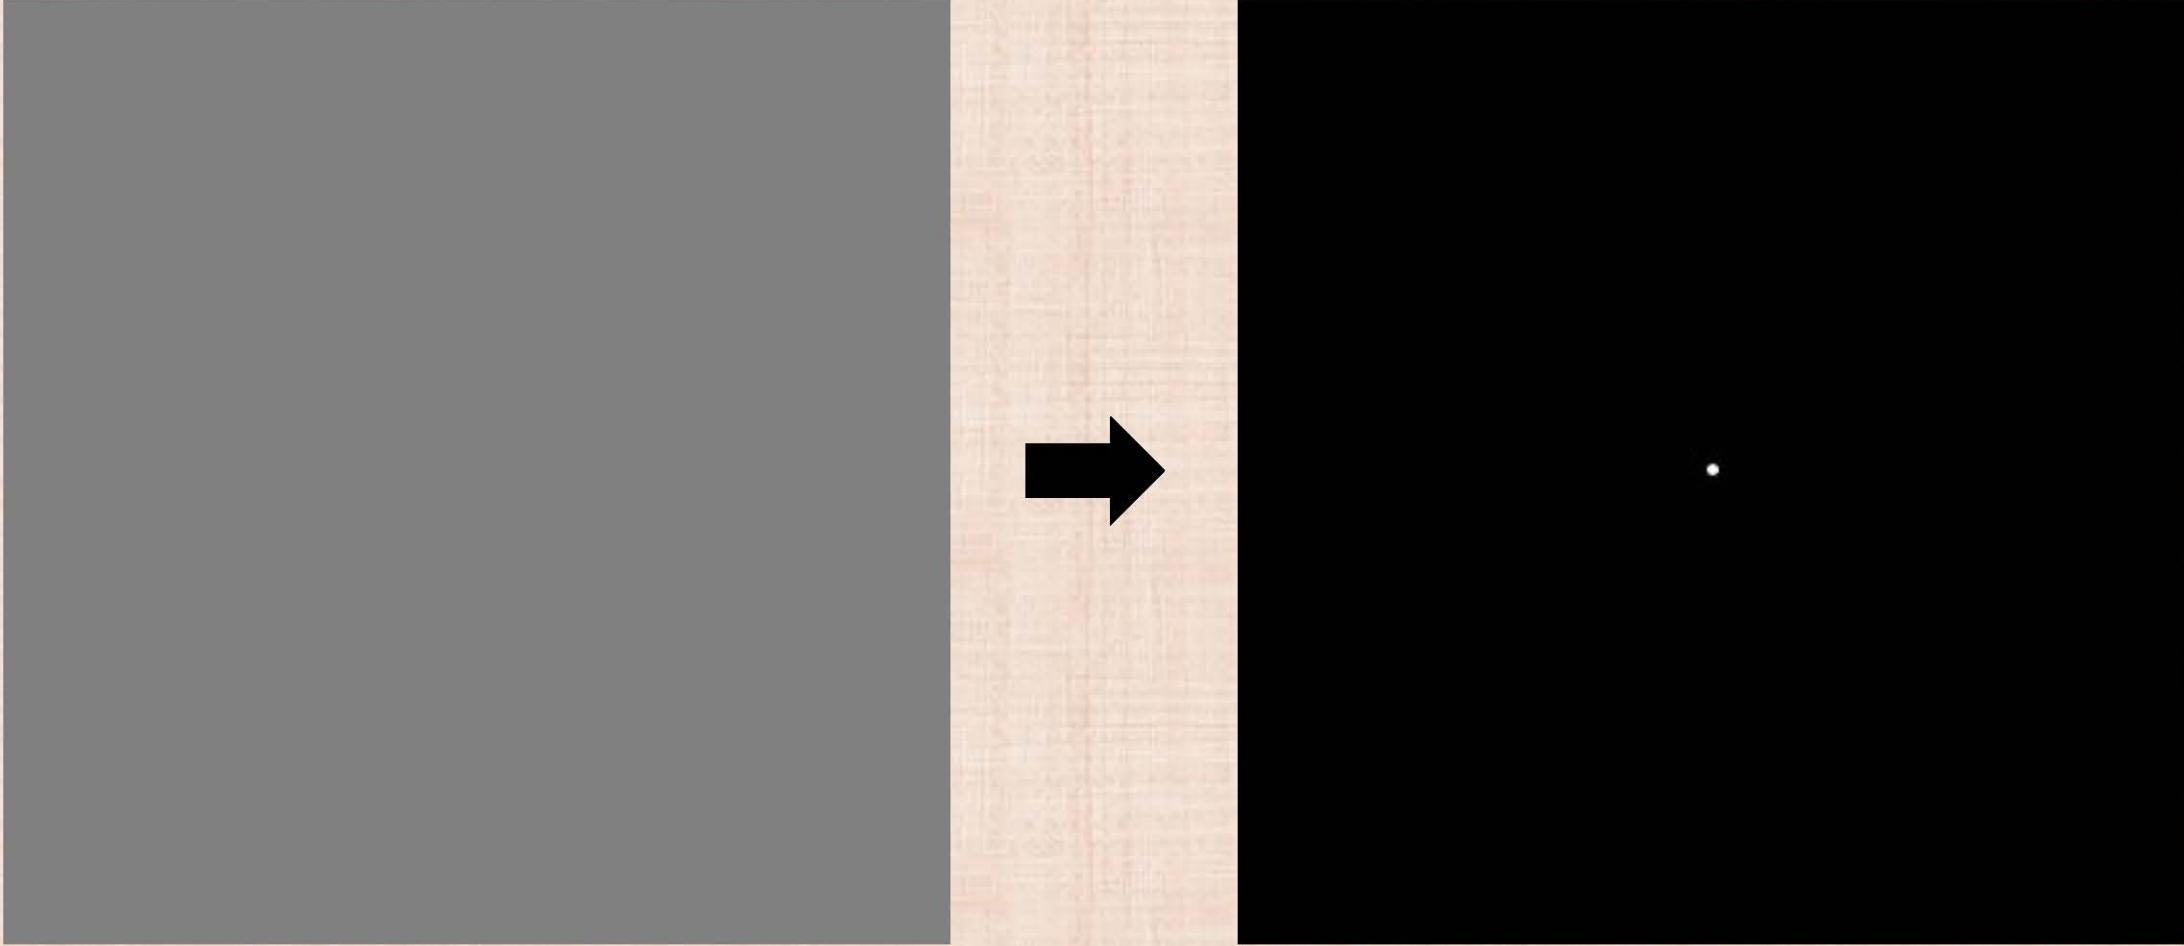


removing high frequencies before sampling

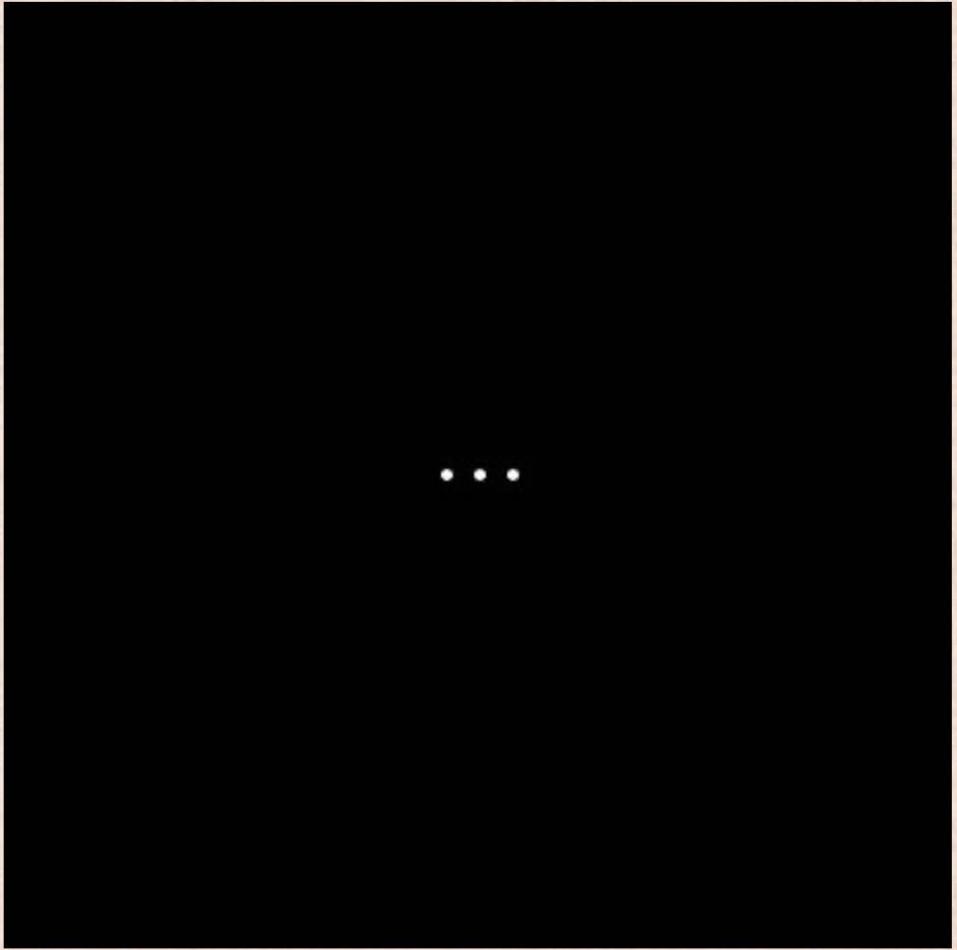
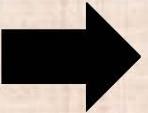
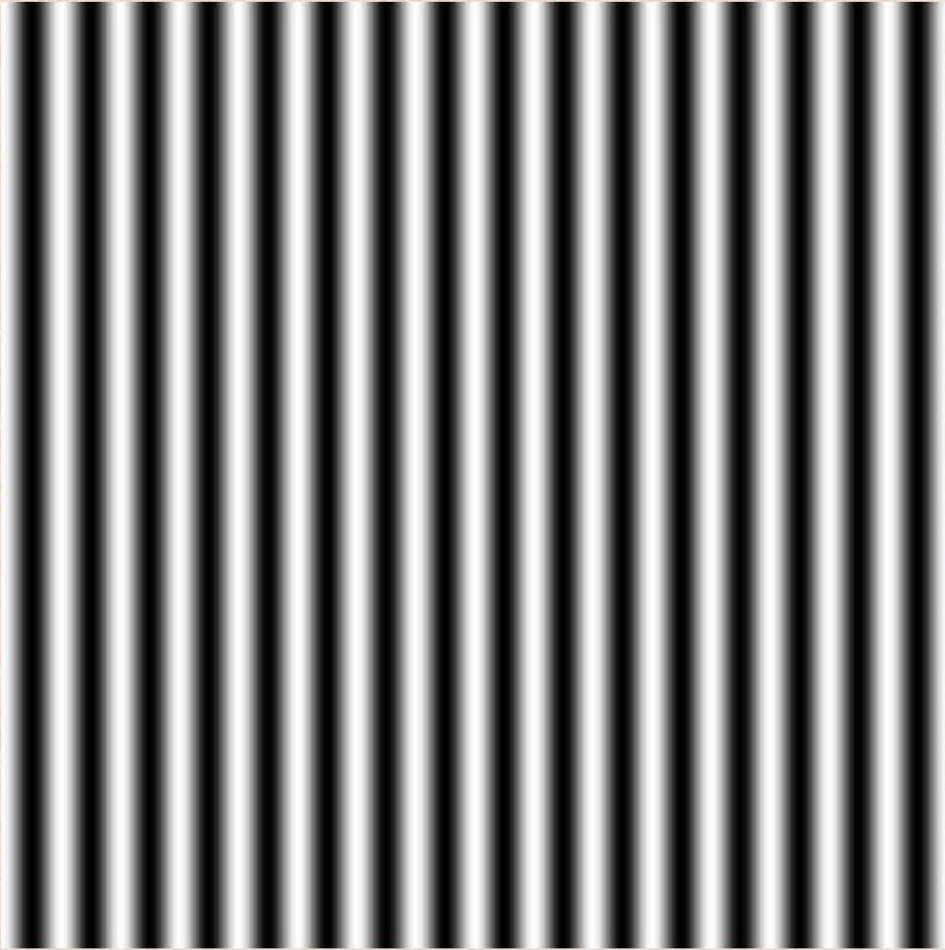
# Images

- Images have discrete values (and are not continuous functions)
    - Use a discrete version of the Fourier transform
    - The Fast Fourier Transform (FFT) computes the discrete Fourier transform (and its inverse) in  $O(n \log n)$  complexity (where  $n$  is the number of samples)
  - Images are 2D (not 1D)
    - A 2D discrete Fourier transform can be computed using 1D transforms along each dimension
1. Fourier transform (into the frequency domain)
    - Discrete image values are transformed into another array of discrete values
  2. Remove high frequencies
  3. Inverse Fourier transform (back out of the frequency domain)

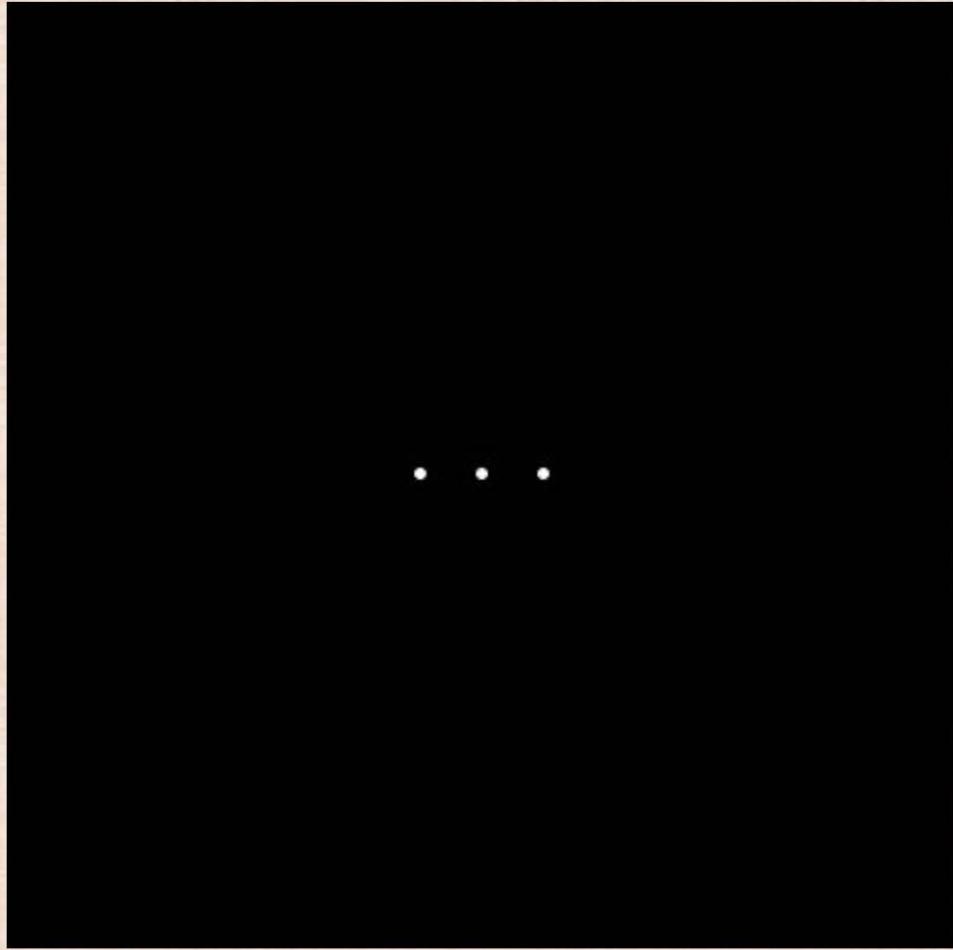
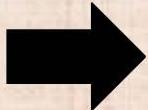
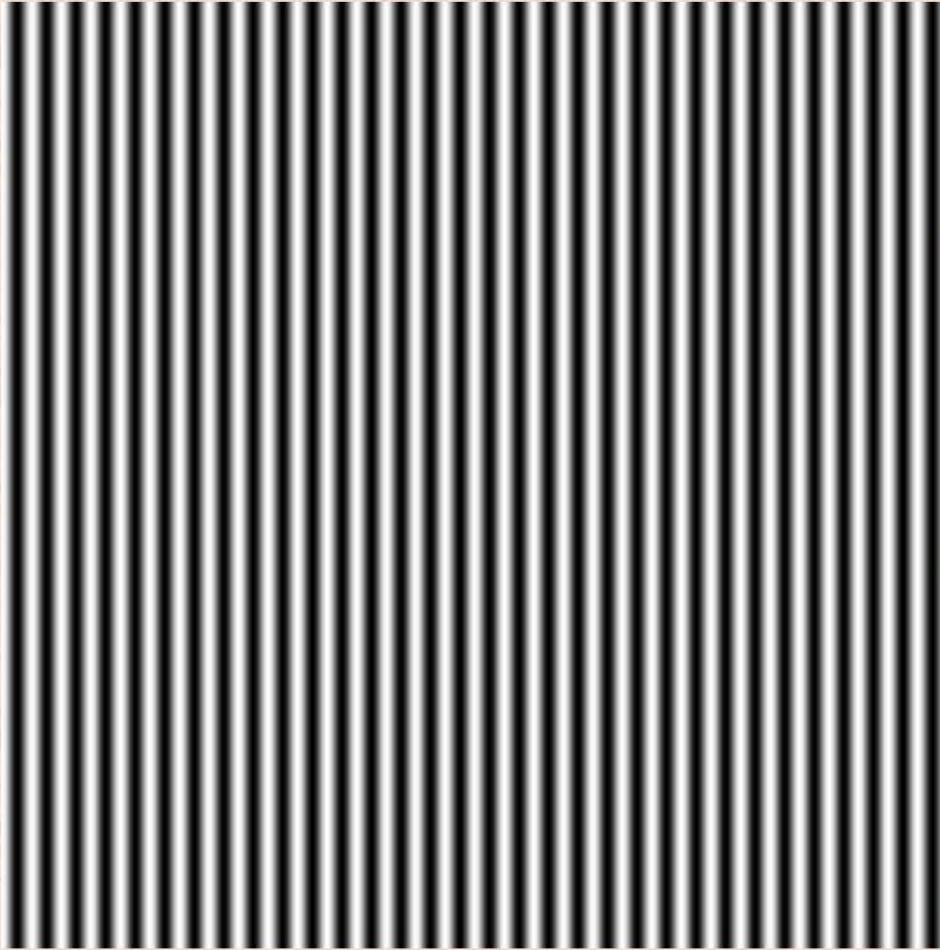
# Constant Function



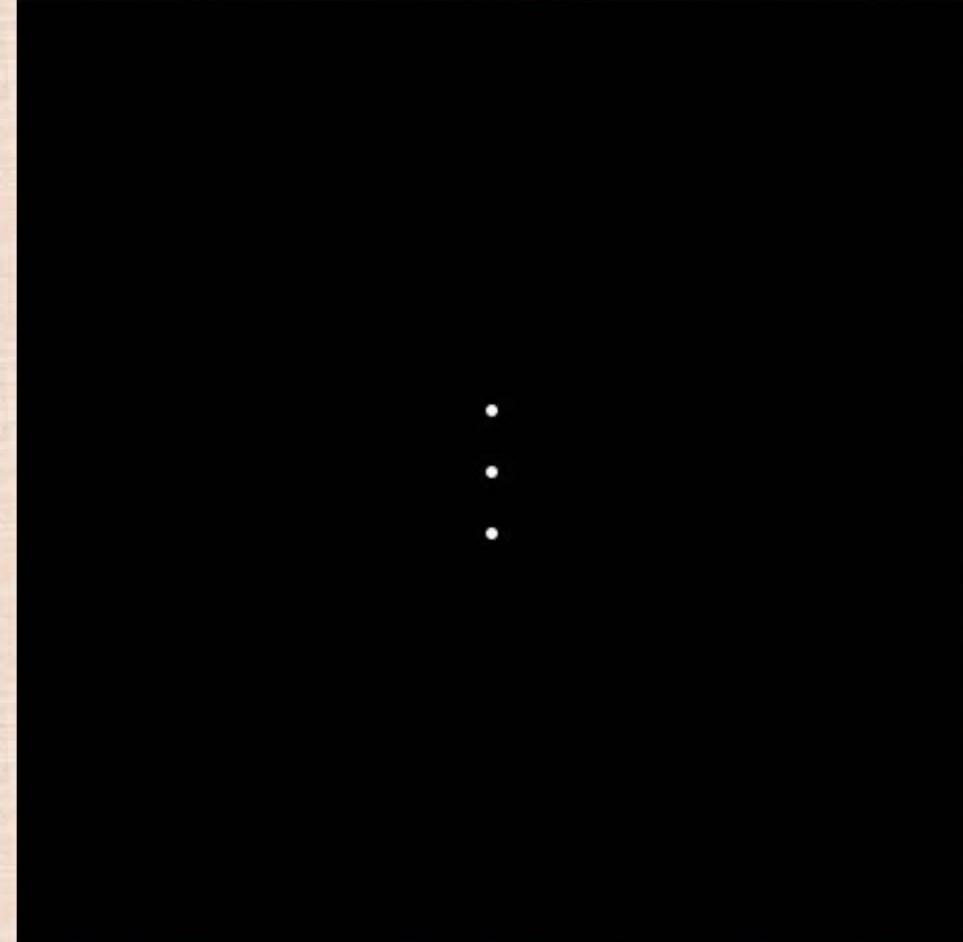
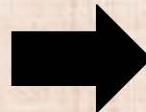
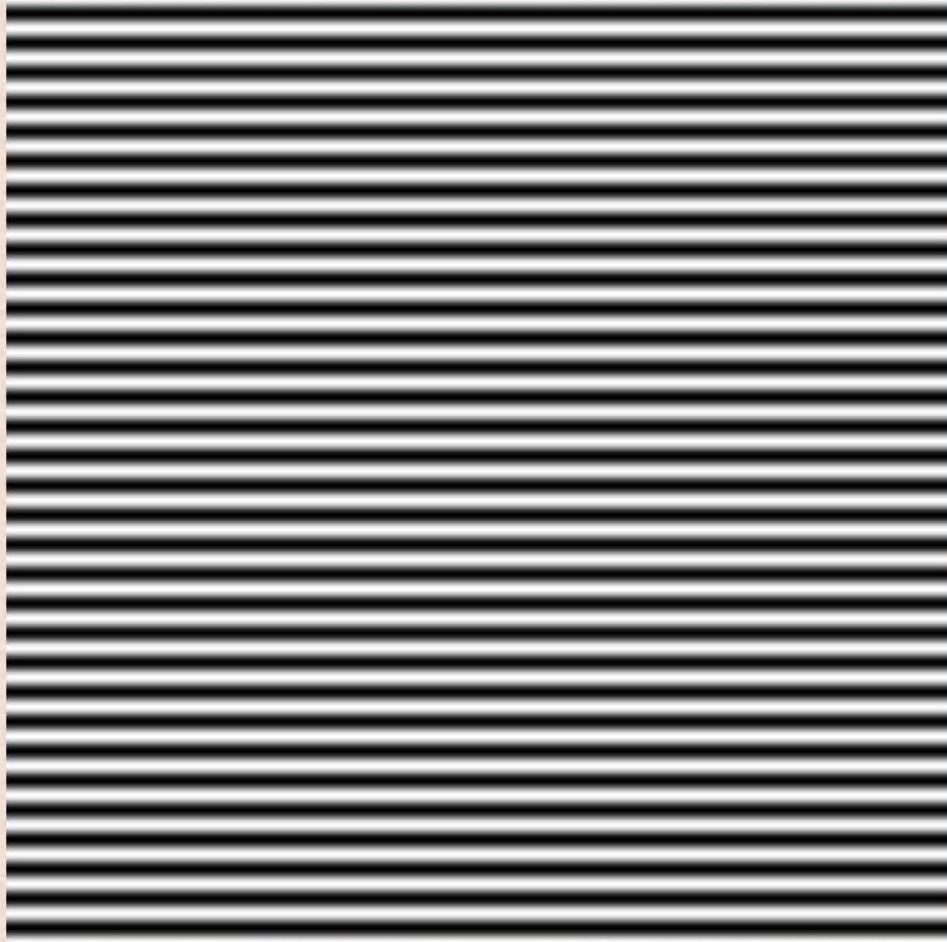
$$\sin(2\pi/32)x$$



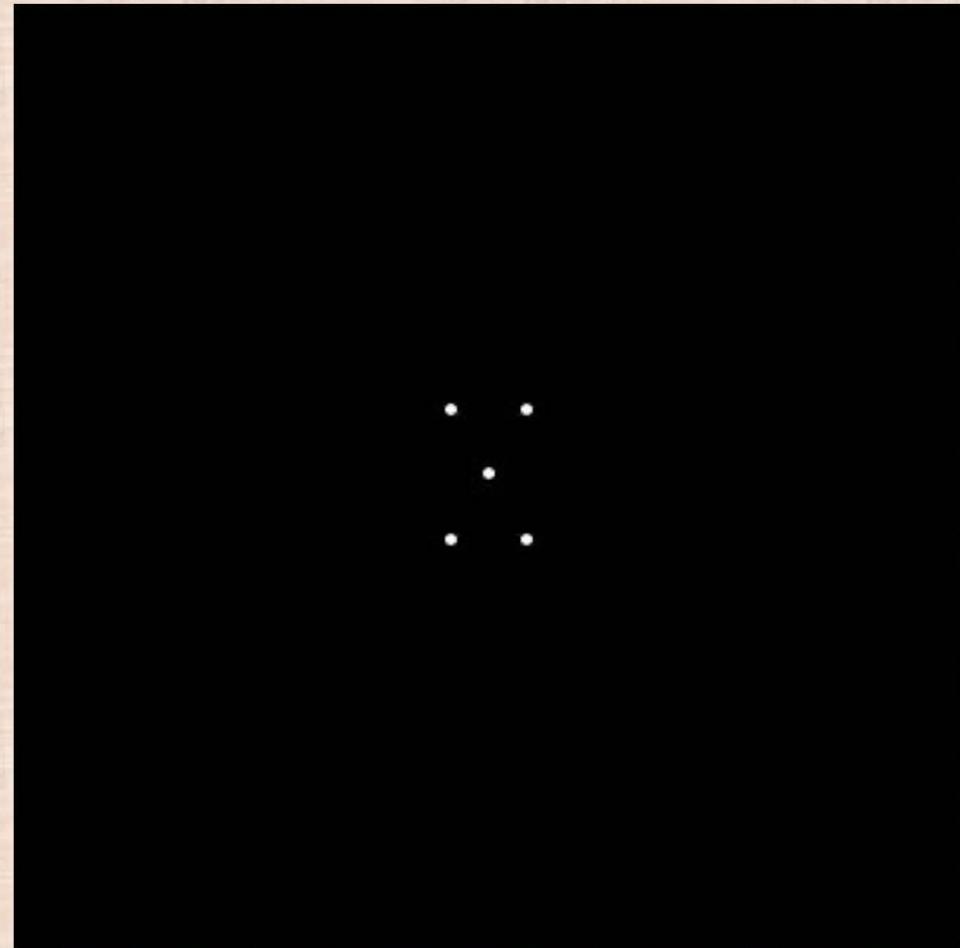
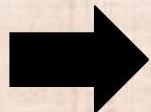
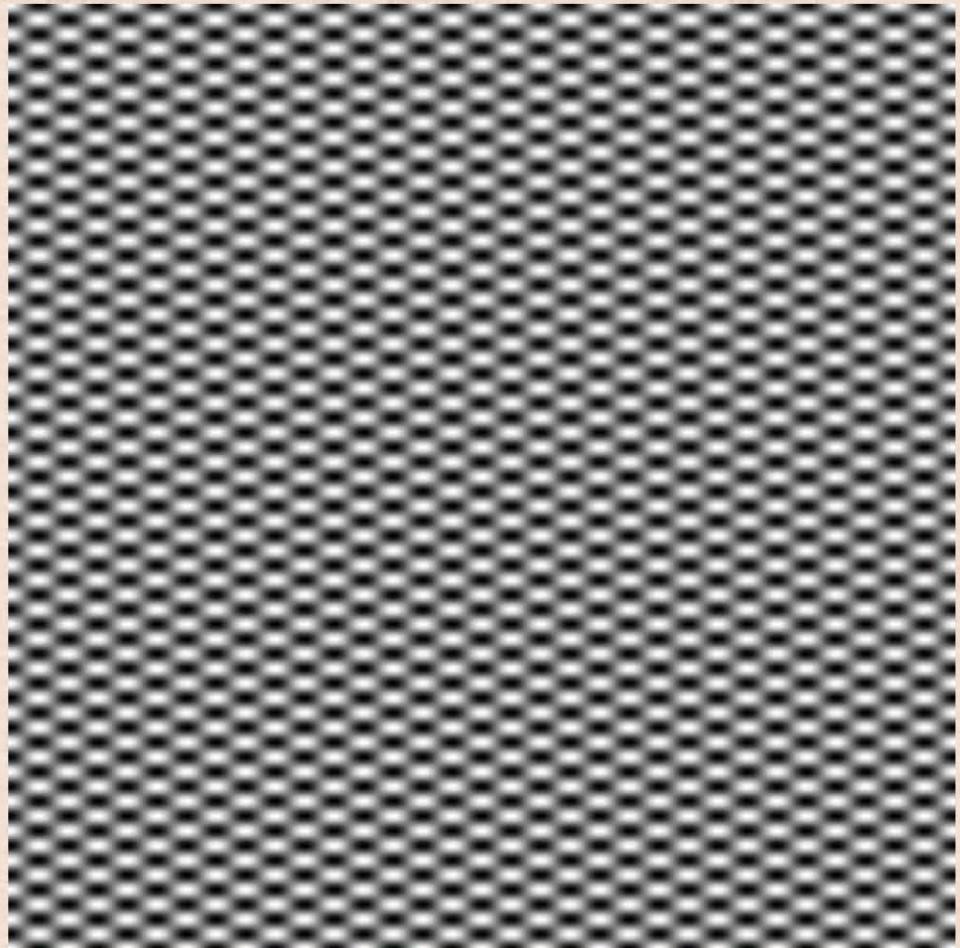
$$\sin(2\pi/16)x$$



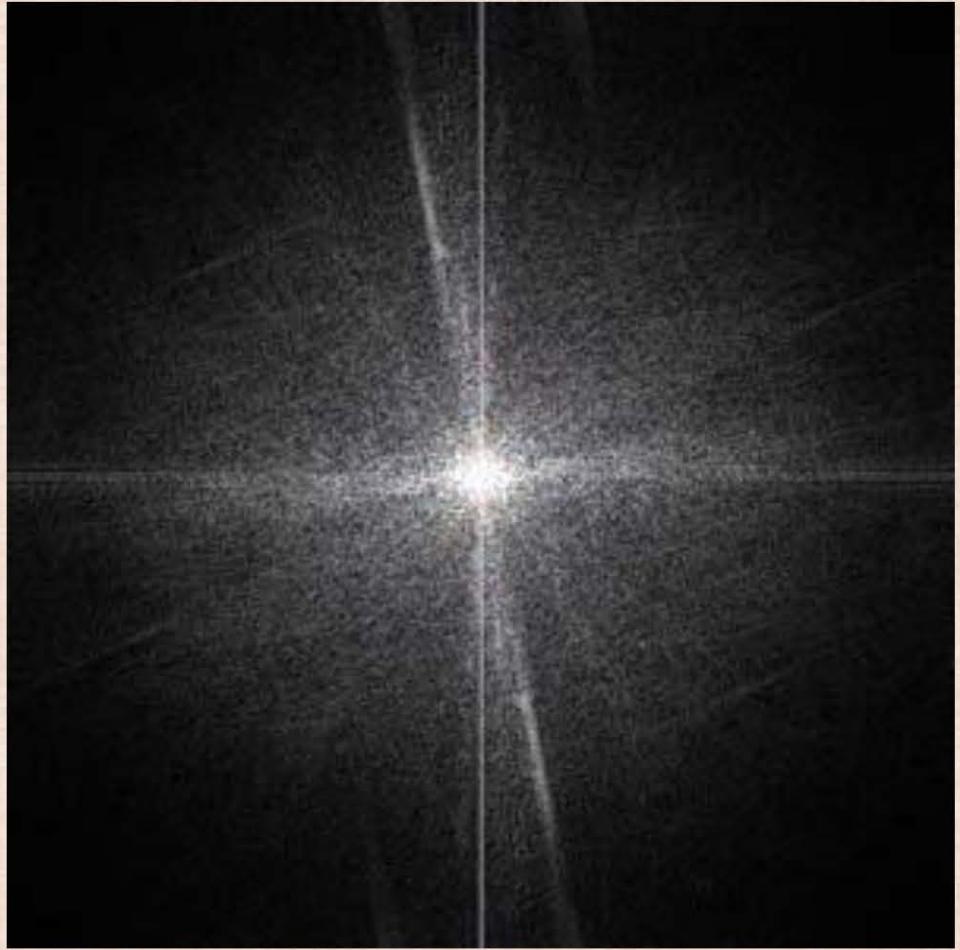
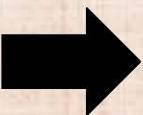
$$\sin(2\pi/16) y$$



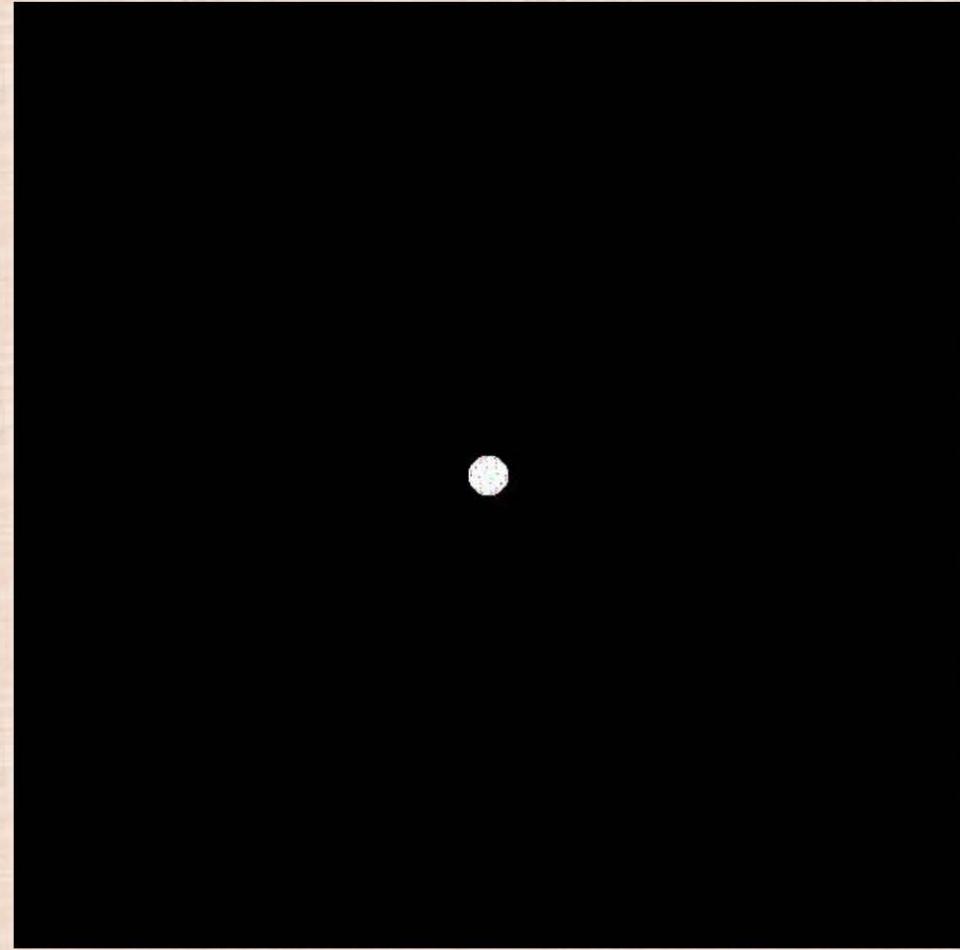
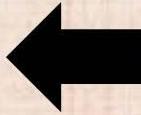
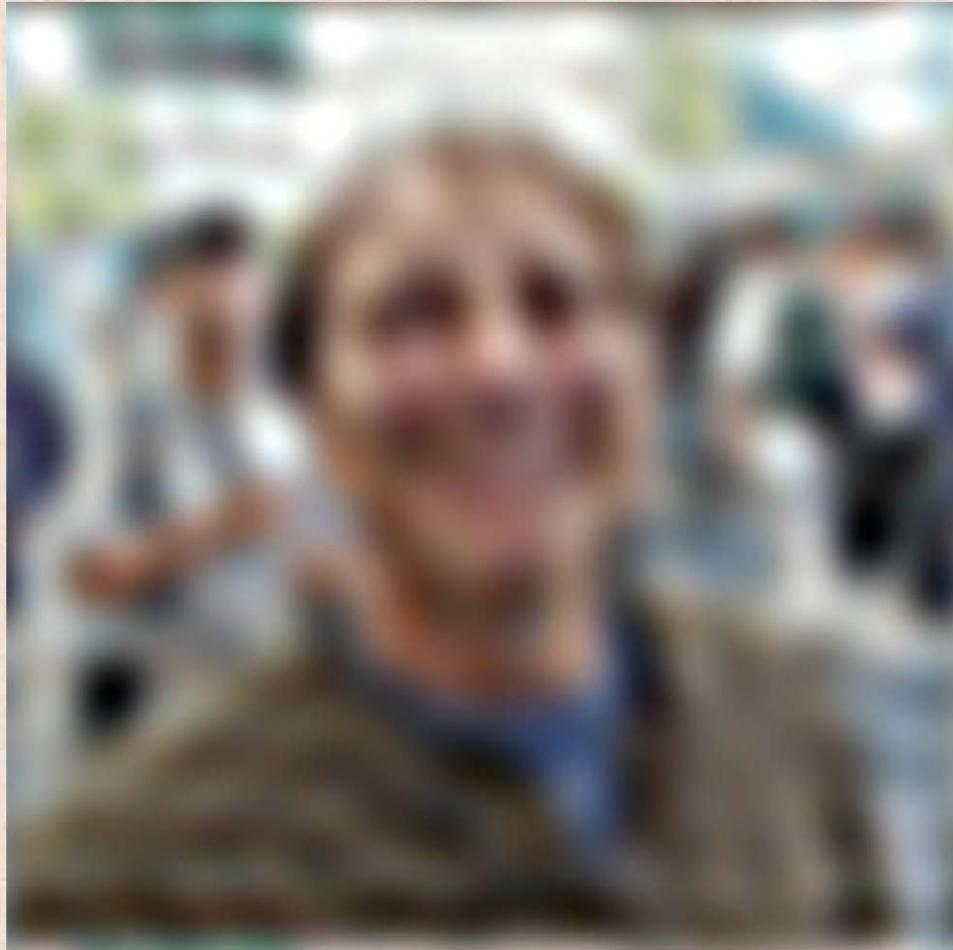
$$\sin(2\pi/32)x * \sin(2\pi/16)y$$



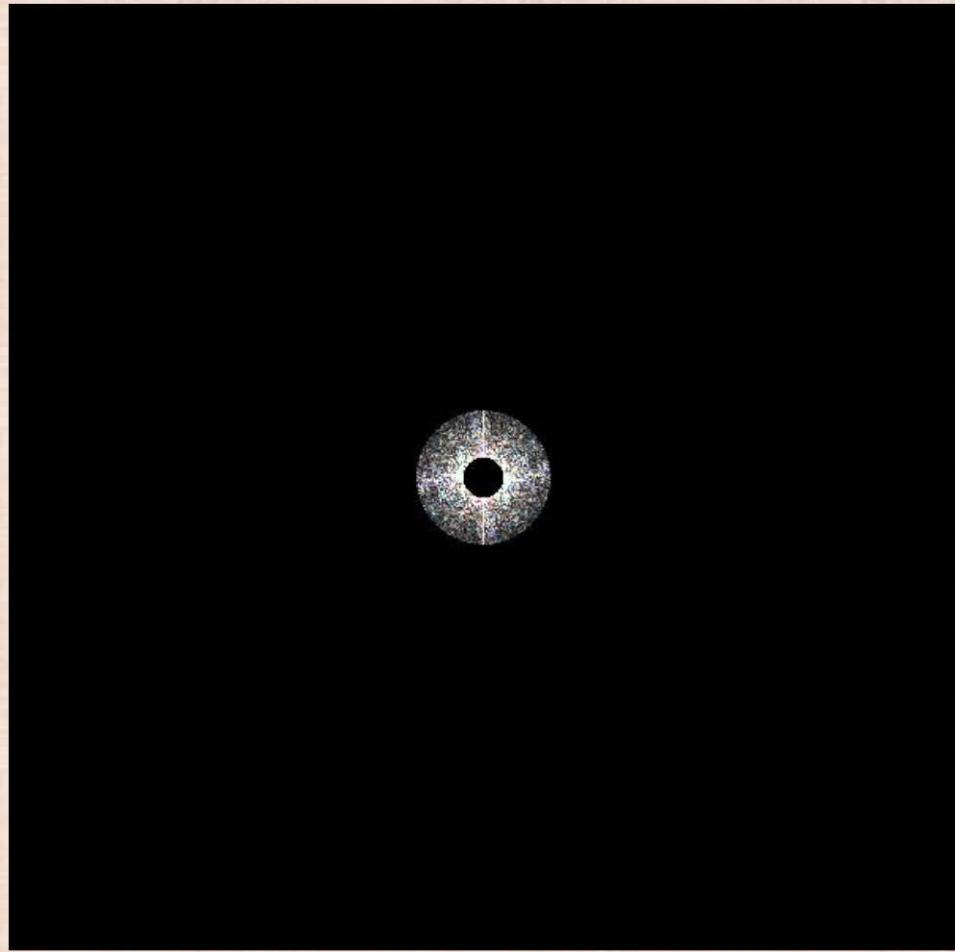
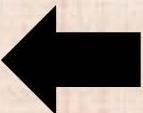
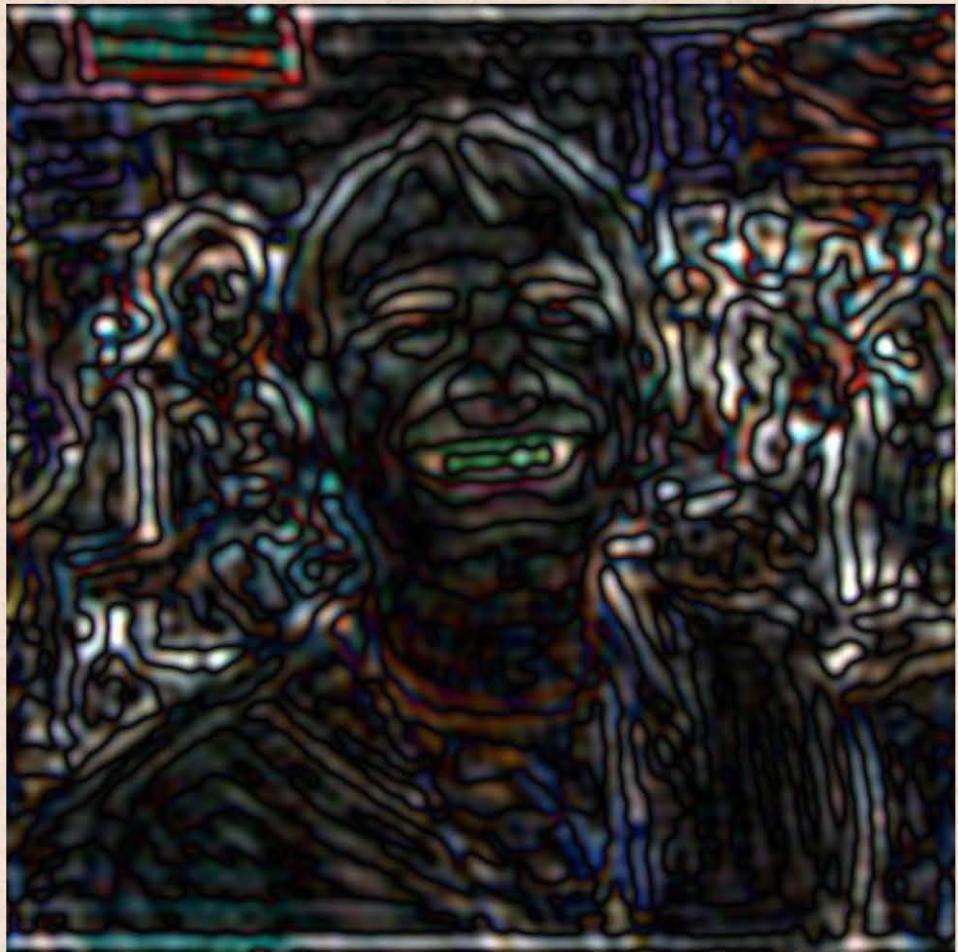
# An obvious star!



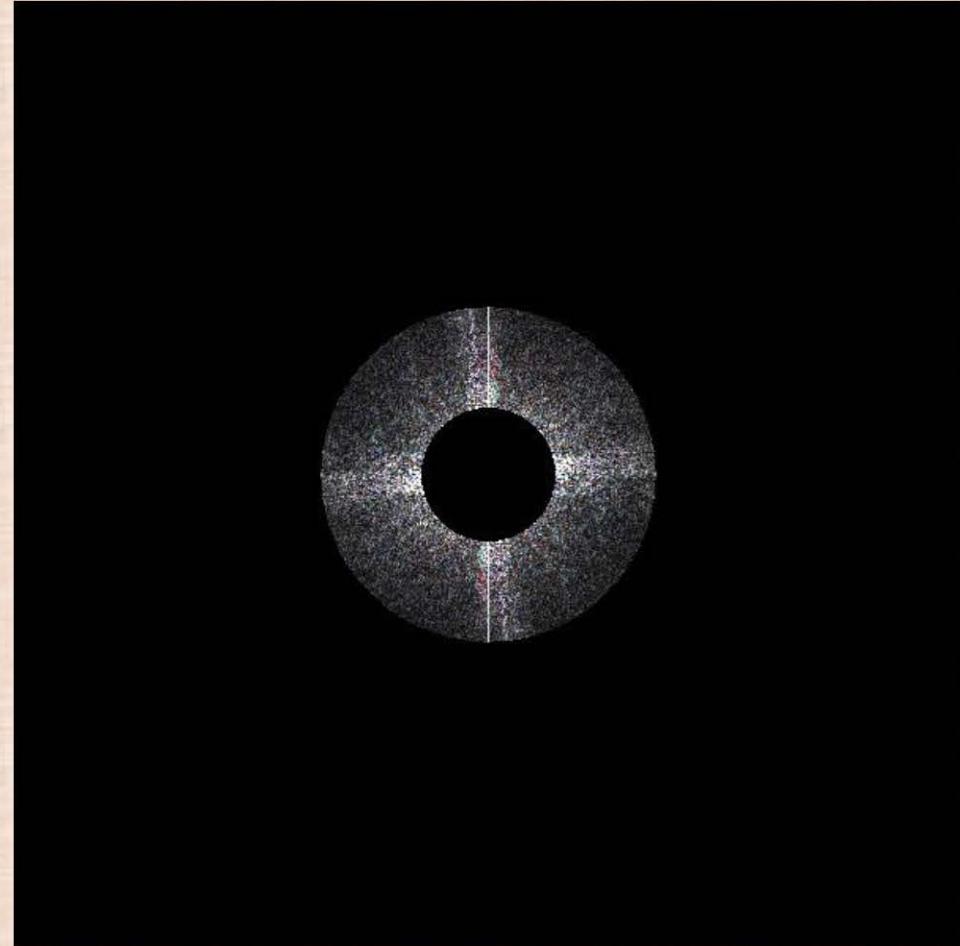
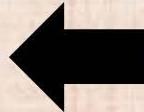
# lowest frequencies



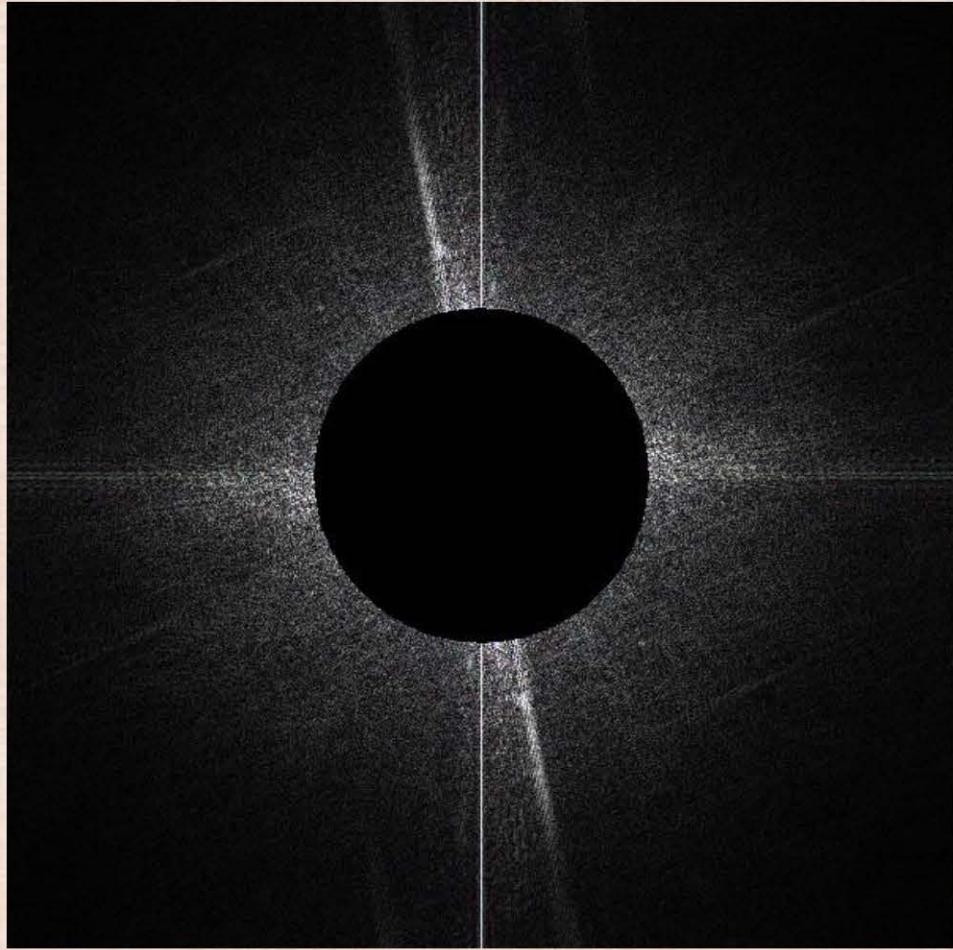
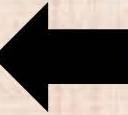
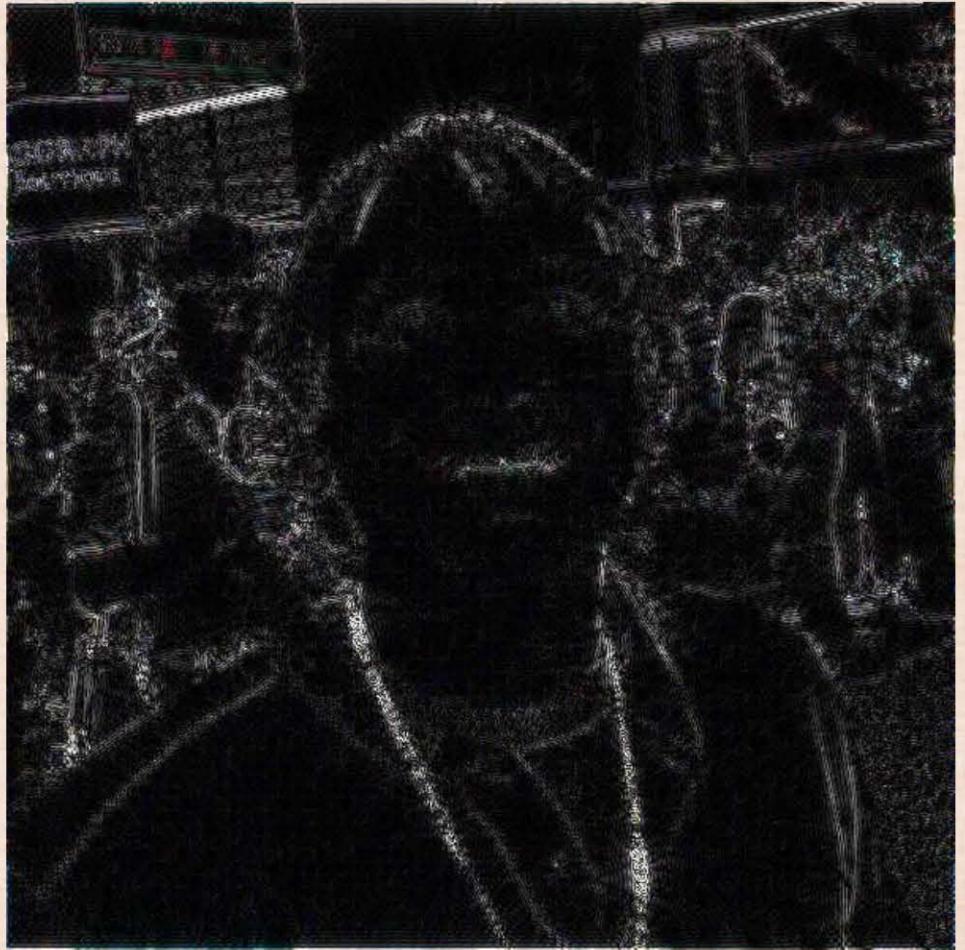
# intermediate frequencies



(larger) intermediate frequencies



# highest frequencies (edges)



# Convolution

- Let  $f$  and  $g$  be functions in the spatial domain (e.g. images), and  $F(f)$  and  $F(g)$  be transformations of  $f$  and  $g$  into the frequency domain
  - In our prior examples:  $f$  was the image (to the left),  $F(f)$  was the frequency domain version of the image (to the right)
- Removing higher frequencies of  $F(f)$  is equivalent to multiplying by a Heaviside function  $F(g)$  ( $=1$  for smaller frequencies,  $=0$  for larger frequencies)
- Then, the inverse transform  $F^{-1}(F(f)F(g))$  gives the final result
- This entire process is called the **convolution** of  $f$  and  $g$ :

$$f * g = F^{-1}(F(f)F(g))$$

# Convolution Integral

- Convolution can be achieved without the Fourier Transform:

$$(f * g)(t) = \int_{-\infty}^{\infty} f(\tau)g(t - \tau)d\tau = \int_{-\infty}^{\infty} f(t - \tau)g(\tau)d\tau$$

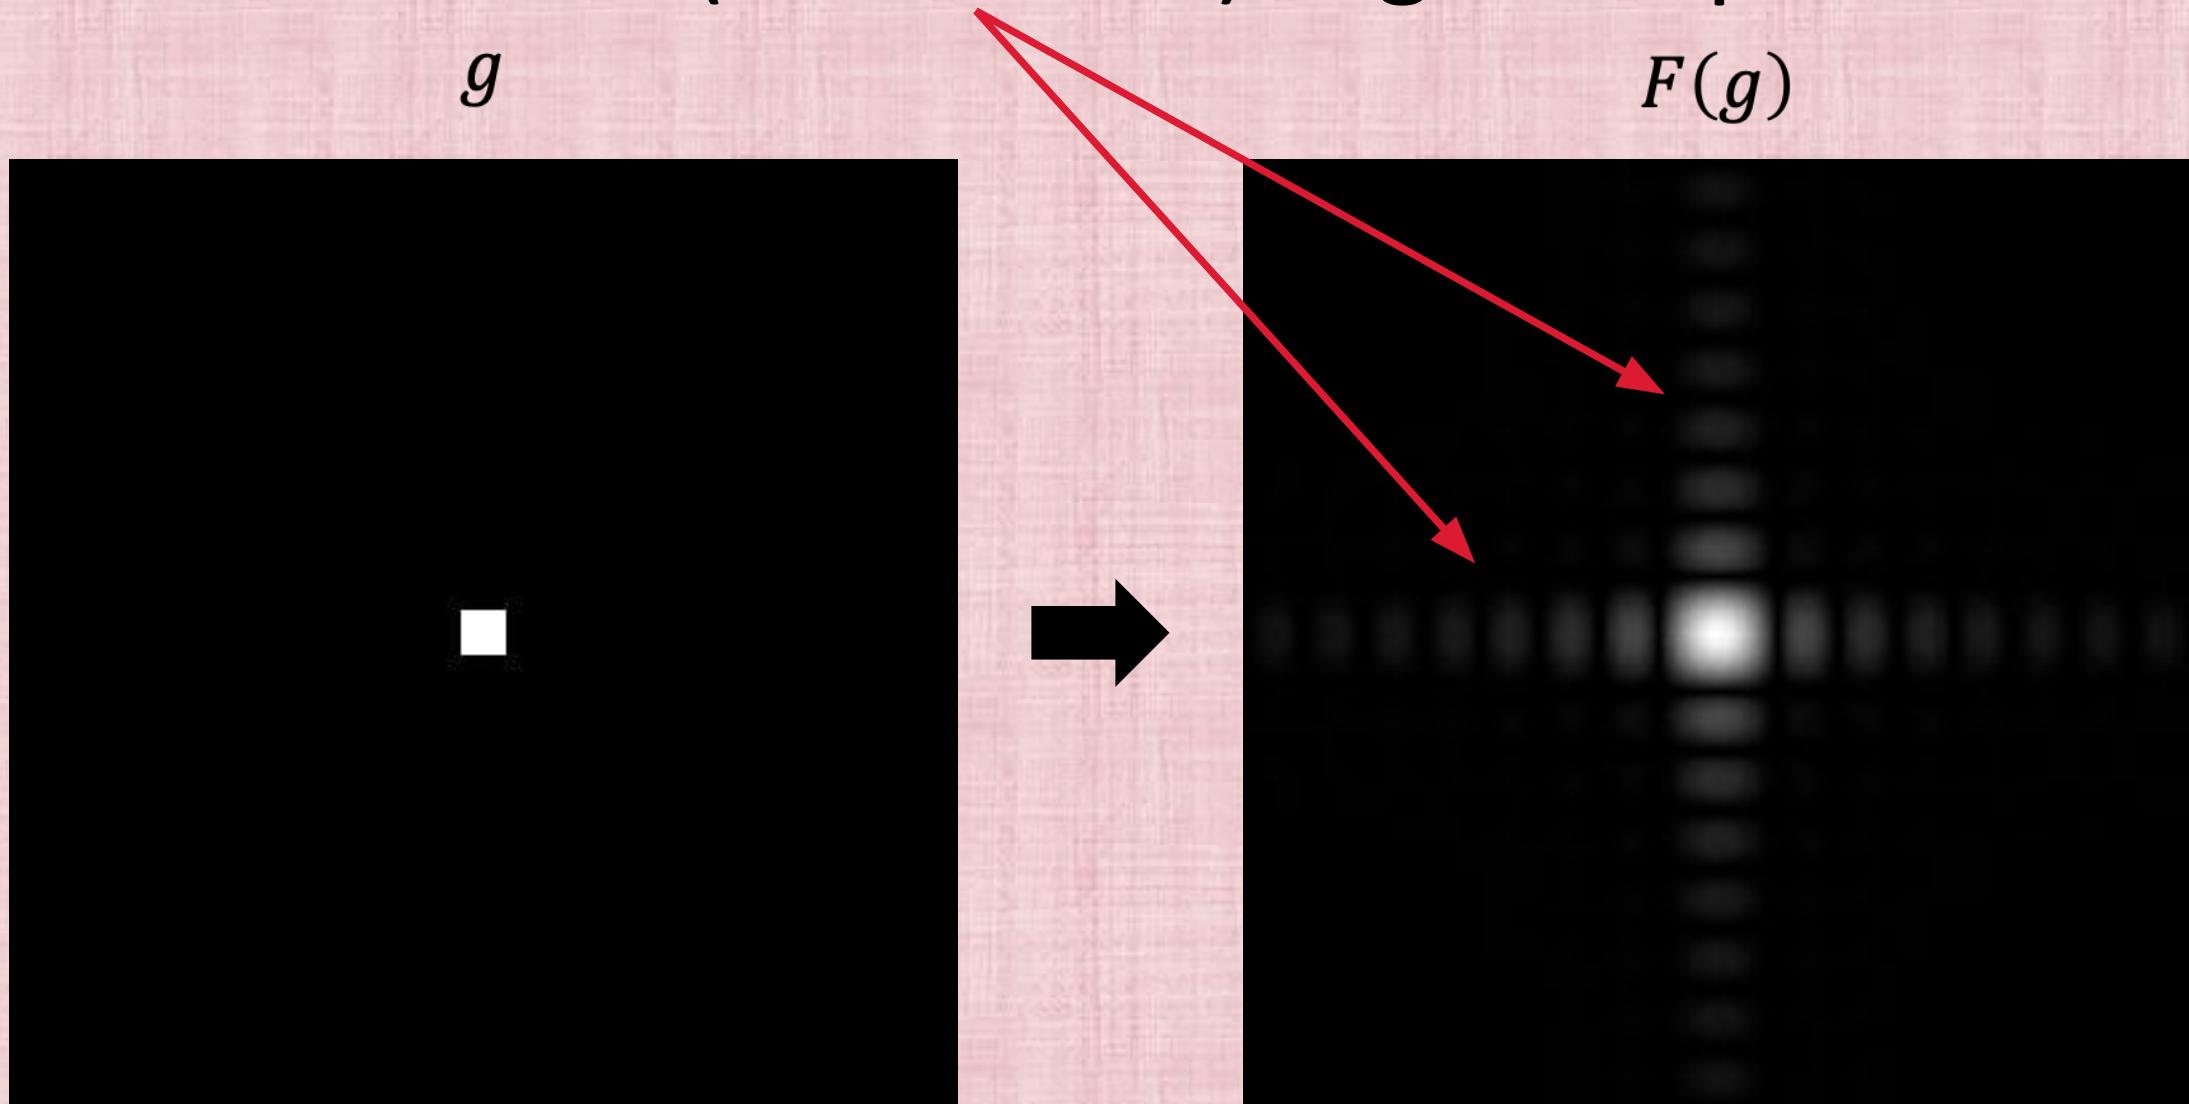
- A narrower  $g$  makes the integral more efficient to compute
- A narrower  $F(g)$  better removes high frequencies
- But, they can't both be narrow
  - Recall: the narrower Gaussian had wider frequencies, and the wider Gaussian had narrower frequencies

# Box Filter

- Let  $g$  have nonzero values in an  $N \times N$  block of pixels (surrounding the origin), and be zero elsewhere
- The discrete convolution (integral) is computed by:
  - overlay the filter  $g$  on the image, multiply the corresponding entries, and sum the results
- The final result is (typically) defined at the center of the filter

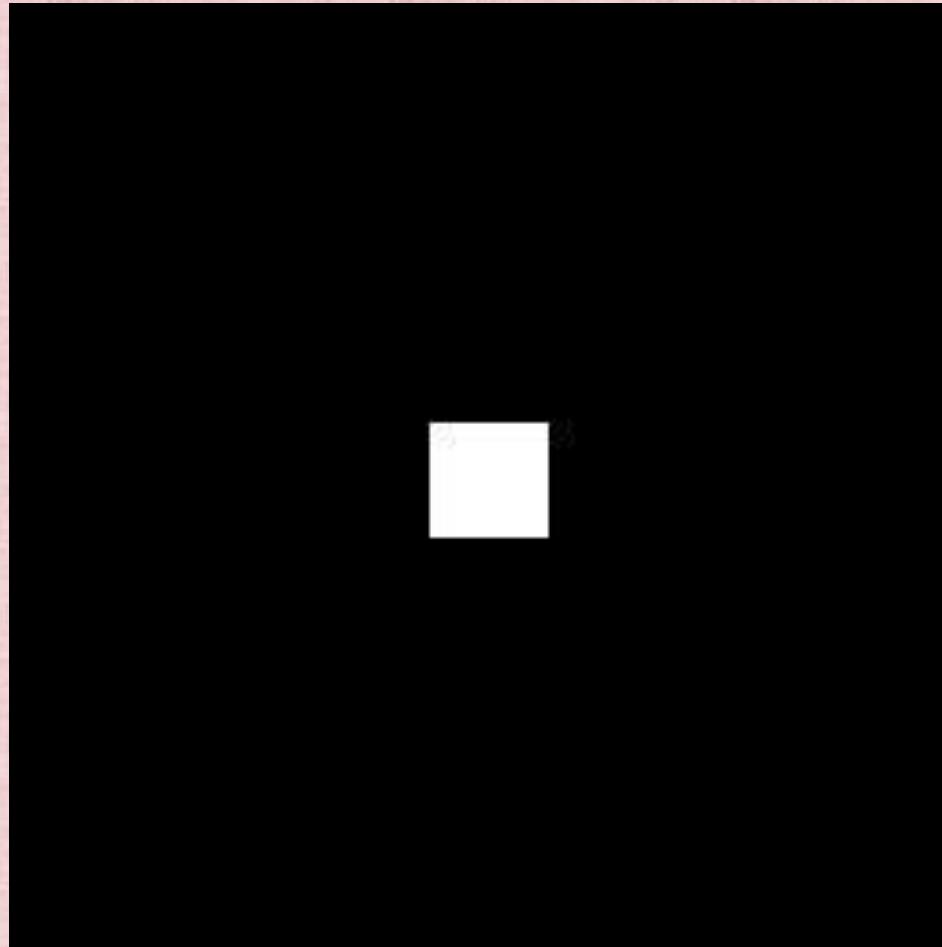
$\frac{1}{16}$	$\frac{1}{16}$	$\frac{1}{16}$	$\frac{1}{16}$
$\frac{1}{16}$	$\frac{1}{16}$	$\frac{1}{16}$	$\frac{1}{16}$
$\frac{1}{16}$	$\frac{1}{16}$	$\frac{1}{16}$	$\frac{1}{16}$
$\frac{1}{16}$	$\frac{1}{16}$	$\frac{1}{16}$	$\frac{1}{16}$

# Filters Most (but not all) High Frequencies



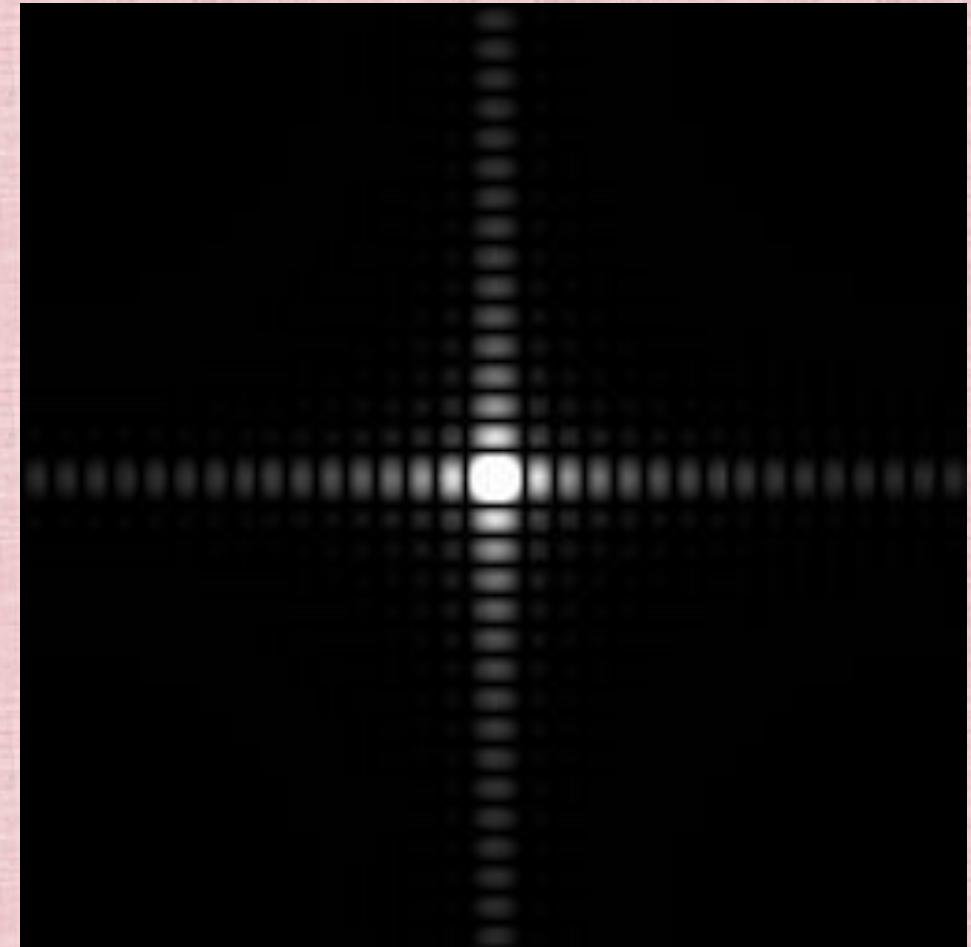
# Wider Box Filter

$g$



more expensive convolution integral

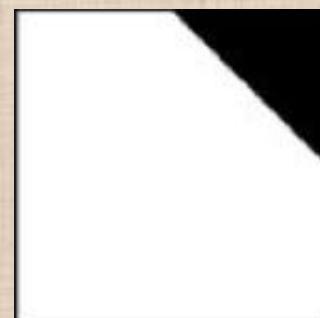
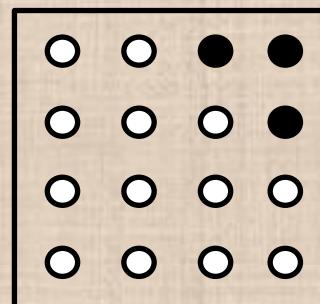
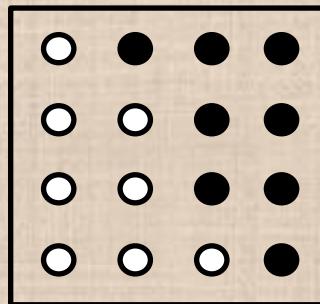
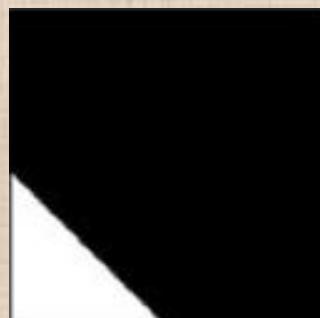
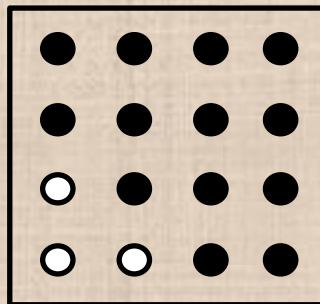
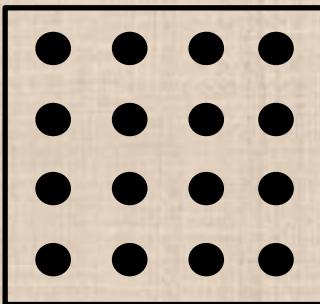
$F(g)$



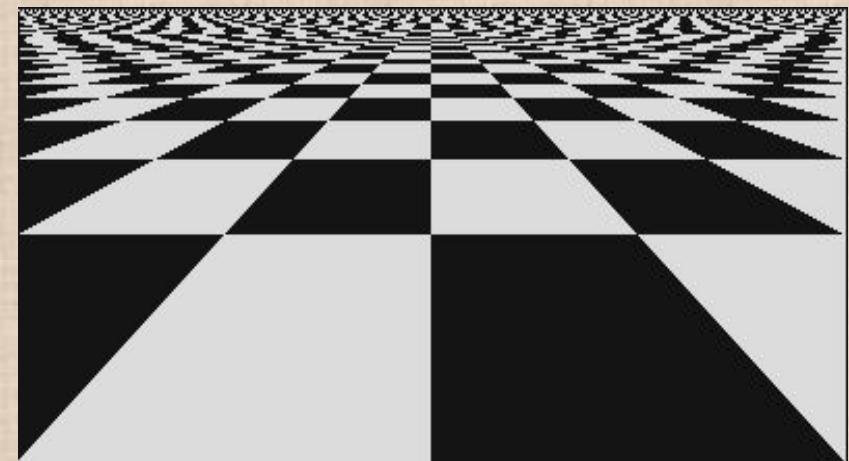
removes more of the high frequencies

# Super-Sampling

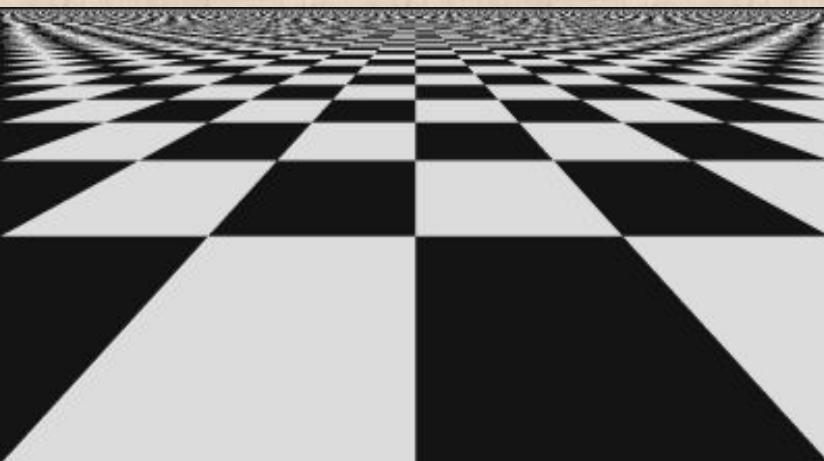
- Collect extra information/samples (in each pixel), and average the result (e.g. with a box filter)
  - E.g. render a 100 by 100 image with 4 by 4 super-sampling (equivalent to rendering a 400 by 400 image)
  - This properly represents (without aliasing) frequencies up to 4 times higher (than the original image)
  - Apply a 4 by 4 box filter to remove as much of those extra frequencies as possible
- As the number samples per pixel increases, this converges to the area coverage integral
  - Computational Cost: only super-sample pixels that have high frequencies (e.g., edges)
  - N.B. use pseudo-random Monte-Carlo super-sampling strategies (instead of uniform n by n super-sampling)



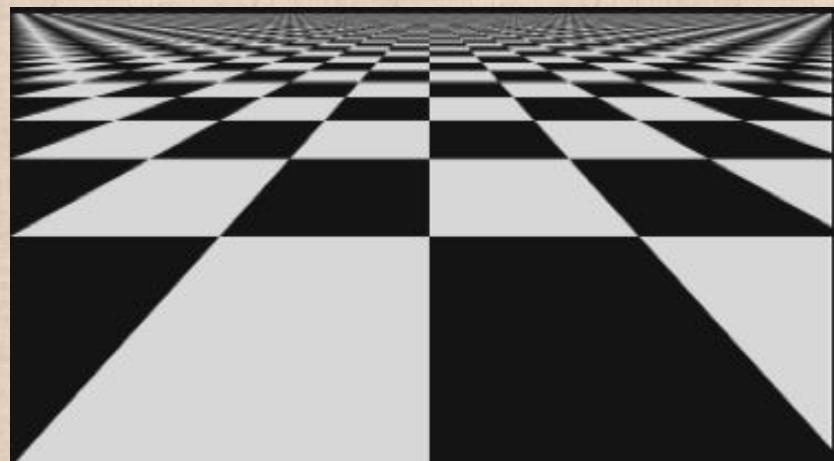
# Super-Sampling



Point Sampling

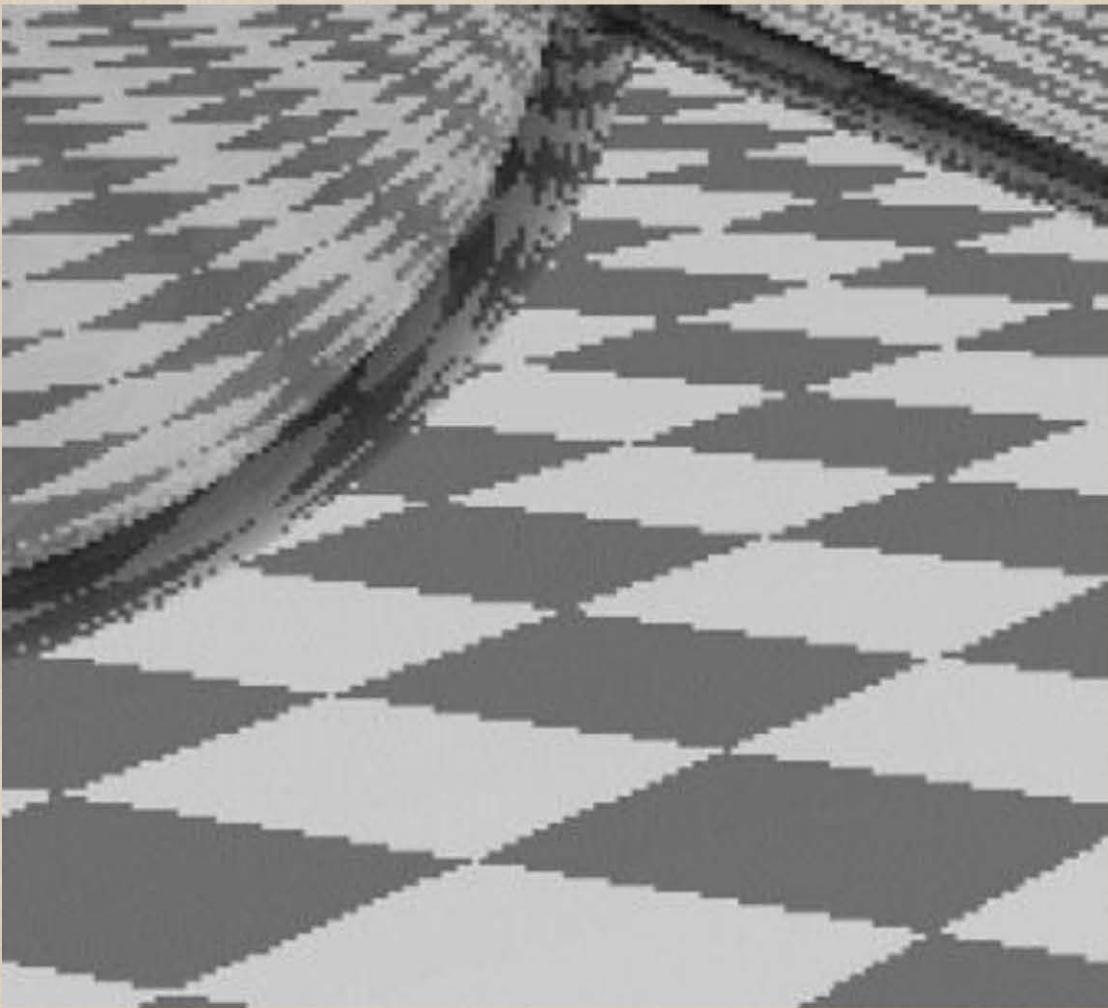


4 by 4 Super-Sampling

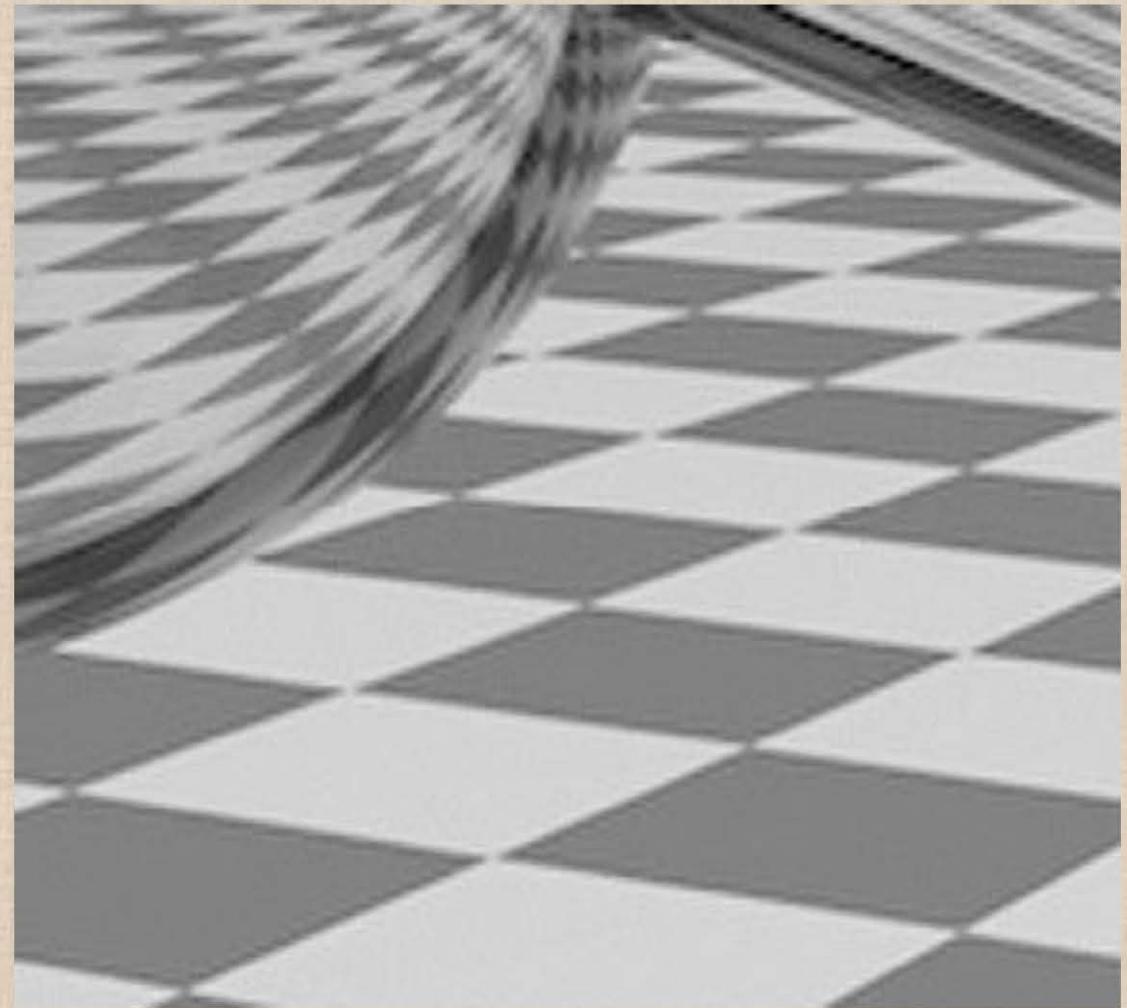


Exact Area Coverage

# Super-Sampling



Jaggies



Anti-Aliased